

3. Making Prediction: Decoding and Inference

3.1. Linear generative models: \Rightarrow PFSA

- Decoding: **Viterbi**,
- Inference: **Forward**

3.2. Linear discriminative models: \Rightarrow CRFs

- Decoding: **Viterbi**,
- Inference: **Forward**

3.3. Non-linear generative models: \Rightarrow PCFGs

- Decoding: probability of most likely parse tree: **Viterbi**
- Inference: probability of a string: **Inside**

3.4. Word Graph based N-best search

- Word graphs
- N-best search

➤ Probability of a path

$$P_A(x, \pi) = I(s_0) \cdot \left(\prod_{i=1}^k P(s_{i-1}, x_i, s_i) \right) \cdot F(s_k)$$

➤ Probability of a string

$$P_A(x) = \sum_{\pi \in \Pi_A(x)} P_A(x, \pi)$$

➤ Probability of the best path

$$\widehat{P}_A(x) = \max_{\pi \in \Pi_A(x)} P_A(x, \pi)$$

➤ Best path

$$\widehat{\pi}(x) = \arg \max_{\pi \in \Pi_A(x)} P_A(x, \pi)$$

PFSA: VITERBI ALGORITHM

- **Definition:** The probability of generating the prefix x_1^t through the best path and reaching state q is:

$$\gamma_x(t, q) \stackrel{\text{def}}{=} \max_{\substack{s_0^t \in \Pi_A(x_1^t): \\ s_t = q}} I(s_0) \cdot \left(\prod_{i=1}^t P(s_{i-1}, x_i, s_i) \right) \quad \forall q \in Q, 0 \leq t \leq |x|$$

- **Initialization:** $\gamma_x(0, q) = I(q) \quad \forall q \in Q$

- **Recursion:** $\forall q \in Q, 1 \leq t \leq |x|$


$$\gamma_x(t, q) = \max_{q' \in Q} \{ \gamma_x(t-1, q') \cdot P(q', x_t, q) \}$$

- **Final result:** The probability of the best state sequence for x given A is:

$$\widehat{P}_A(x) = \max_{q \in Q} \{ \gamma_x(|x|, q) \cdot F(q) \}$$

Backpointers can be used to recover the optimal sequence of states

ALGORITMO DE VITERBI: EJEMPLO



PFSA: FORWARD ALGORITHM

➤ **Definition:** The probability of generating the prefix x_1^t and reaching state q is:

$$\alpha_x(t, q) \stackrel{\text{def}}{=} \sum_{\substack{s_0 \in \Pi_A(x_1^t): \\ s_t = q}} I(s_0) \cdot \left(\prod_{i=1}^t P(s_{i-1}, x_i, s_i) \right) \quad \forall q \in Q, 0 \leq t \leq |x|$$

➤ **Initialization:** $\alpha_x(0, q) = I(q) \quad \forall q \in Q$

➤ **Recursion:** $\forall q \in Q, 1 \leq t \leq |x|$

$$\alpha_x(t, q) = \sum_{q' \in Q} \alpha_x(t-1, q') \cdot P(q', x_t, q)$$

➤ **Final result:** The probability of string x is:

$$P_A(x) = \sum_{q \in Q} \alpha_x(|x|, q) \cdot F(q)$$

PFSA: BACKWARD ALGORITHM

➤ **Definition:** The probability of generating the suffix $x_{t+1}^{|x|}$ from the state q is:

$$\beta_x(t, q) \stackrel{\text{def}}{=} \sum_{\substack{s_t \in \Pi_A(x_{t+1}^{|x|}): \\ s_t = q}} \left(\prod_{i=t+1}^{|x|} P(s_{i-1}, x_i, s_i) \right) \cdot F(s_{|x|}) \quad \forall q \in Q, 0 \leq t \leq |x|$$

➤ **Initialization:** $\beta_x(|x|, q) = F(q) \quad \forall q \in Q$

➤ **Recursion:** $\forall q \in Q, 0 \leq t \leq |x|-1$

$$\beta_x(t, q) = \sum_{q' \in Q} \beta_x(t+1, q') \cdot P(q, x_{t+1}, q')$$

➤ **Final result:** The probability of string x is:

$$P_A(x) = \sum_{q \in Q} I(q) \cdot \beta_x(0, q)$$

CONDITIONAL RANDOM FIELDS

Given $x = x_1, x_2, \dots, x_T \in \mathcal{X}^*$ and $y = y_1, y_2, \dots, y_T \in \mathcal{Y}^*$

Discriminative models: Conditional Random Fields (CRF)

[Lafferty, McCallum, Pereira, 2001].

$$p(y|x; \theta) = \frac{\exp\{\sum_{t=1}^T \sum_{k=1}^K \theta_k f_k(y_{t-1}, y_t, x_t)\}}{Z(x; \theta)} \quad (4)$$

Where

$$Z(x; \theta) = \sum_{y' \in \mathcal{Y}^*} \exp \left\{ \sum_{t=1}^T \sum_{k=1}^K \theta_k f_k(y'_{t-1}, y'_t, x_t) \right\} \quad (5)$$

CRFs AS FACTOR GRAPHS

We can define a transition (factor) as:

$$\Psi_t(y_{t-1}, y_t, x_t) \stackrel{\text{def}}{=} \exp \left\{ \sum_{k=1}^K \theta_k f_k(y_{t-1}, y_t, x_t) \right\}$$

From (4) CRFs can be defined as:

$$p(y|x; \theta) = \frac{1}{Z(x; \theta)} \prod_{t=1}^T \Psi_t(y_{t-1}, y_t, x_t)$$

Also from (5)

$$Z(x; \theta) = \sum_{y' \in \mathcal{Y}^*} \prod_{t=1}^T \Psi_t(y'_{t-1}, y'_t, x_t)$$

CRFs AS FACTOR GRAPHS

➤ **Decoding:** Given θ and x , predict the best output sequence for x .

$$\hat{y} = \arg \max_y p(y|x; \theta) = \arg \max_y \prod_{t=1}^T \Psi_t(y_{t-1}, y_t, x_t)$$

➤ The probability of an output sequence

$$p(y|x; \theta) = \frac{1}{Z(x; \theta)} \prod_{t=1}^T \Psi_t(y_{t-1}, y_t, x_t)$$

DECODING WITH CRFs: VITERBI ALGORITHM

➤ **Definition:** Score of optimal sequence for $x_1 \dots x_t$ ending at $y_t = s \in \mathcal{Y}$

$$\gamma_t(s) \stackrel{\text{def}}{=} \max_{y_1^t; y_t=s} \prod_{i=1}^t \Psi_i(y_{i-1}, y_i, x_i)$$

➤ **Initialization:** $\forall s \in \mathcal{Y}$

$$\gamma_1(s) = \Psi_1(y_0 = \text{null}, y_1 = s, x_1)$$

➤ **Recursion:** $\forall t = 2 \dots T$; and $\forall s \in \mathcal{Y}$

$$\gamma_t(s) = \max_{s' \in \mathcal{Y}} \left\{ \gamma_{t-1}(s') \cdot \Psi_t(y_{t-1} = s', y_t = s, x_t) \right\}$$

➤ **Final result:** The optimal score for x is:

$$\max_s \gamma_T(s)$$

Backpointers can be used to recover the optimal sequence \hat{y}

INFERENCE WITH CRFs: FORWARD ALGORITHM

➤ **Definition:** Score for $x_1 \dots x_t$ ending at $y_t = s \in \mathcal{Y}$

$$\alpha_t(s) \stackrel{\text{def}}{=} \sum_{y_1^t; y_t=s} \prod_{i=1}^t \Psi_i(y_{i-1}, y_i, x_i)$$

➤ **Initialization:** $\forall s \in \mathcal{Y}$

$$\alpha_1(s) = \Psi_1(y_0 = \text{null}, y_1 = s, x_1)$$

➤ **Recursion:** $\forall t = 2 \dots T$; and $\forall s \in \mathcal{Y}$

$$\alpha_t(s) = \sum_{s' \in \mathcal{Y}} \alpha_{t-1}(s') \cdot \Psi_t(y_{t-1} = s', y_t = s, x_t)$$

➤ **Final result:** The optimal score for x is:

$$\sum_s \alpha_T(s)$$

CRFs AS FACTOR GRAPHS

➤ The probability of an output sequence

$$p(y|x; \theta) = \frac{1}{Z(x; \theta)} \prod_{t=1}^T \Psi_t(y_{t-1}, y_t, x_t)$$

Where we can compute $Z(x; \theta)$ efficiently using the forward algorithm

$$Z(x; \theta) = \sum_{y_1^T} \prod_{t=1}^T \Psi_t(y'_{t-1}, y'_t, x_t) = \sum_s \alpha_T(s)$$

PROBABILISTIC CONTEXT-FREE GRAMMARS

Generative models: Probabilistic Context-Free Grammars

$$P_{\theta}(x, t_x) = \prod_{i=1}^m p(r_i) \quad \text{where } t_x : r_1, \dots, r_m$$

➤ Probability of an observation

$$P_{\theta}(x) = \sum_{t_x \in T_x} P_{\theta}(x, t_x)$$

➤ Probability of the best parse tree

$$\hat{P}_{\theta}(x) = \max_{t_x \in T_x} P_{\theta}(x, t_x)$$

➤ Finding the best parse

$$\hat{t}_x = \arg \max_{t_x \in T_x} P_{\theta}(x, t_x)$$

PCFG: VITERBI ALGORITHM

➤ **Definition:** Given $x = x_1 \dots x_T \in \Sigma^*$ and $A \in N$

$$\hat{e}(A, i, i+l) \stackrel{\text{def}}{=} \hat{P}_{\theta}(A \Rightarrow^* x_{i+1} \dots x_{i+l})$$

➤ **Initialization:** $\forall A \in N; \quad \forall i : 0 \dots T-1;$

$$\hat{e}(A, i, i+1) = p(A \rightarrow b) \cdot \delta(b, x_{i+1})$$

➤ **Recursion:** $\forall A \in N; \quad \forall l : 2 \dots T; \quad \forall i : 0 \dots T-l;$

$$\hat{e}(A, i, i+l) = \max_{B, C \in N} \left\{ p(A \rightarrow BC) \cdot \max_{k=1, \dots, l-1} \{ \hat{e}(B, i, i+k) \cdot \hat{e}(C, i+k, i+l) \} \right\}$$

➤ **Final result:** The probability of the best parse tree is:

$$\hat{P}_{\theta}(x) = \hat{e}(S, 0, T)$$

Backpointers can be used to recover the optimal sequence \hat{y}

PCFG: INSIDE ALGORITHM

➤ **Definition:** Given $x = x_1 \dots x_T \in \Sigma^*$ and $A \in N$

$$e(A, i, i+l) \stackrel{\text{def}}{=} P_{\theta}(A \Rightarrow^* x_{i+1} \dots x_{i+l})$$

➤ **Initialization:** $\forall A \in N; \quad \forall i : 0 \dots T-1;$

$$e(A, i, i+1) = p(A \rightarrow b) \cdot \delta(b, x_{i+1})$$


➤ **Recursion:** $\forall A \in N; \quad \forall l : 2 \dots T; \quad \forall i : 0 \dots T-l;$

$$e(A, i, i+l) = \sum_{B, C \in N} \left\{ p(A \rightarrow BC) \cdot \sum_{k=1, \dots, l-1} e(B, i, i+k) \cdot e(C, i+k, i+l) \right\}$$


➤ **Final result:** The sentence probability is:

$$P_{\theta}(x) = e(S, 0, T)$$

WORD GRAPH BASED N-BEST SEARCH



WORD GRAPHS



WORD GRAPHS

➤ Why do we need word graphs?

- Trellis-based decoding is very expensive in time and space, even though pruning techniques were used.
- The Word Graph (WG) allow for a reduction of the search space to the most probable hypotheses.

➤ Word Graph (WG) definition

- A WG is a labeled weighted directed acyclic graph.
- A WG is a data structure that represents a large finite sample of word sequences very efficiently.
- A WG represents the interpretations with the highest posterior probabilities.
- The WG is (a pruned version of) the Viterbi search trellis obtained during recognition.

WORD GRAPHS: NORMALIZATION

- The WG obtained after the search process is not a probabilistic and consistent graph.
- However, for some WG-based operations, it is necessary to normalize the WG.
- There are various graph normalization criteria.

First, we need to define the forward and backward functions:

$$\alpha_x(t, q) = \begin{cases} I(q) & \text{if } t = 0 \\ \sum_{q'} \alpha_x(t-1, q') \cdot P(q', x_t, q) & \text{otherwise} \end{cases}$$

$$\beta_x(t, q) = \begin{cases} F(q) & \text{if } t = |x| \\ \sum_{q'} \beta_x(t+1, q') \cdot P(q, x_{t+1}, q') & \text{otherwise} \end{cases}$$

Finally, the following statement can be proven:

$$P_A(x) = \sum_{q \in Q} \alpha_x(|x|, q) \cdot F(q) = \sum_{q \in Q} I(q) \cdot \beta_x(0, q)$$

WORD GRAPHS NORMALIZATION: PROPERTIES

The normalized weight for an arc (q', q) can be calculated as:

$$\psi(q', q) = \frac{P(q', q) \cdot \beta(q)}{\beta(q')}$$

➤ Property 1: Well-formed

$$\sum_{q \in Q} \psi(q', q) = 1 \quad \forall q' \in (Q - F);$$


➤ Property 2: Consistency

$$\sum_{\phi \in \Phi} P(\phi) = 1,$$


where Φ are all the paths of the WG.

➤ Property 3: The weight distribution of the sentences in the WG is not changed.

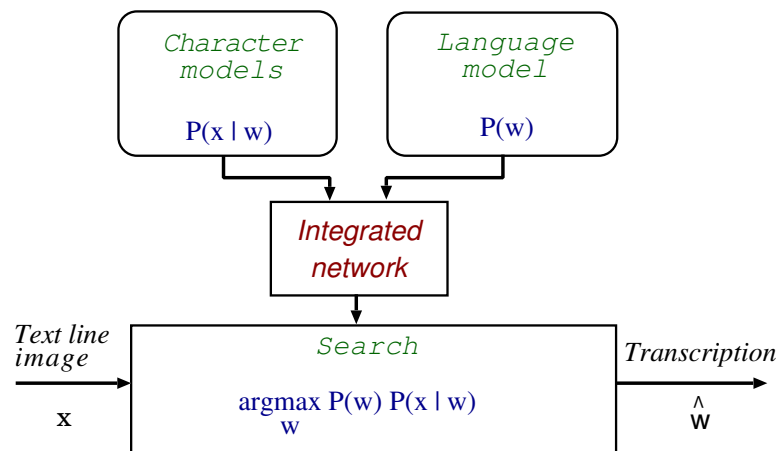
WORD GRAPHS: NORMALIZATION



WORD GRAPHS: NORMALIZATION



WORD GRAPHS: HTR INTEGRATION AND DECODIFICATION



Search engine:

THE VITERBI ALGORITHM (+ beam search + ...)

WORD GRAPHS: HTR INTEGRATION AND DECODIFICATION

Handwritten Text Recognition (HTR)

Given an input handwritten text image, find its most likely written transcription.

- Units combine into higher level units: **morphological** → **lexical** → **syntactics**.
- Relationships between levels can be modeled by weighted graphs.
- Recognition: find the best path in a suitable product graph.
- Morphological, lexical and syntactical are modelled by homogeneous PFS models.
- All these PFS models can be easily integrated into a single global (huge) PFS model.
- This global PFS model accepts sequences of raw feature vectors and outputs strings of recognized words.


M Connectionist Temporal Classification (CTC) and hybrid NN/HMMs.

L Pronunciation dictionary transducer mapping character transcriptions to sequences.

S Language model weighted automaton.

$$\min (\det (M \circ L \circ G))$$

WORD GRAPHS: HTR INTEGRATION AND DECODIFICATION



WORD GRAPHS BASED N-BEST SEARCH

Preliminaries

- The first step consists of calculating the probability of **the best path** from each state $q \in Q$ **to the final state** $f \in F$.

$$\Phi[q] = \max_{\pi \in P(q,f)} \{P(\pi) \cdot \rho(f)\}$$

- We consider pairs (q, p) of a state $q \in Q$ and a cumulative probability p .
- The algorithm uses a **priority queue** S containing the set of pairs (q, p) to examine next. The queue's ordering is based on Φ and defined by:

$$(q, p) > (q', p') \iff (p \cdot \Phi[q] > p' \cdot \Phi[q'])$$

- For each state $q \in Q$, $r[q]$ gives the number of times a pair (q, p) with first state q has been extracted from S .
- $\pi[(q, p)]$ defines the predecessor in the path for pair (q, p) .
- $E[q]$ provides all the edges that start from state q .

WORD GRAPHS BASED N-BEST SEARCH


Algorithm 2: N-best algorithm

```

for  $q : 0 \dots |Q| - 1$  do  $r[q] = 0$  ;
 $\pi[(i, 1)] = \text{NIL}$  ;
 $S = \{(i, 1)\}$  ;
while  $S \neq \emptyset$  do
     $(q, p) = \text{DEQUEUE}(S)$  ;
     $r[q] = r[q] + 1$  ;
    if ( $r[q] = N$  and  $q = f \in F$ ) then EXIT;
    if ( $r[q] \leq N$ ) then
        for  $e \in E(q)$  do
             $p' = p \cdot w[e]$  ;
             $\pi[(n(e), p')] = (q, p)$  ;
             $\text{ENQUEUE}(S, (n(e), p'))$  ;
    
```

[M.Mohri and M.Riley: *An efficient algorithm for the N-best-strings problem*, 2002], for more details.

WORD GRAPHS BASED N-BEST SEARCH



WORD GRAPHS BASED N-BEST SEARCH

		.32	.64	.56	1	.8	.8	1	1	1
Φ	0	1	2	3	4	5	6	7	8	
r	1	1	1	0	0	2	0	2	2	


N=2

S

(0, 1)	1	1	· 0.5	$\pi(1, 0.5) \leftarrow (0, 1)$
	2	1	· 0.5	$\pi(2, 0.5) \leftarrow (0, 1)$
(1, 0.5)	3	0.5	· 0.2	$\pi(3, 0.1) \leftarrow (1, 0.5)$
	5	0.5	· 0.8	$\pi(5, 0.4) \leftarrow (1, 0.5)$
(5, 0.4)	6	0.4	· 0.2	$\pi(6, 0.08) \leftarrow (5, 0.4)$
	7	0.4	· 0.8	$\pi(7, 0.32) \leftarrow (5, 0.4)$
(7, 0.32)	8	0.32	· 1	$\pi(8, 0.32) \leftarrow (7, 0.32)$
(8, 0.32)				
(2, 0.5)	4	0.5	· 0.3	$\pi(4, 0.15) \leftarrow (2, 0.5)$
	5	0.5	· 0.7	$\pi(5, 0.35) \leftarrow (2, 0.5)$
(5, 0.35)	6	0.35	· 0.2	$\pi(6, 0.07) \leftarrow (5, 0.35)$
	7	0.35	· 0.8	$\pi(7, 0.28) \leftarrow (5, 0.35)$
(7, 0.28)	8	0.28	· 1	$\pi(8, 0.28) \leftarrow (7, 0.28)$
(8, 0.28)				

(0, 1)	(1, 0.5)	(2, 0.5)
	0.32	0.28
(3, 0.1)	(5, 0.4)	(6, 0.08)
0.1	0.32	0.08
(7, 0.32)	(8, 0.32)	(4, 0.15)
0.32	0.32	0.12
(5, 0.35)	(6, 0.07)	(7, 0.28)
0.28	0.07	0.28
(8, 0.28)		
0.28		

WORD GRAPHS BASED N-BEST SEARCH



WORD GRAPHS BASED N-BEST SEARCH

