

AI Planning.

Plan-space planning



Eva Onaindia

Universitat Politècnica de València

Acknowledgements

Most of the slides used in this course are taken or are modifications from Dana Nau's lecture slides for the textbook *Automated Planning*, licensed under the Creative Commons Attribution-NonCommercial-ShareAlike License:

<http://creativecommons.org/licenses/by-nc-sa/2.0/>

<http://creativecommons.org/licenses/by/3.0/es/>


I would like to gratefully acknowledge Dana Nau's contributions and thank him for generously permitting me to use aspects of his presentation material.

Plan-space planning. Outline

- Motivation
- POP tree (plan space search)
- Structure of a partial plan
 - Example *transportation*
- The PSP algorithm
- Example
- POP and the Sussman anomaly
- Summary
- Heuristics in plan-space planning
 - Flaw-selection heuristics
 - Resolver-selection heuristics

Plan-space planning. Motivation.

- Problem with state-space search
 - In some cases we may try many different orderings of the same actions before realizing there is no solution



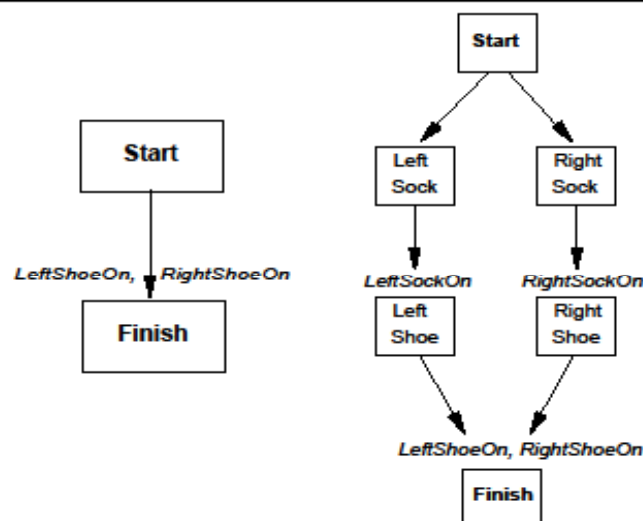
- *Least-commitment strategy*: don't commit to orderings, instantiations, etc., until necessary

Plan-space planning. Motivation.

- State-space planning:
 - Actions are planned in the same order as they will be later executed
 - Sequential plans (total order)
 - Multiple relationships among actions. Committing to an specific ordering may yield a non-efficient planning process
- Plan-space planning:
 - Not commitment to a particular action ordering until necessary
 - Plans with parallel actions (POP: partial-order planning)

Plan-space planning. Motivation.

Partially ordered plans



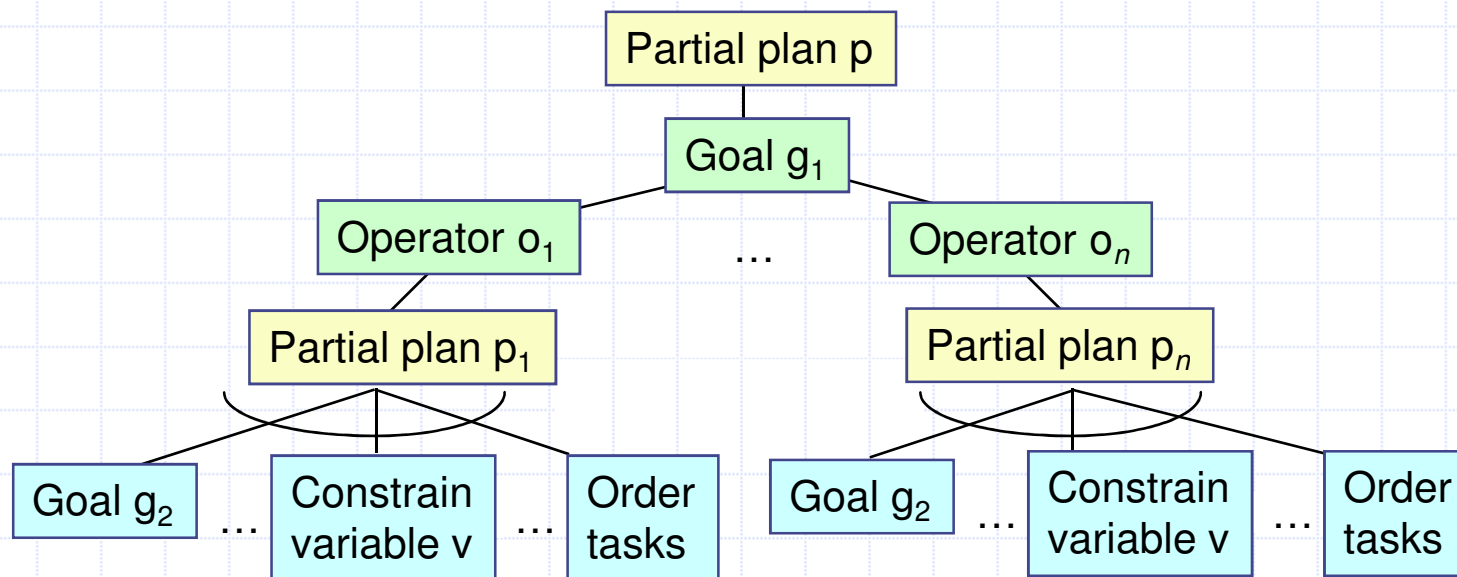
A plan is complete iff every precondition is achieved

A precondition is achieved iff it is the effect of an earlier step
and no possibly intervening step undoes it

POP tree (plan space search)

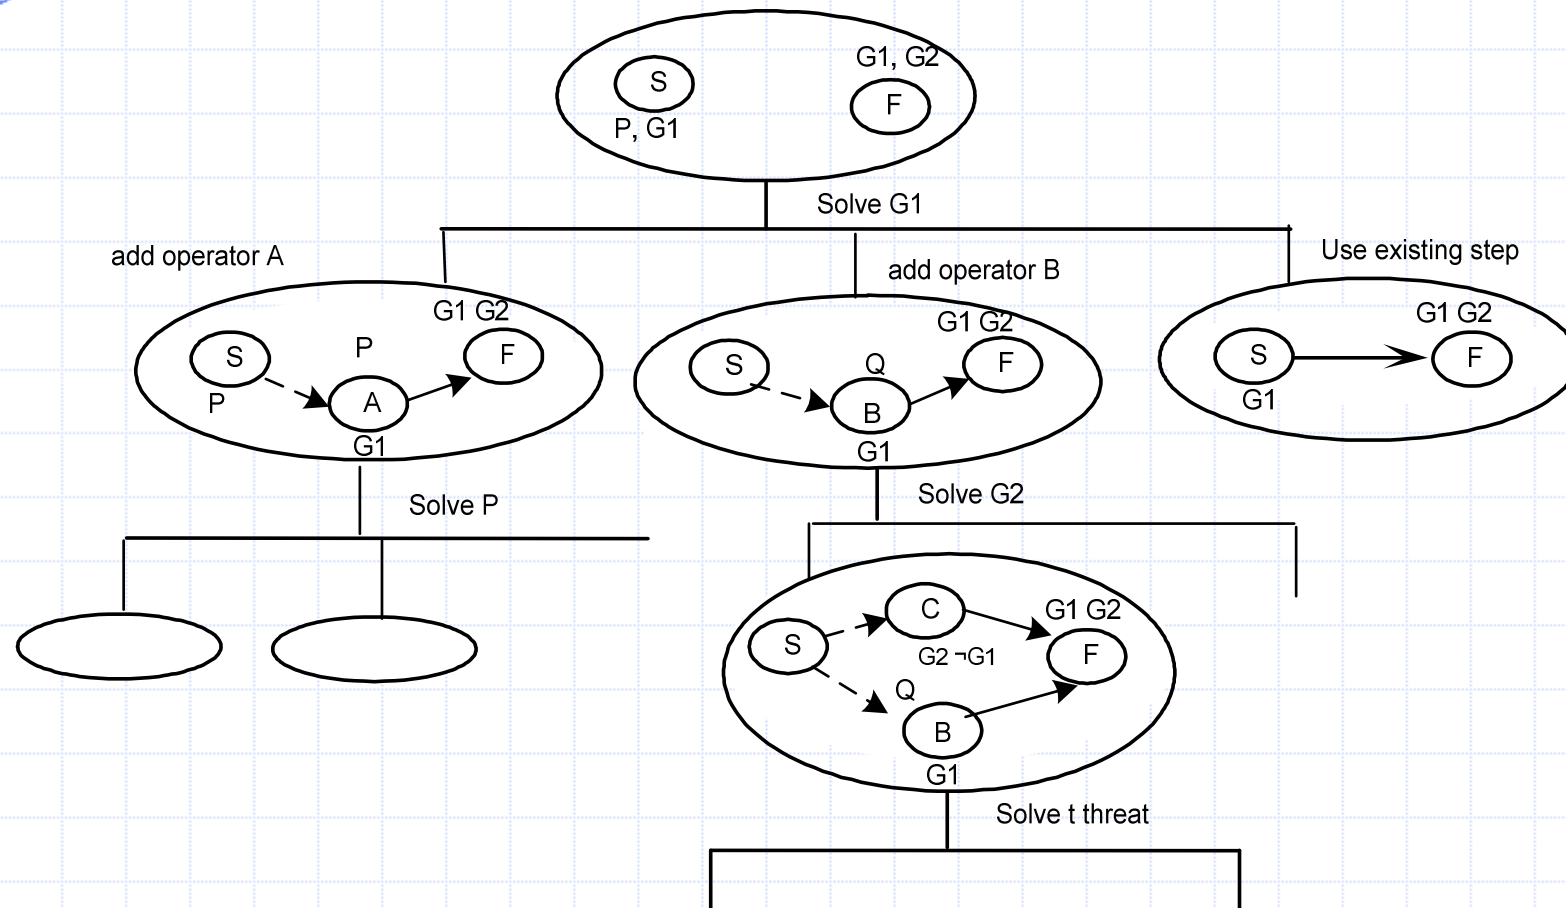
- Backward search from the goal
- Each node of the search space is a *partial plan* (*graph of actions*)
 - A set of partially-instantiated actions
 - A set of constraints
- Make more and more refinements, until we have a solution

POP tree (plan space search)



Nodes in the POP tree are partial plans (yellow boxes)

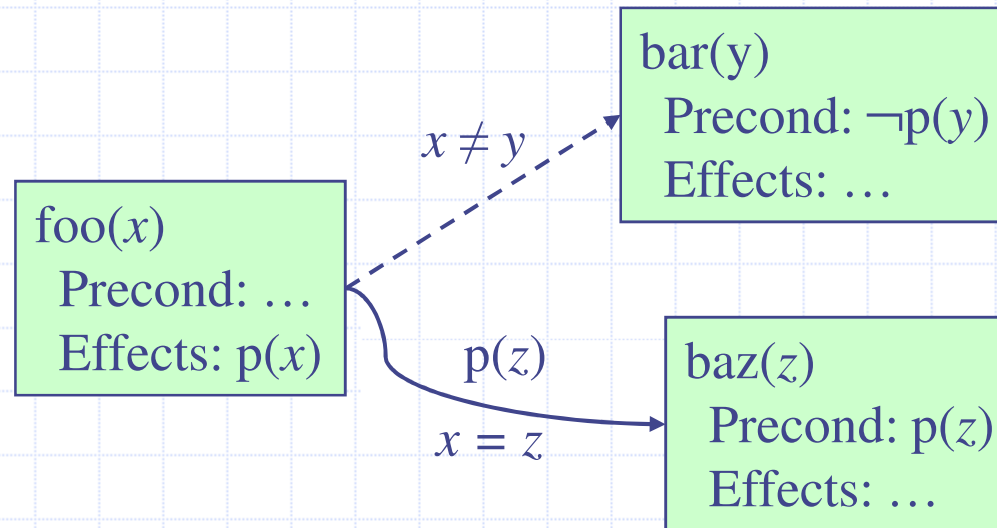
POP tree (plan space search)



Structure of a partial plan

- Plan step: totally or partially-instantiated operator (action)
 - Start step: initial state
 - Finish step: problem goals
- Types of constraints between steps:
 - *causal link*:
 - use action a to establish the precondition p needed by action b $\langle a, p, b \rangle$
 - *precedence constraint*: a must precede b ($a < b$)
 - *binding constraints*:
 - inequality constraints, e.g., $v_1 \neq v_2$ or $v \neq c$
 - equality constraints (e.g., $v_1 = v_2$ or $v = c$) and/or substitutions

Structure of a partial plan (example 1)



Causal link: $\langle \text{foo}(x), p(x), \text{bar}(z) \rangle$

Binding/equality constraint: $x = z$

Precedence constraint: $\text{foo}(x) < \text{bar}(y)$

Inequality constraint: $x \neq y$

Structure of a partial plan (plan refinement)


- Plan refinement: selecting a *flaw* in the plan and branching all over ways of solving the flaw
- Flaw:
 - *open goal*: selecting a pending subgoal to be solved in the plan
 - *binding constraint*: selecting a non-instantiated variable
 - *threat*: selecting a step that *threatens* a causal link
 - causal link: $\langle p1, l, p2 \rangle$
 - $\exists p3 / l \in \text{effects}^-(p3)$
 - $p3$ is unordered with respect to $p1$ and $p2$
 - - promotion: $p3 < p1$
 - - demotion: $p2 < p3$
 - - separation: constraint variables to avoid the threat
- How to tell we have a solution: no more *flaws* in the plan

Flaws: open goals


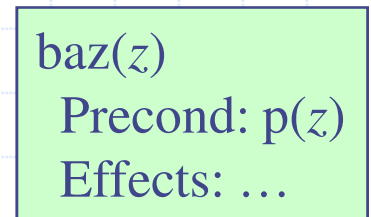
- Open goal:
 - An action/step a has a precondition p that we haven't decided how to establish

- Resolving the flaw:

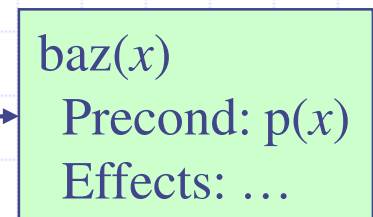
- Find an action/step b
 - (either **already in the plan**, or **insert it**)
- that can be used to establish p
 - can precede a and produce p
- Instantiate variables and/or constrain variable bindings
- Create a causal link



p(z)




p(x)



Flaws: threats

- Threat: a deleted-condition interaction
 - Action/step a establishes a precondition (e.g., $pq(x)$) of action b
 - Another action/step c is capable of deleting p
- Resolving the flaw:
 - impose a constraint to prevent c from deleting p
- Three possibilities:
 - Make b precede c
 - Make c precede a
 - Constrain variable(s) to prevent c from deleting p




Example *transportation* (partial plans)

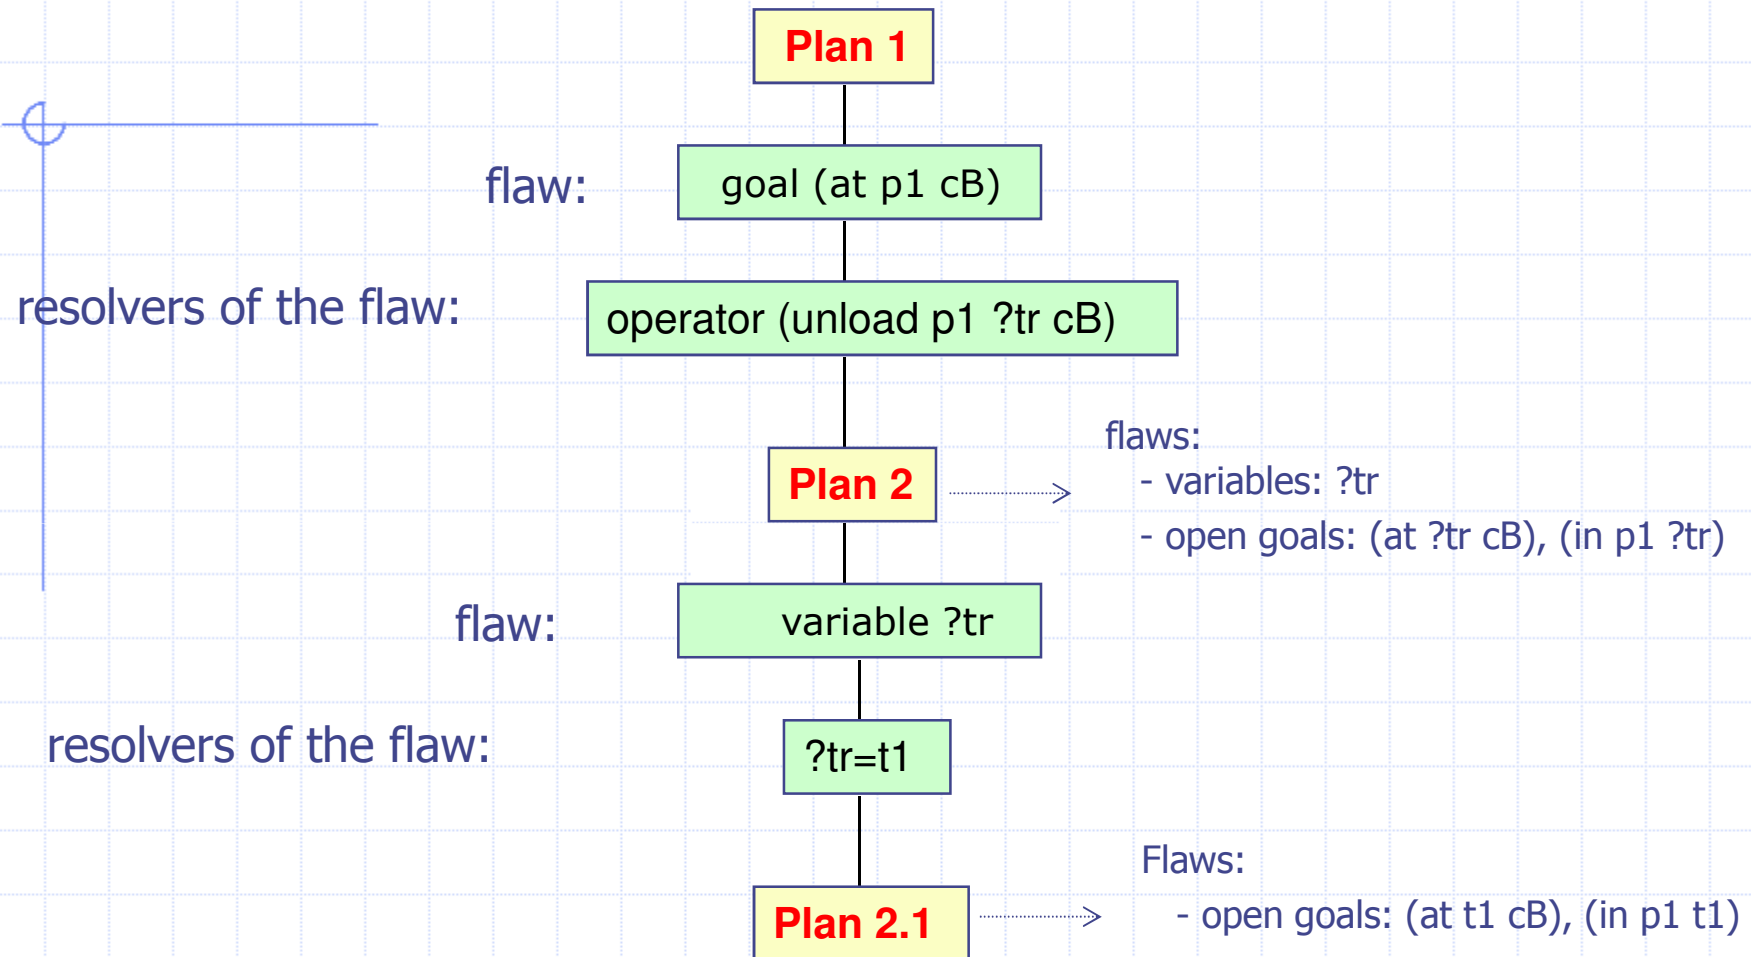
Plan 1: initial plan



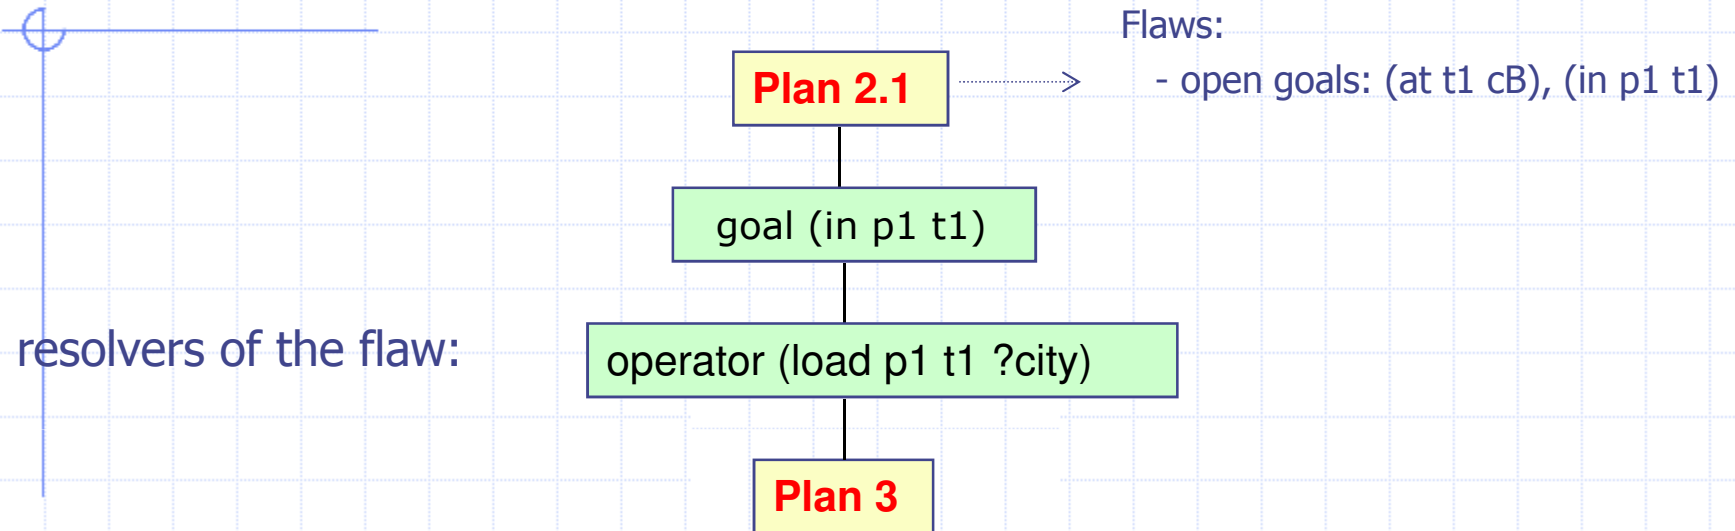
Plan 2: refinement plan over plan 1



Example *transportation* (POP tree)

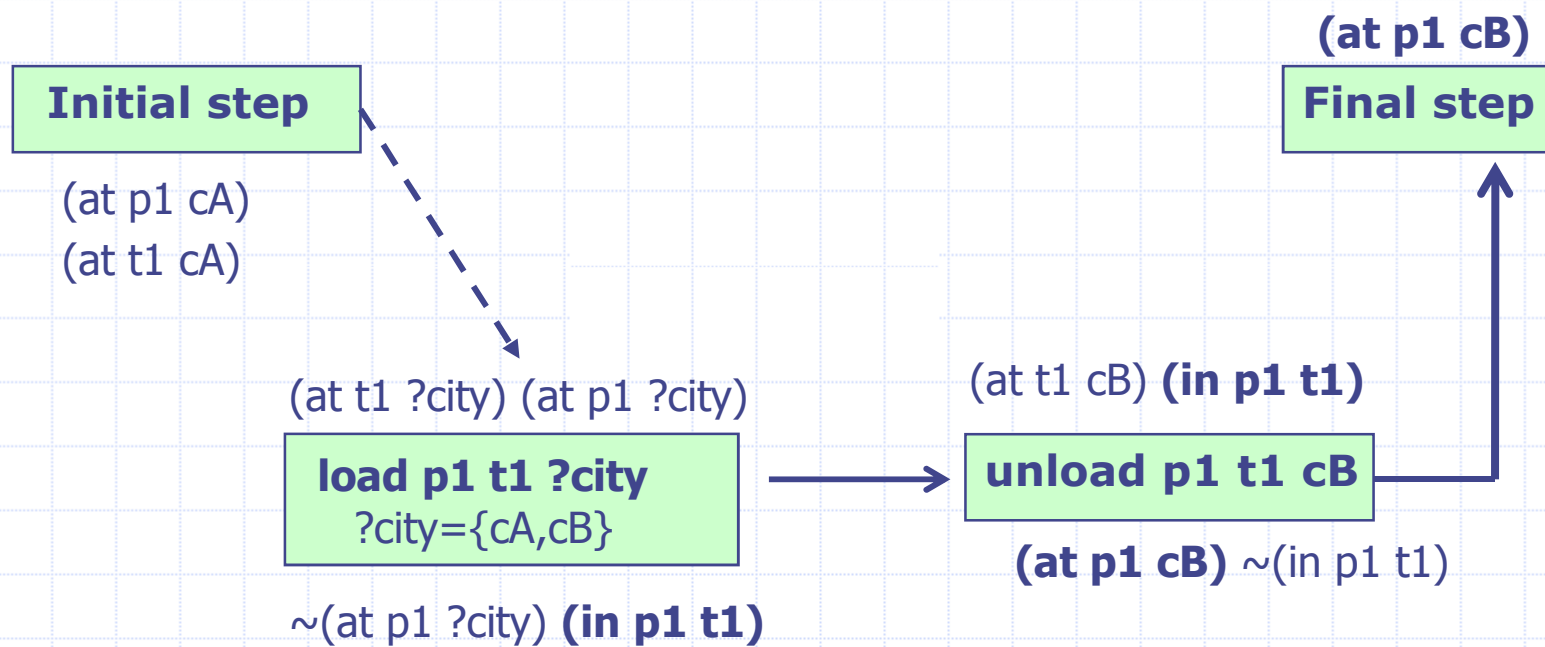


Example *transportation* (POP tree cont.)

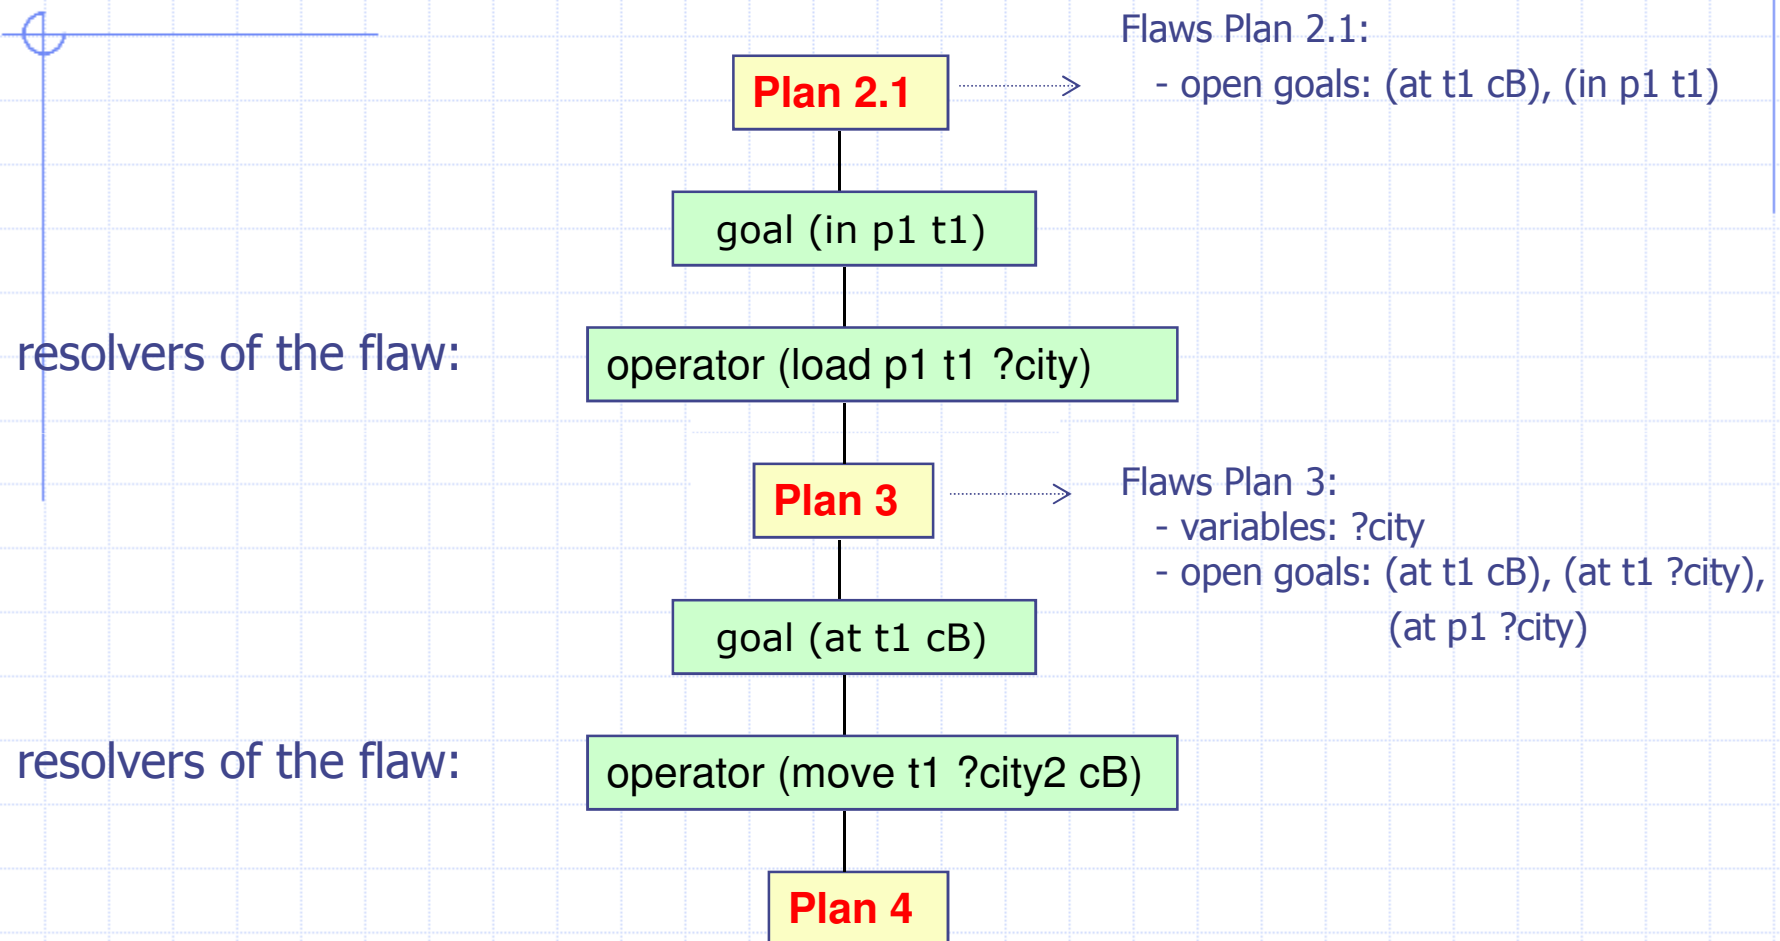


Example *transportation* (partial plans cont.)

Plan 3

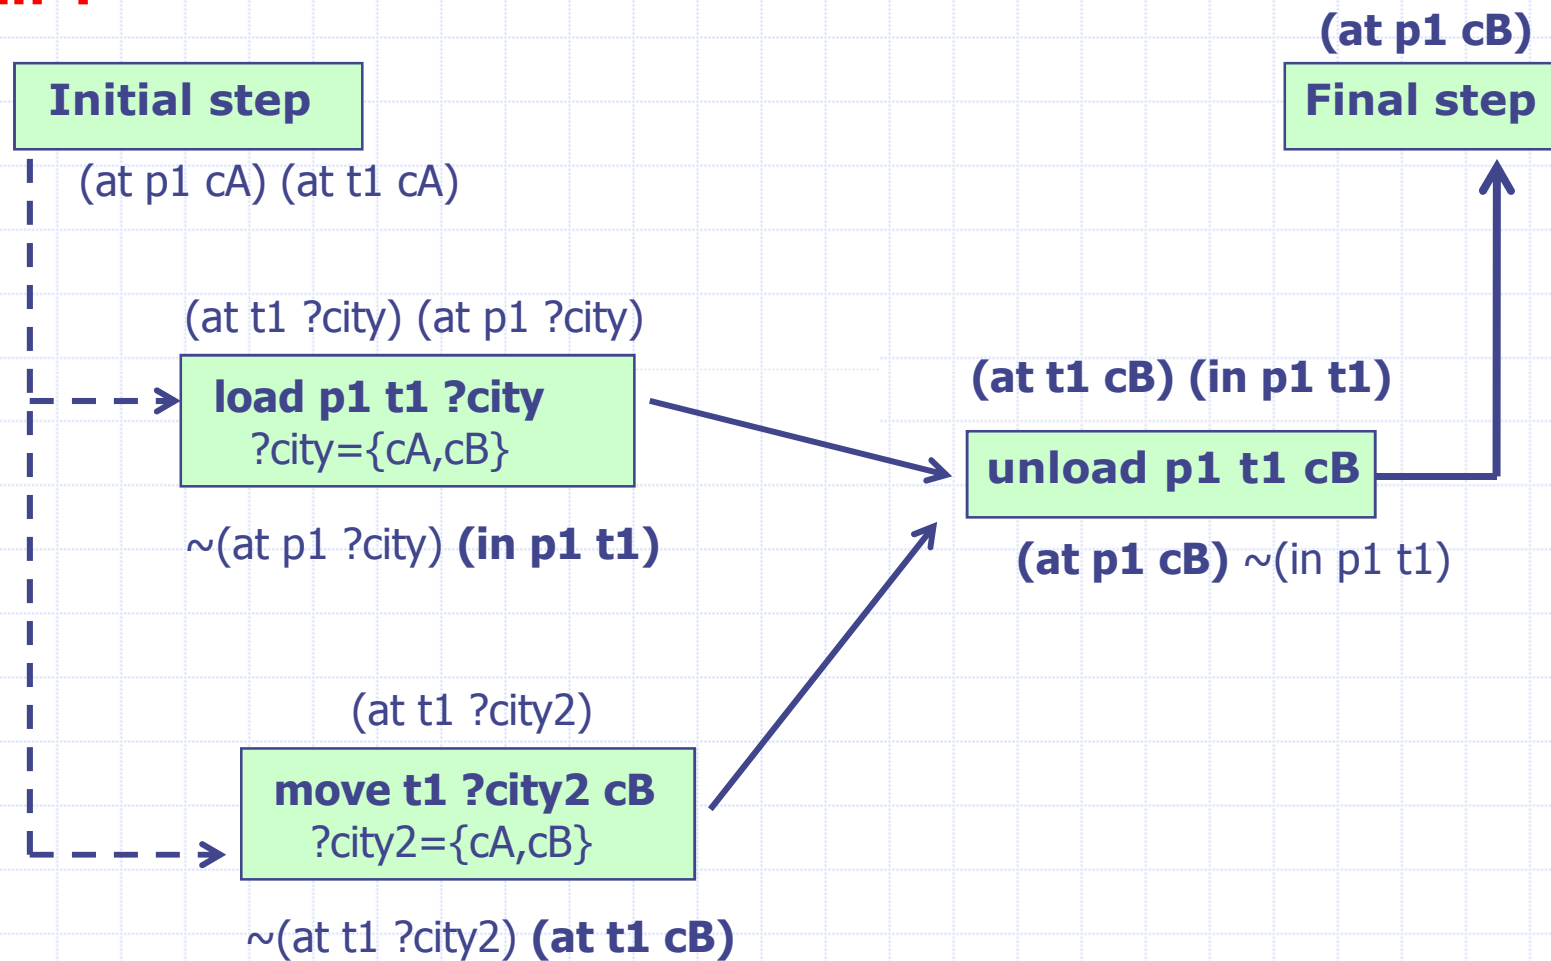


Example *transportation* (POP tree cont.)

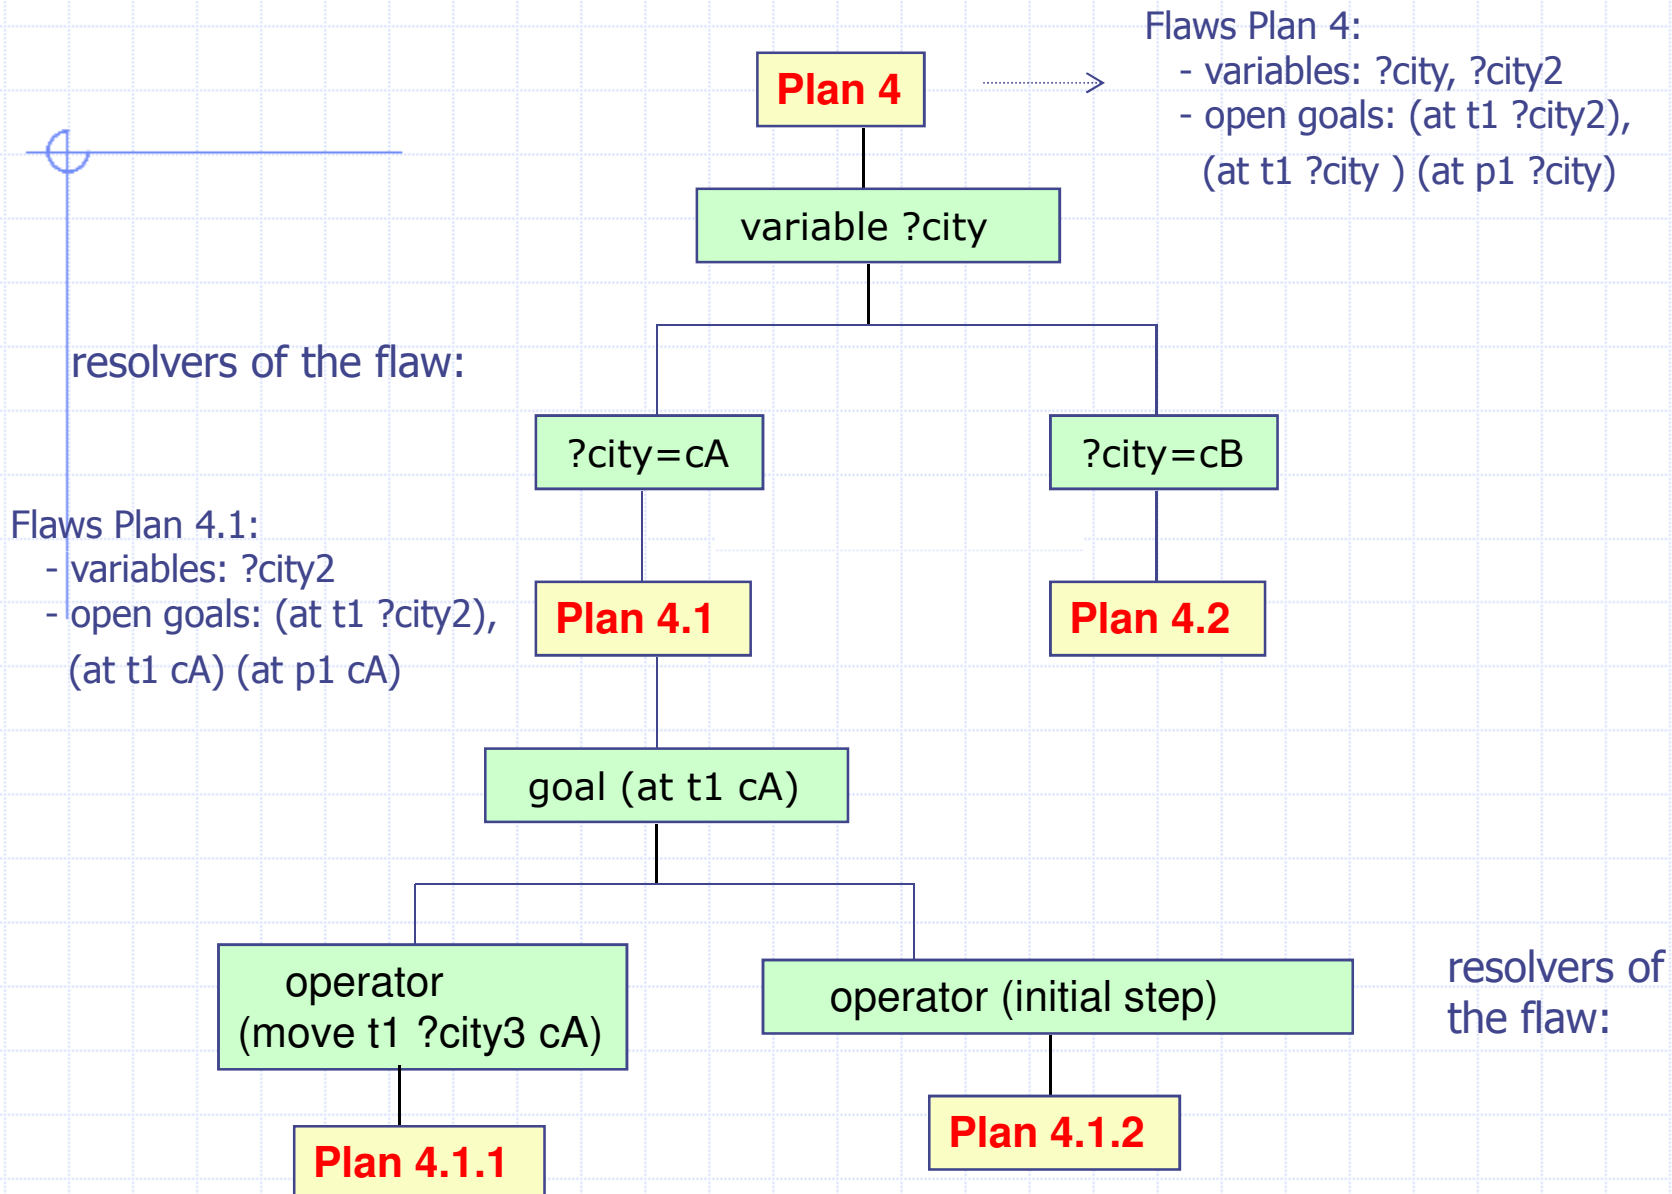


Example *transportation* (partial plans cont.)

Plan 4

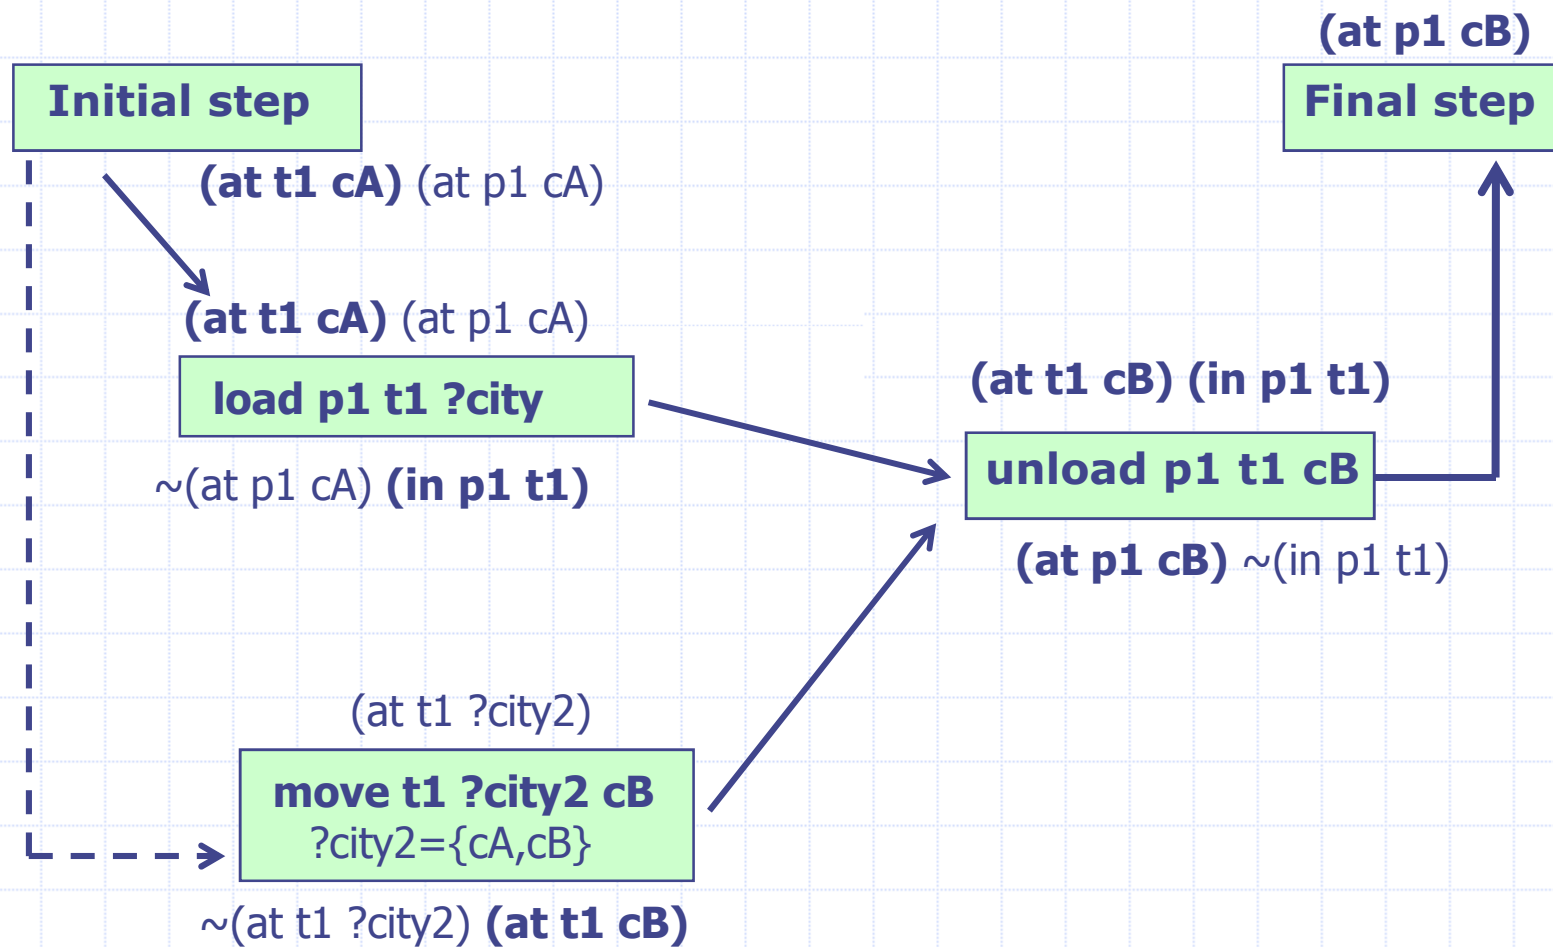


Example *transportation* (POP tree cont.)

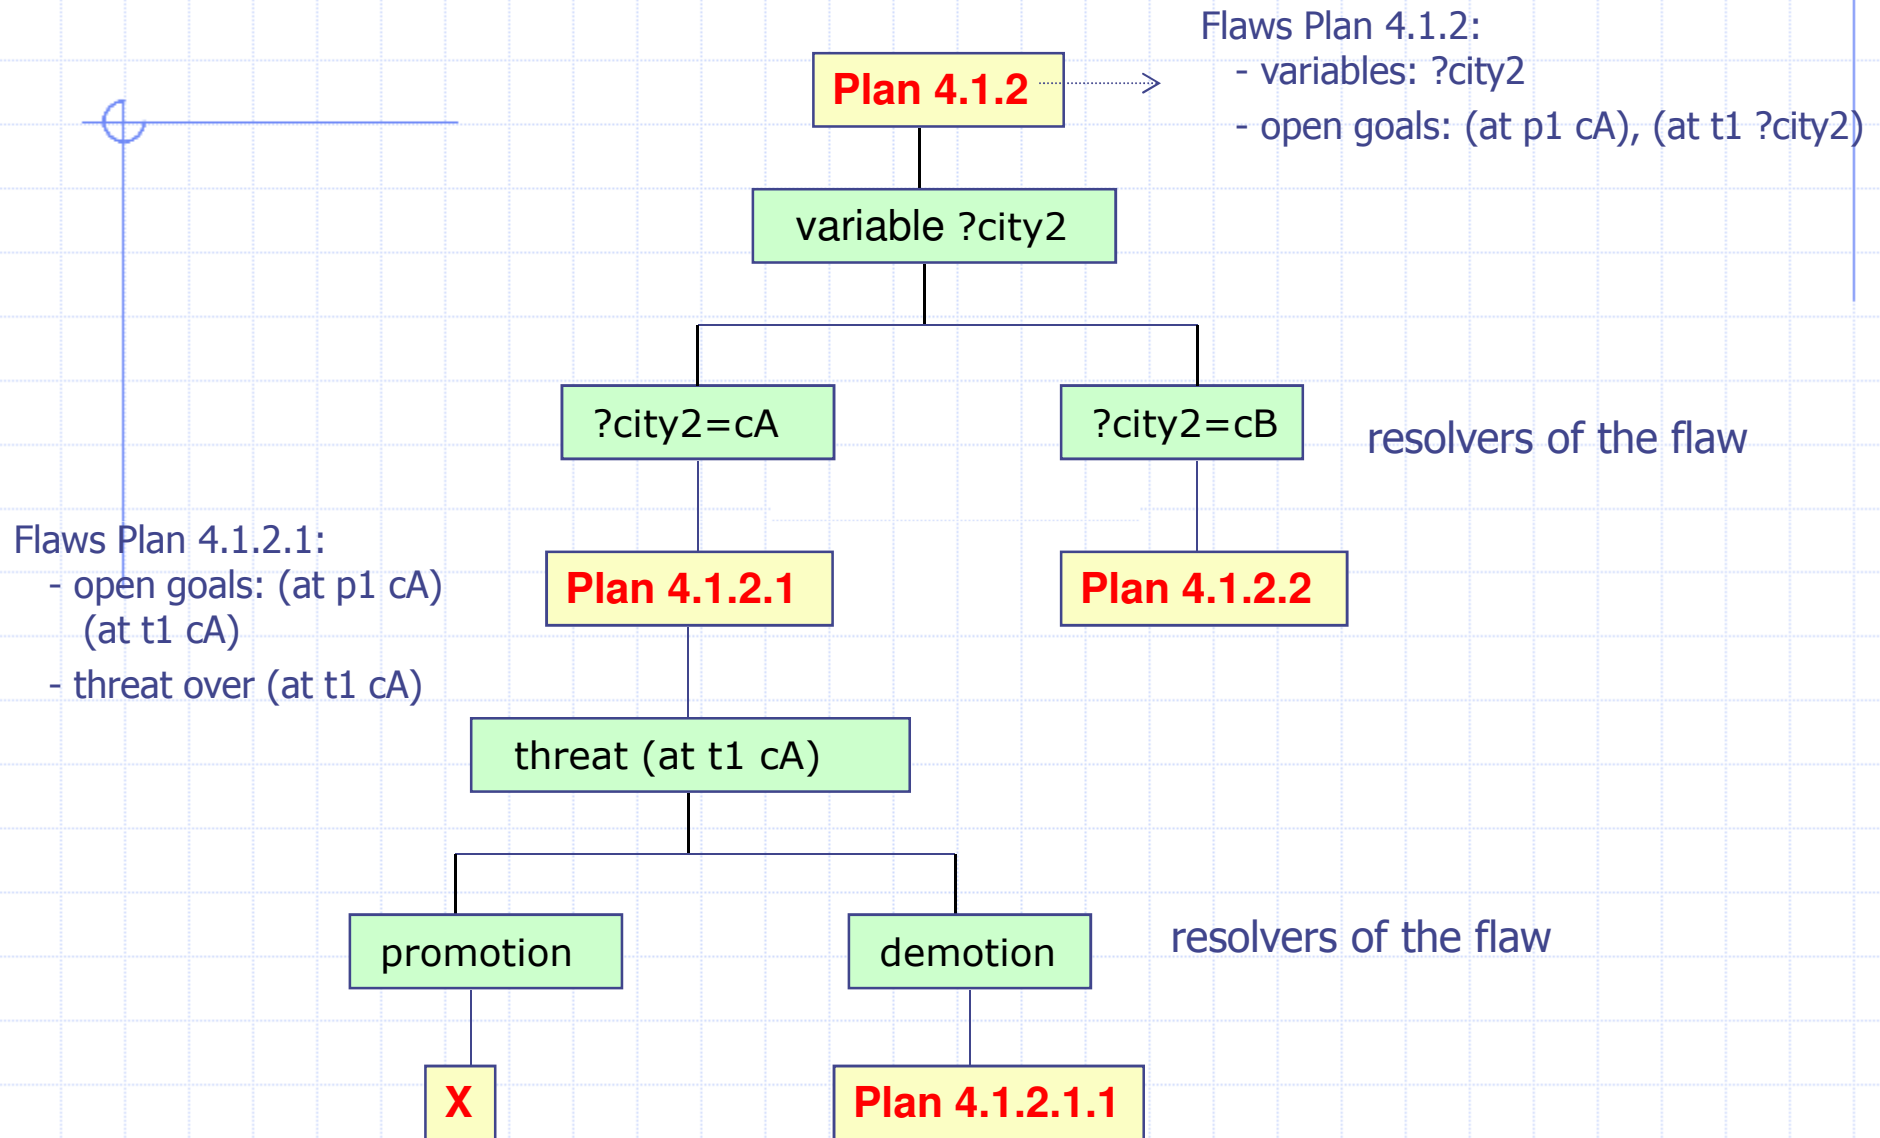


Example *transportation* (partial plans cont.)

Plan 4.1.2

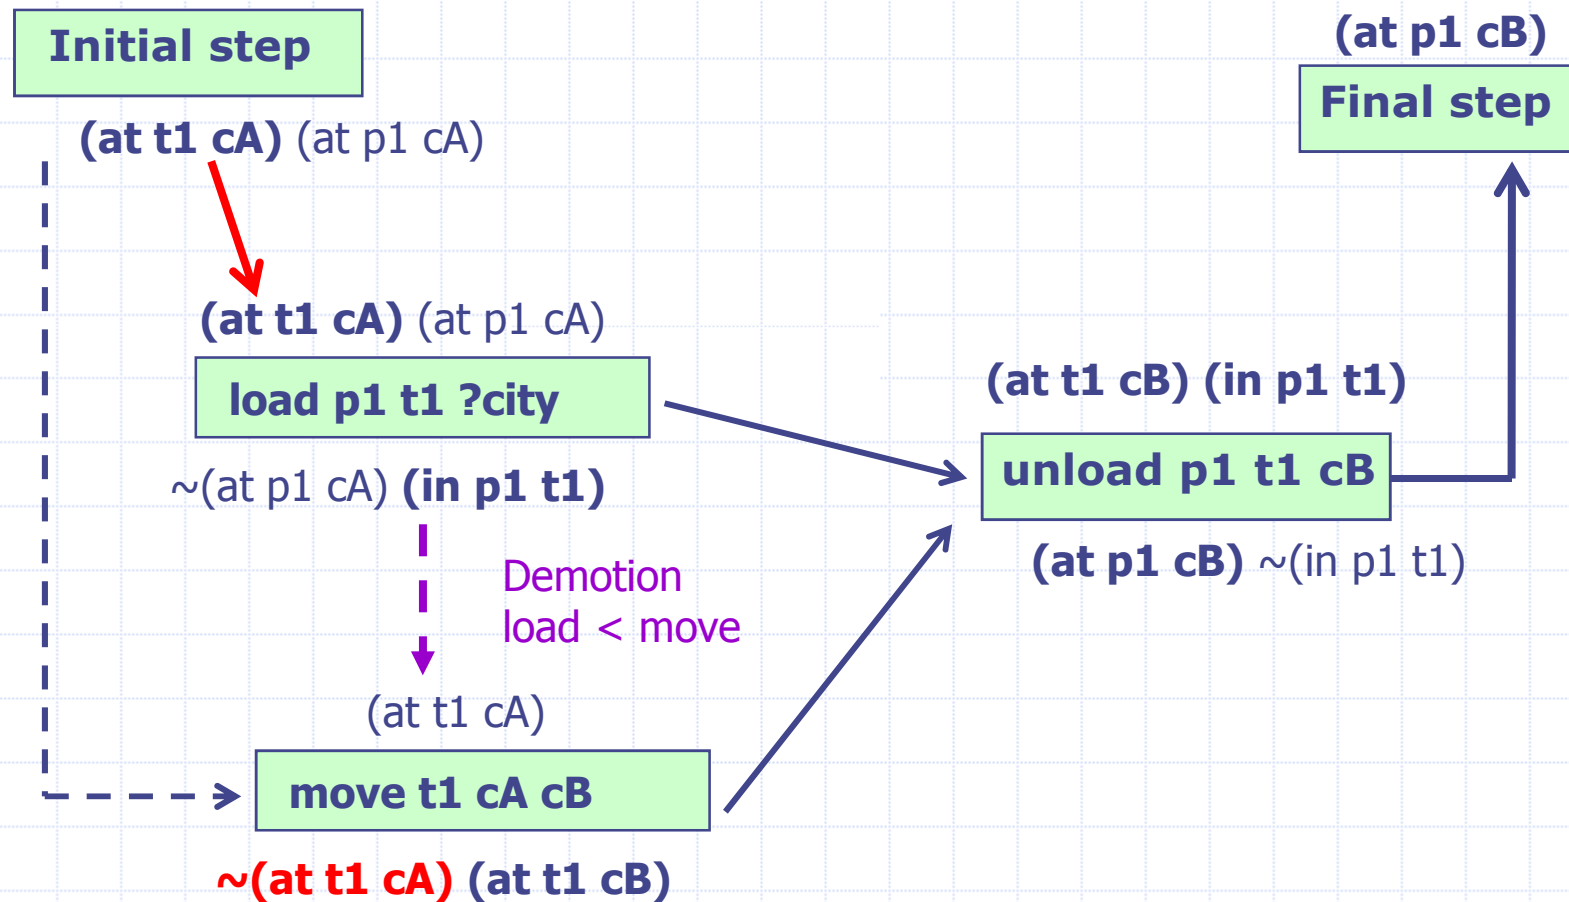


Example *transportation* (POP tree cont.)

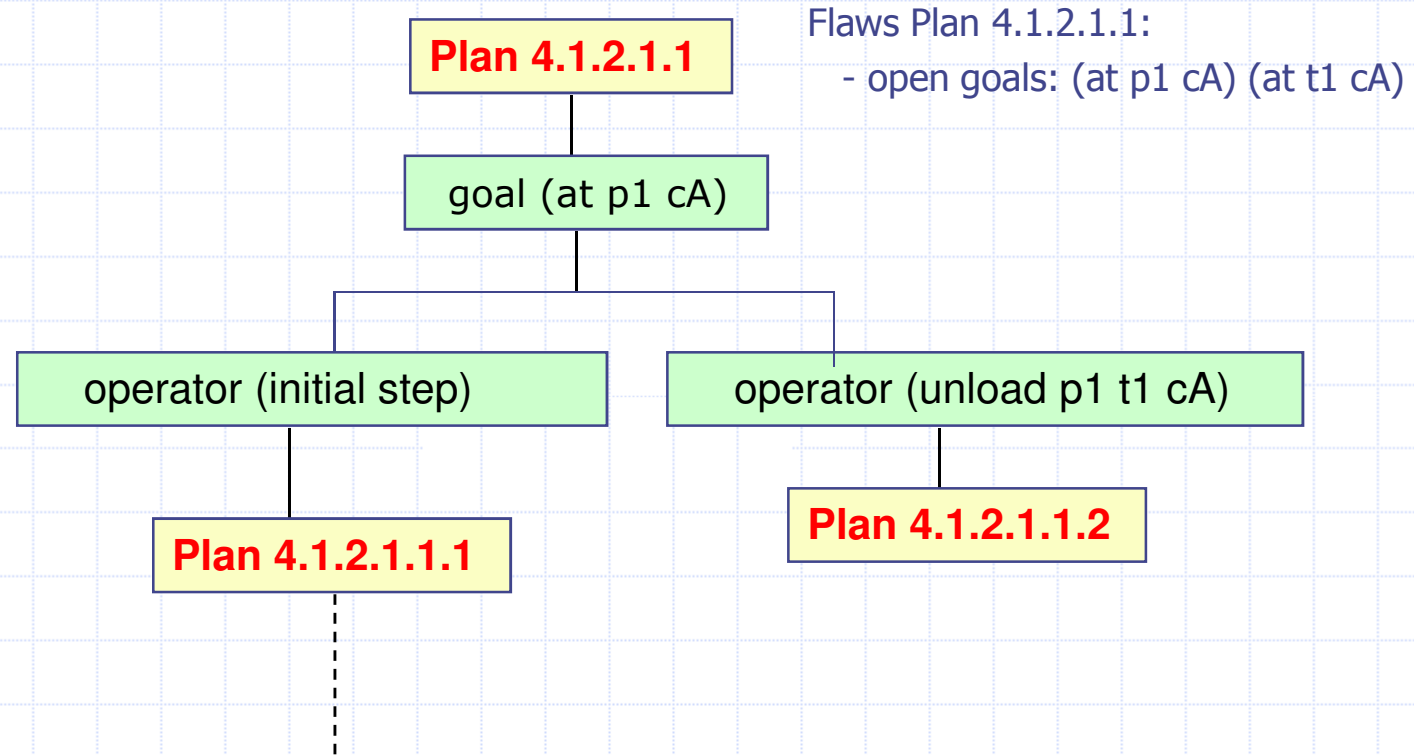


Example *transportation* (partial plans cont.)

Plan 4.1.2.1.1

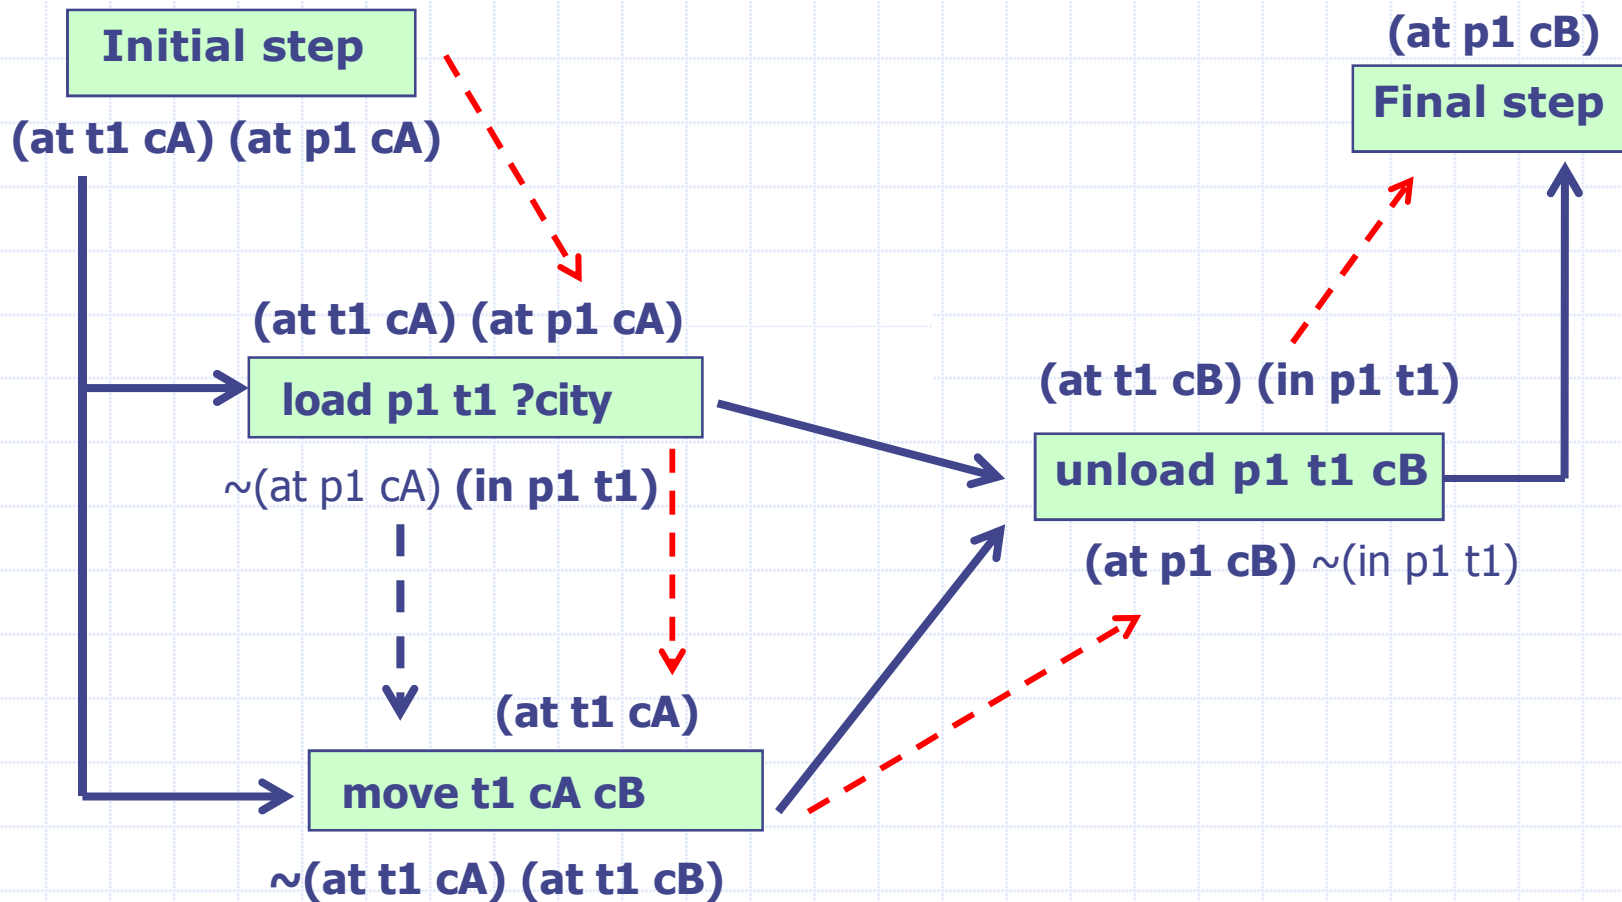


Example *transportation* (POP tree cont.)



Example *transportation* (partial plans cont.)

Plan 4.1.2.1.1.1



The PSP Procedure



PSP(π)

$flaws \leftarrow \text{OpenGoals}(\pi) \sqcup \text{Threats}(\pi)$

if $flaws = \emptyset$ then return(π)

select any flaw $\phi \in flaws$

$resolvers \leftarrow \text{Resolve}(\phi, \pi)$

if $resolvers = \emptyset$ then return(failure)

nondeterministically choose a resolver $\rho \in resolvers$

$\pi' \leftarrow \text{Refine}(\rho, \pi)$

return(PSP(π'))

end

- PSP is both sound and complete
- It returns a partially ordered solution plan
 - Any total ordering of this plan will achieve the goals
 - Or could execute actions in parallel if the environment permits it

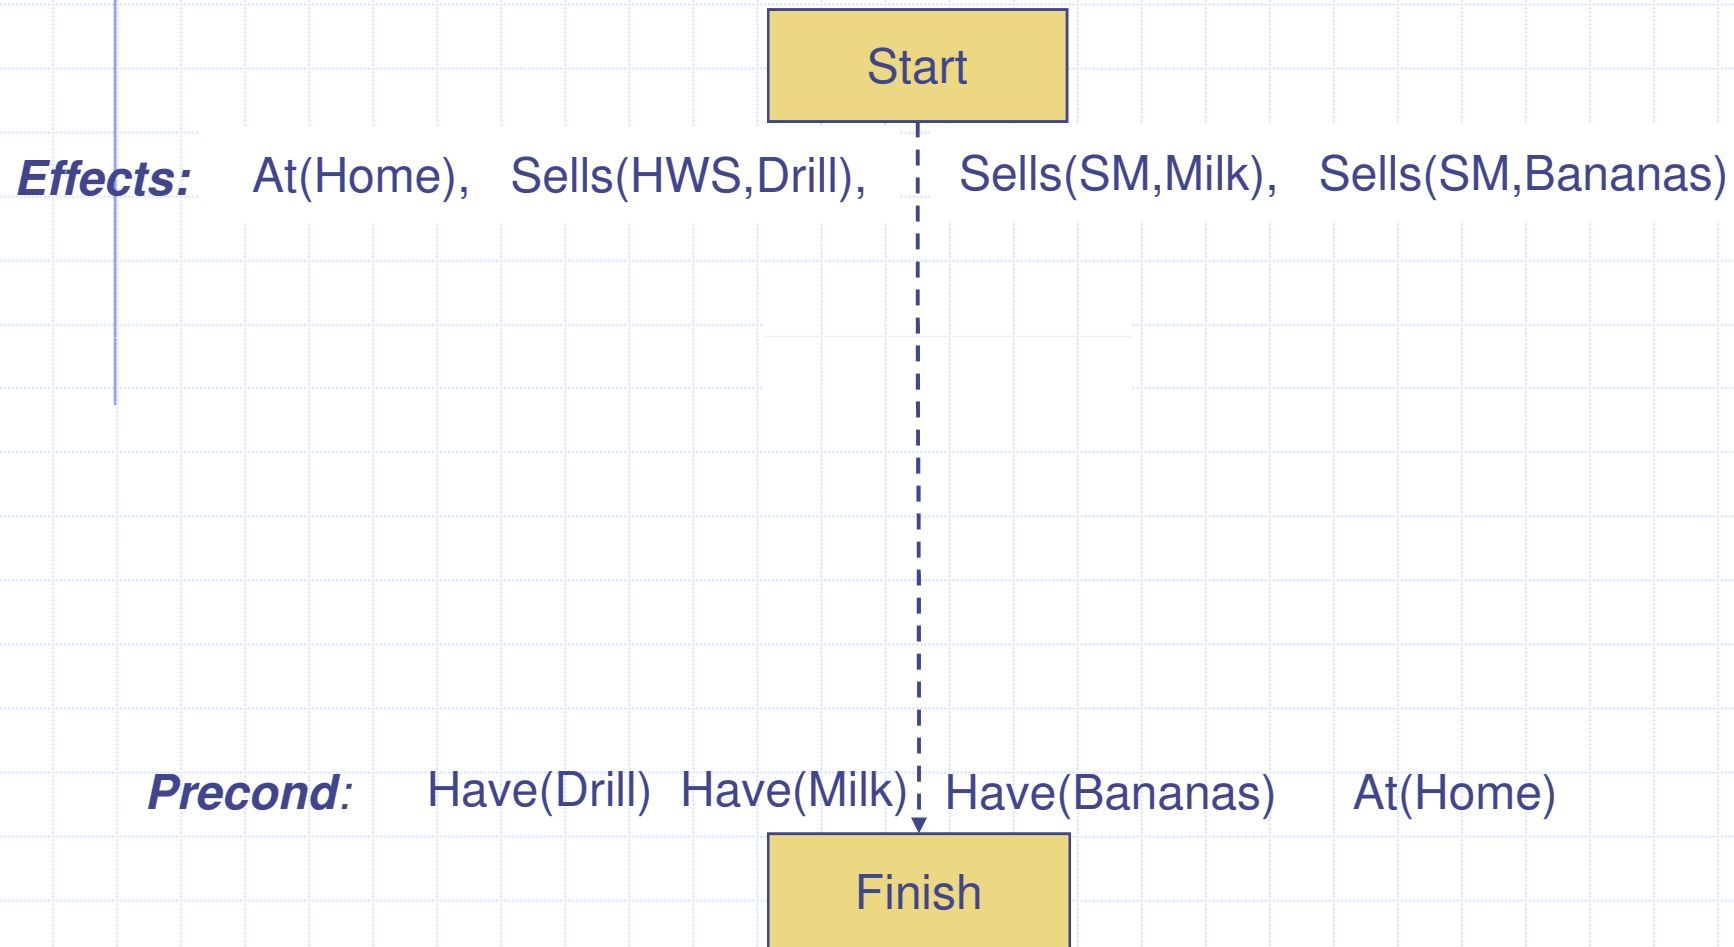
Example

- Similar (but not identical) to an example in Russell and Norvig's *Artificial Intelligence: A Modern Approach* (1st edition)
- Operators:
 - **Start**
Precond: none
Effects: At(Home), sells(HWS,Drill), Sells(SM,Milk), Sells(SM,Banana)
 - **Finish**
Precond: Have(Drill), Have(Milk), Have(Banana), At(Home)
Effects: none
 - **Go(*l,m*)**
Precond: At(*l*)
Effects: At(*m*), \neg At(*l*)
 - **Buy(*p,s*)**
Precond: At(*s*), Sells(*s,p*)
Effects: Have(*p*)

Start and **Finish** are dummy actions that we'll use instead of the initial state and goal

Example (continued)

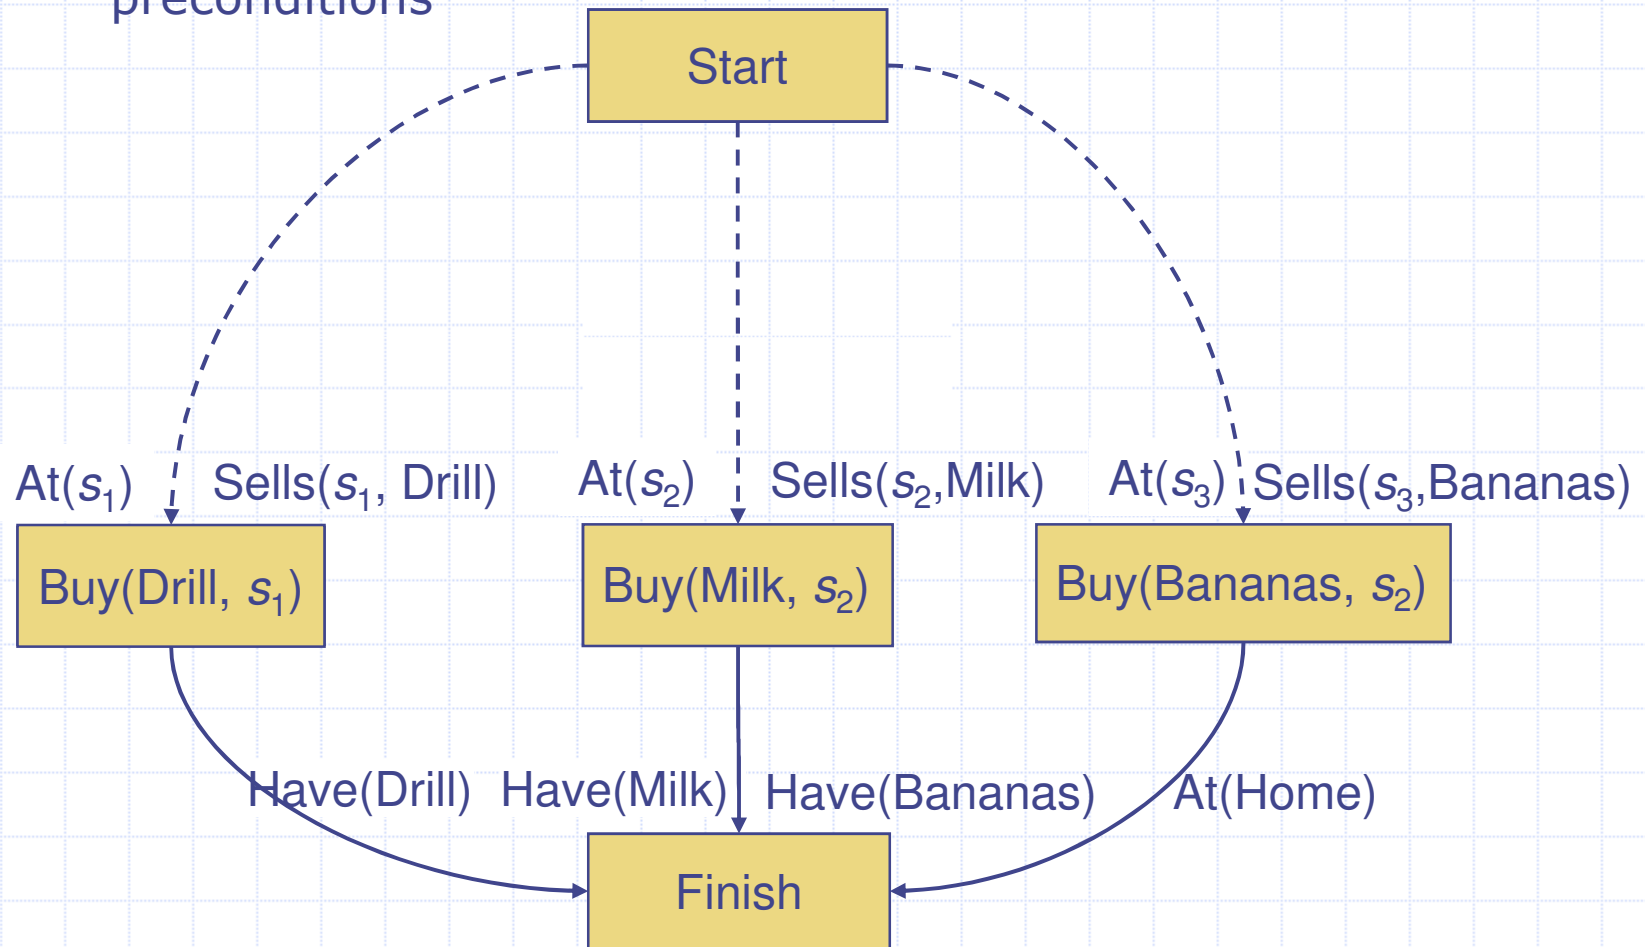
- Need to give PSP a plan π as its argument
 - Initial plan: **Start**, **Finish**, and an ordering constraint



Example (continued)

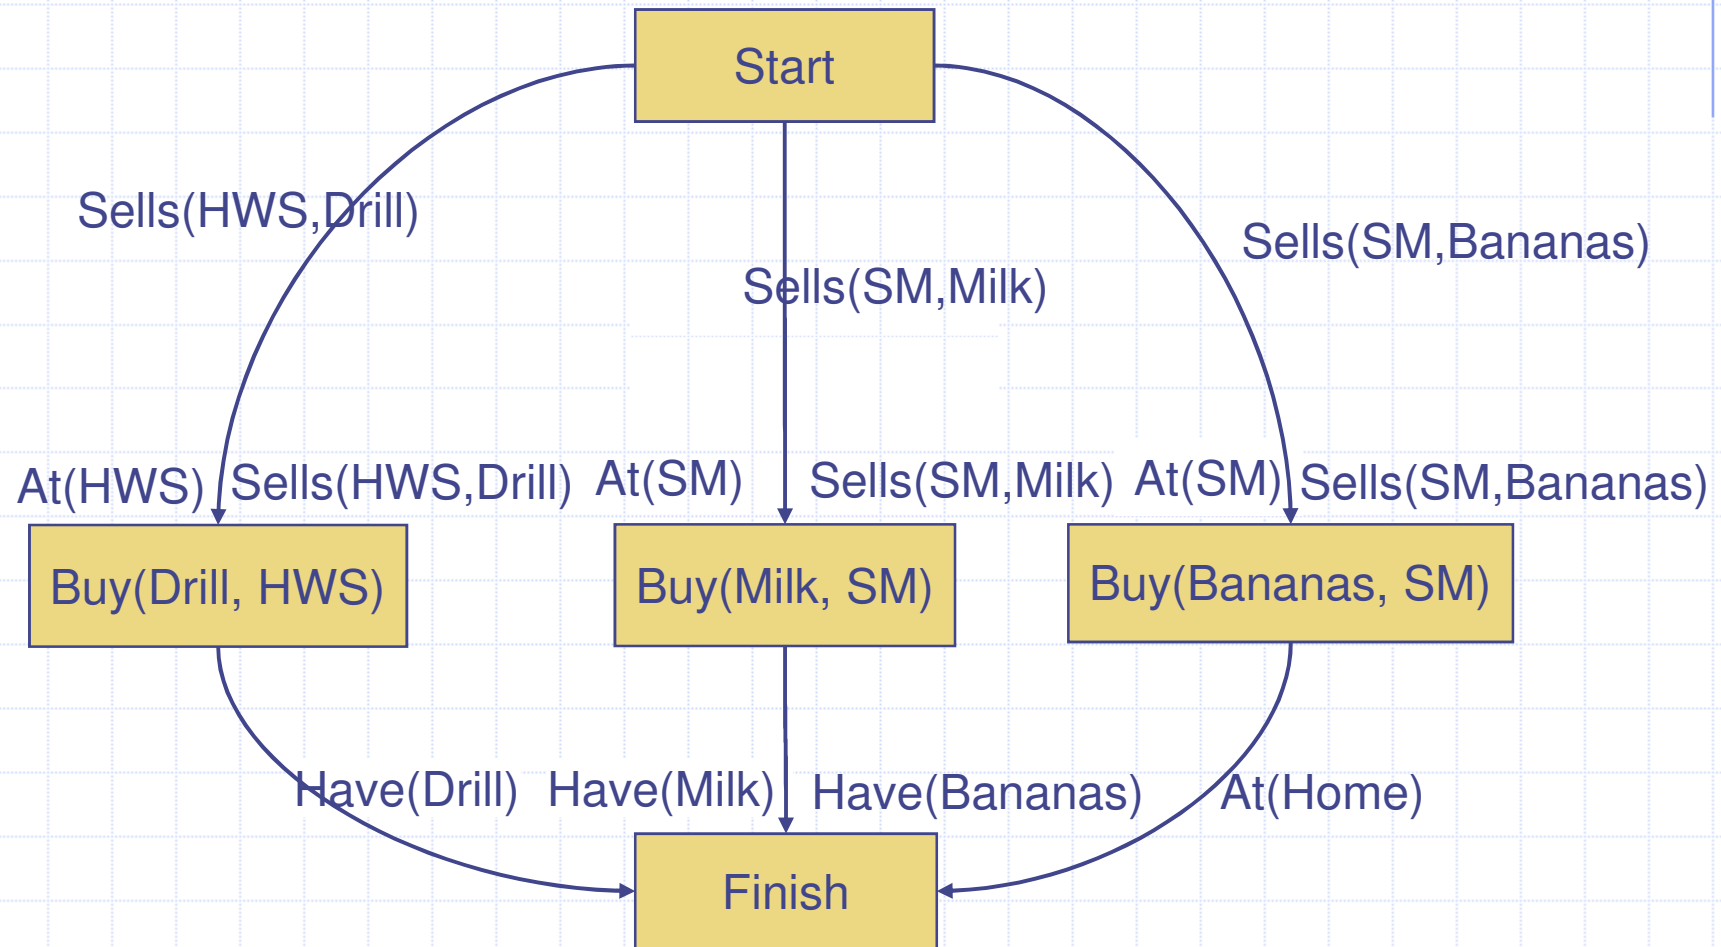
- The first three refinement steps

— These are the only possible ways to establish the Have preconditions



Example (continued)

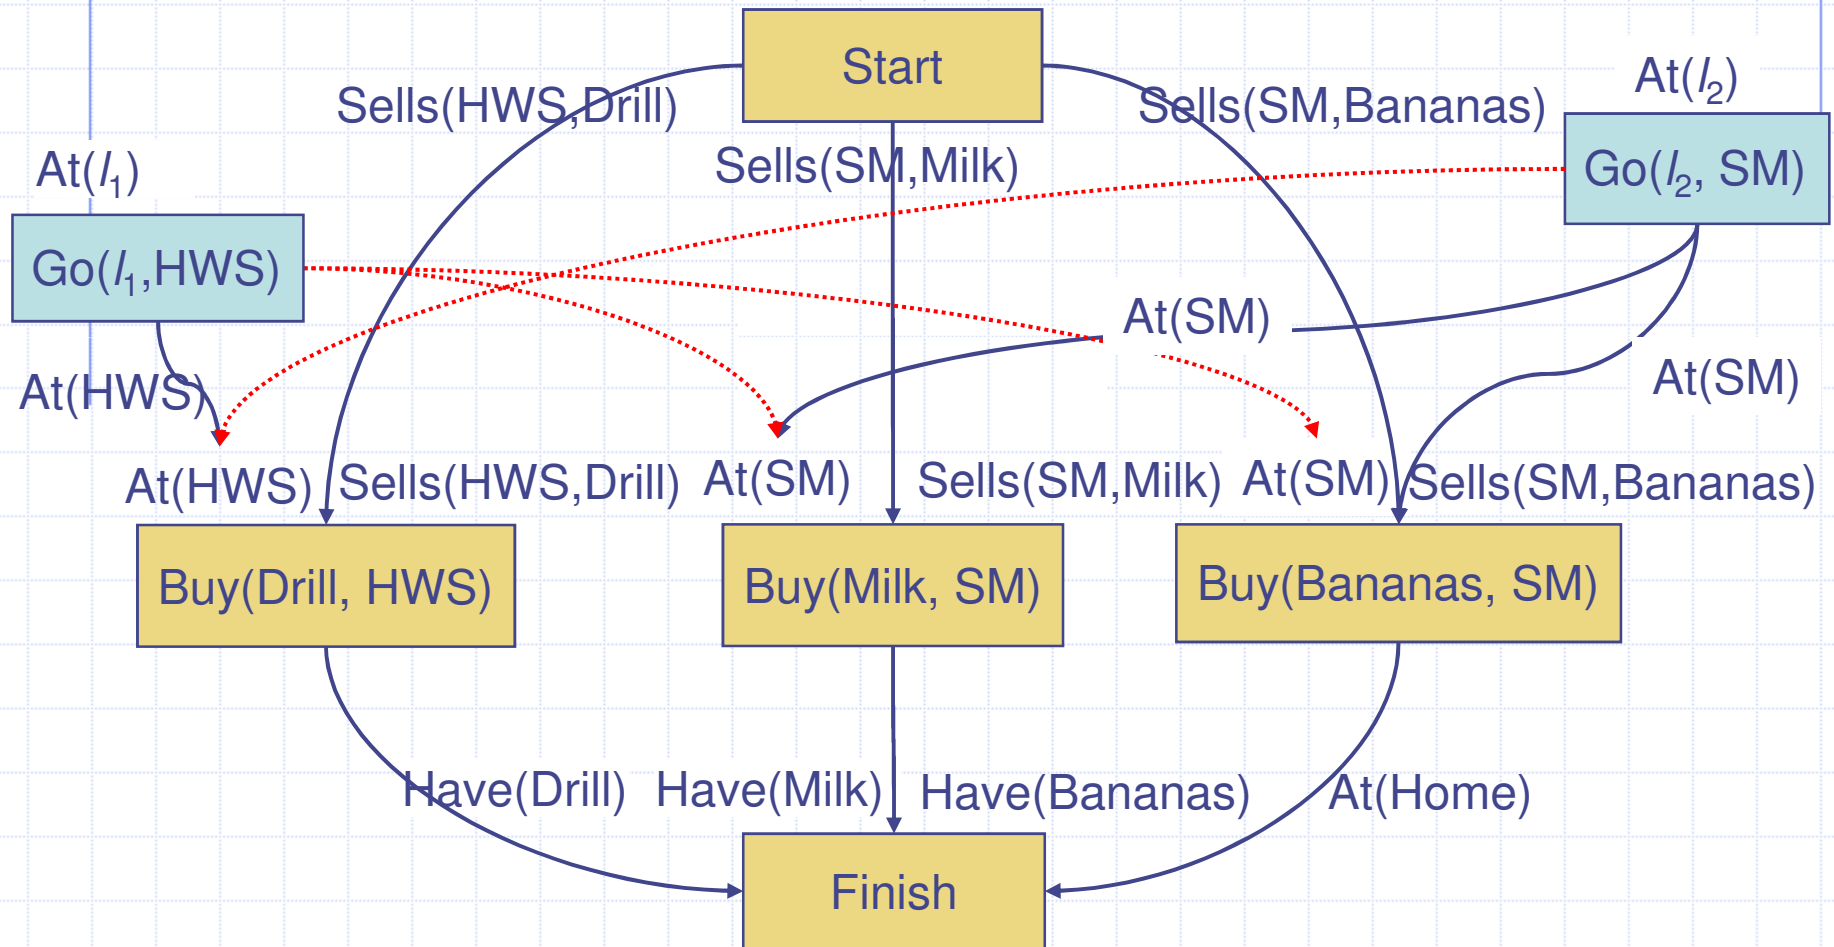
- Three more refinement steps
 - The only possible ways to establish the Sells preconditions



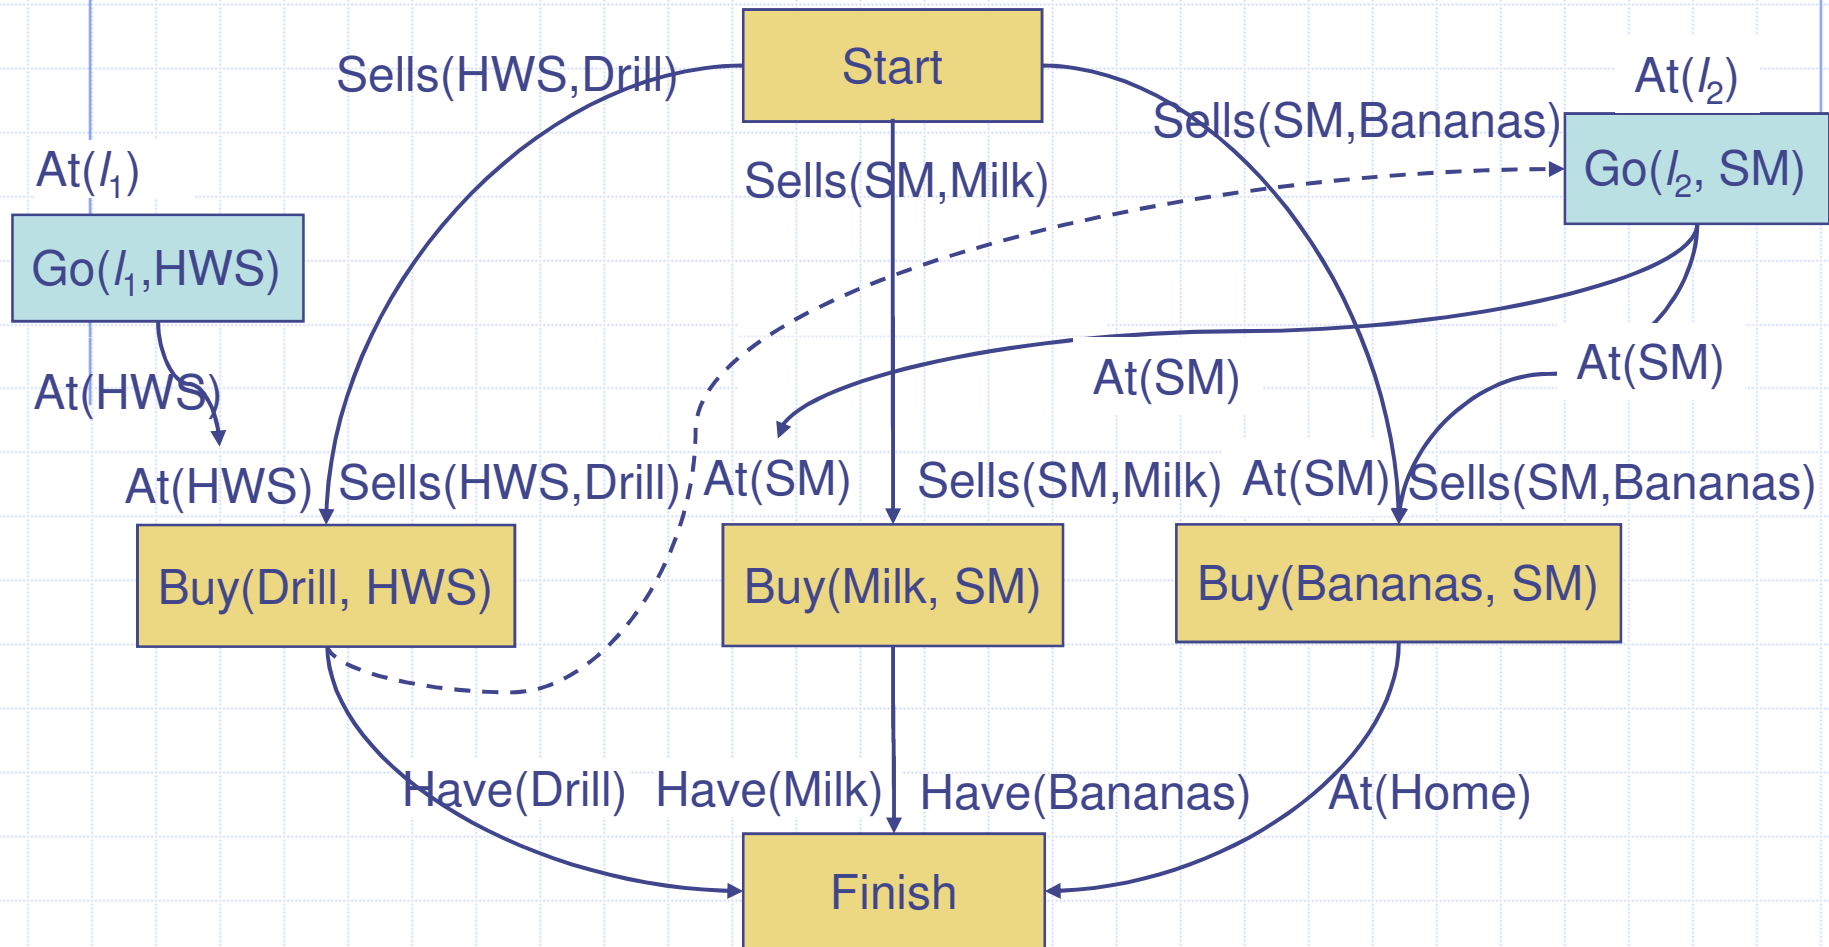
Example (continued)

- Two more refinements: the only ways to establish $At(HWS)$ and $At(SM)$

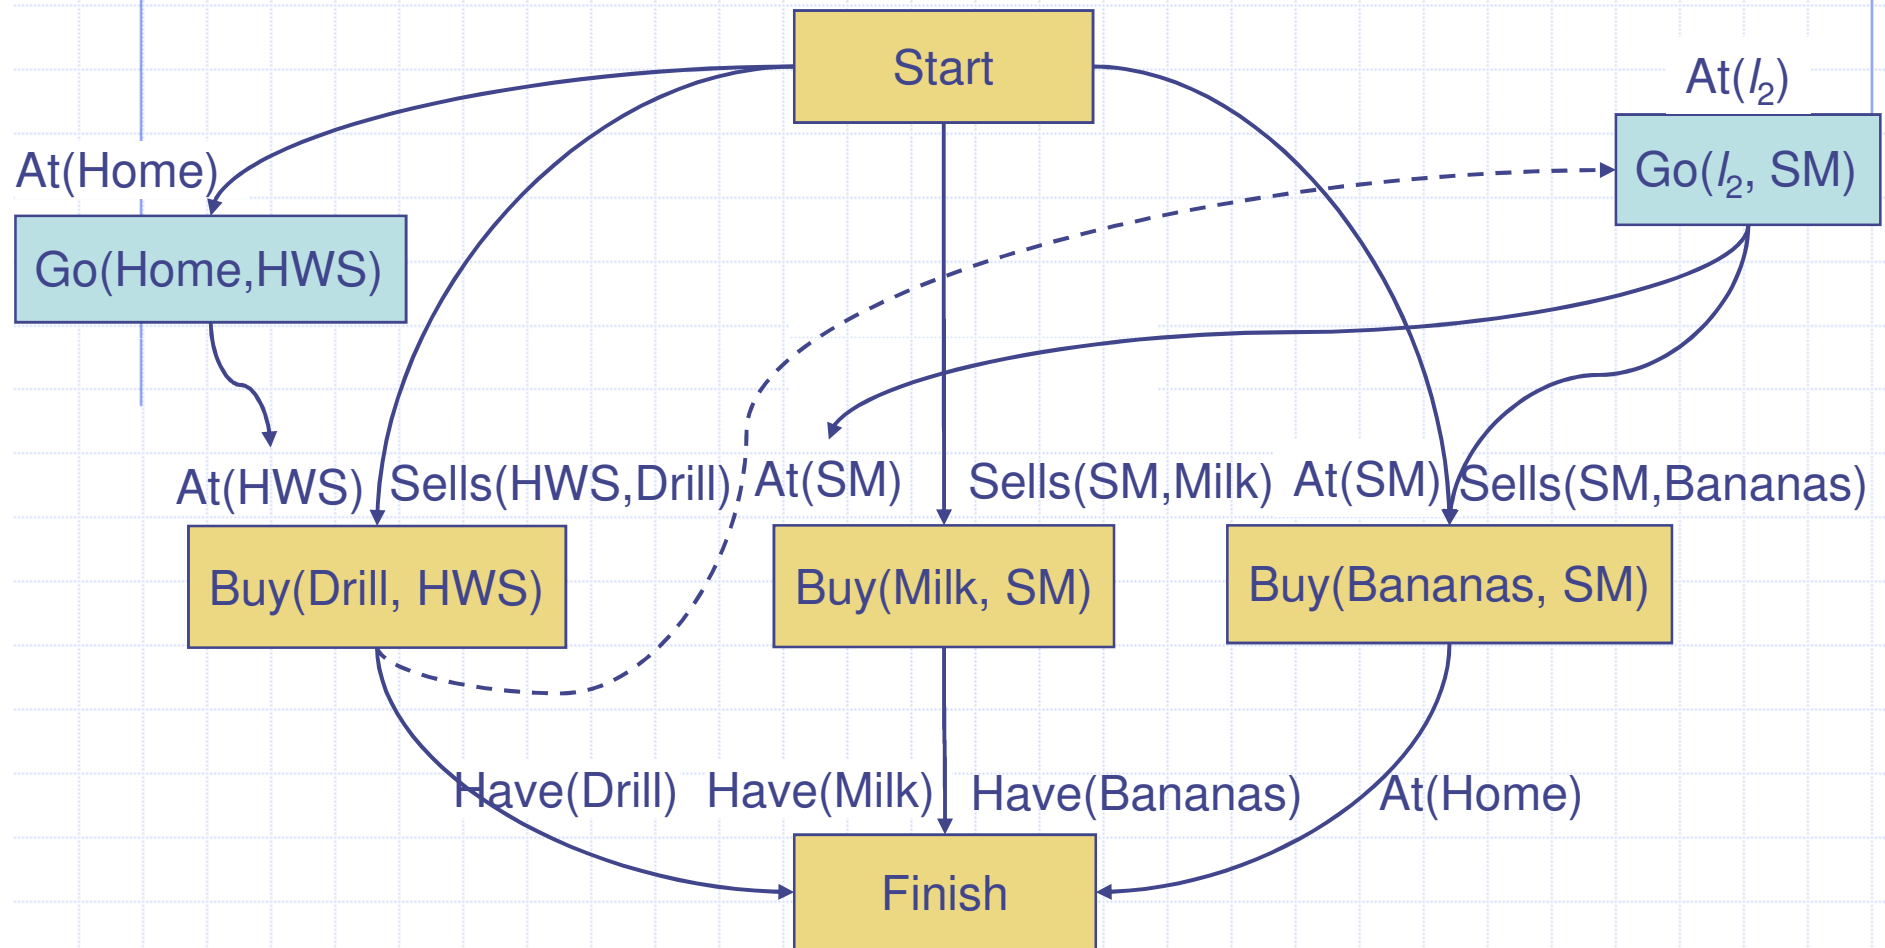
– This time, several threats occur



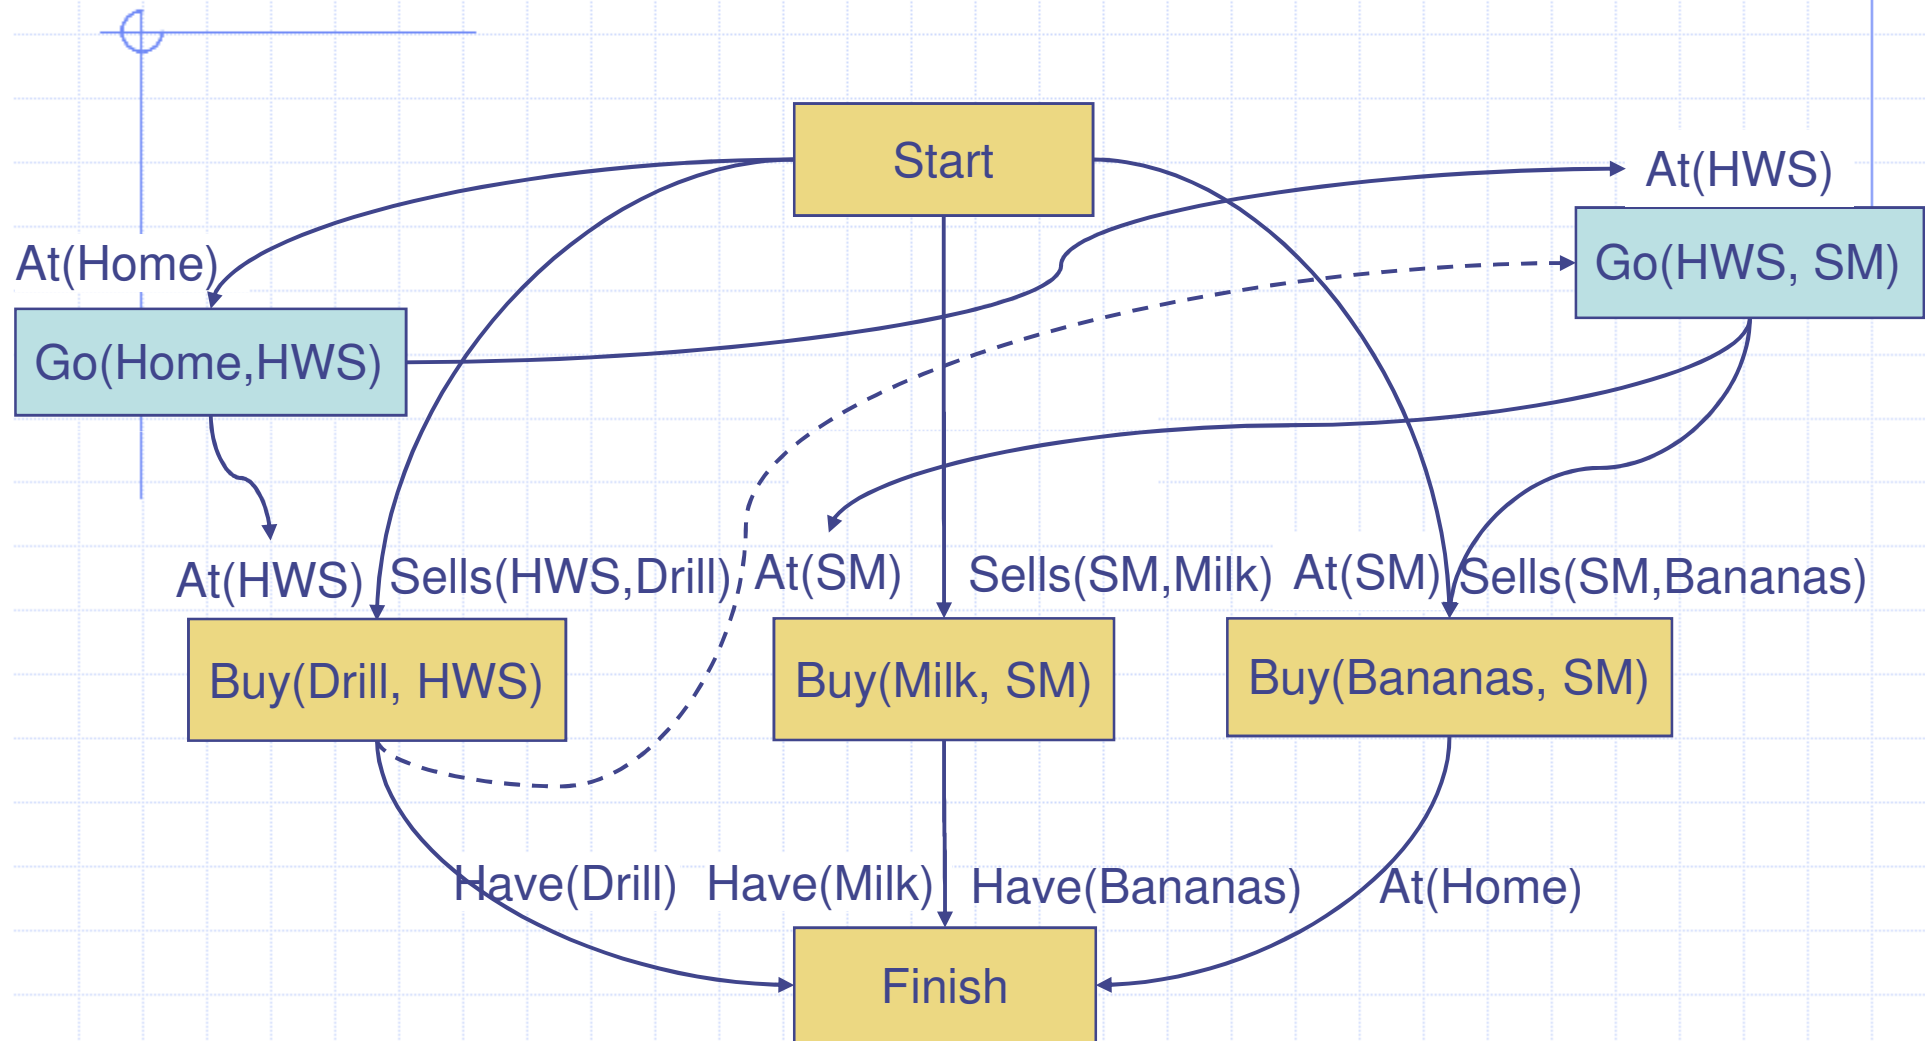
- Finally, a nondeterministic choice: how to resolve the threat to $At(s_1)$?
 - Our choice: make $Buy(Drill)$ precede $Go(SM)$
 - This also resolves the other two threats



- Nondeterministic choice: how to establish $At(I_1)$?
 - We'll do it from Start, with $I_1 = \text{Home}$

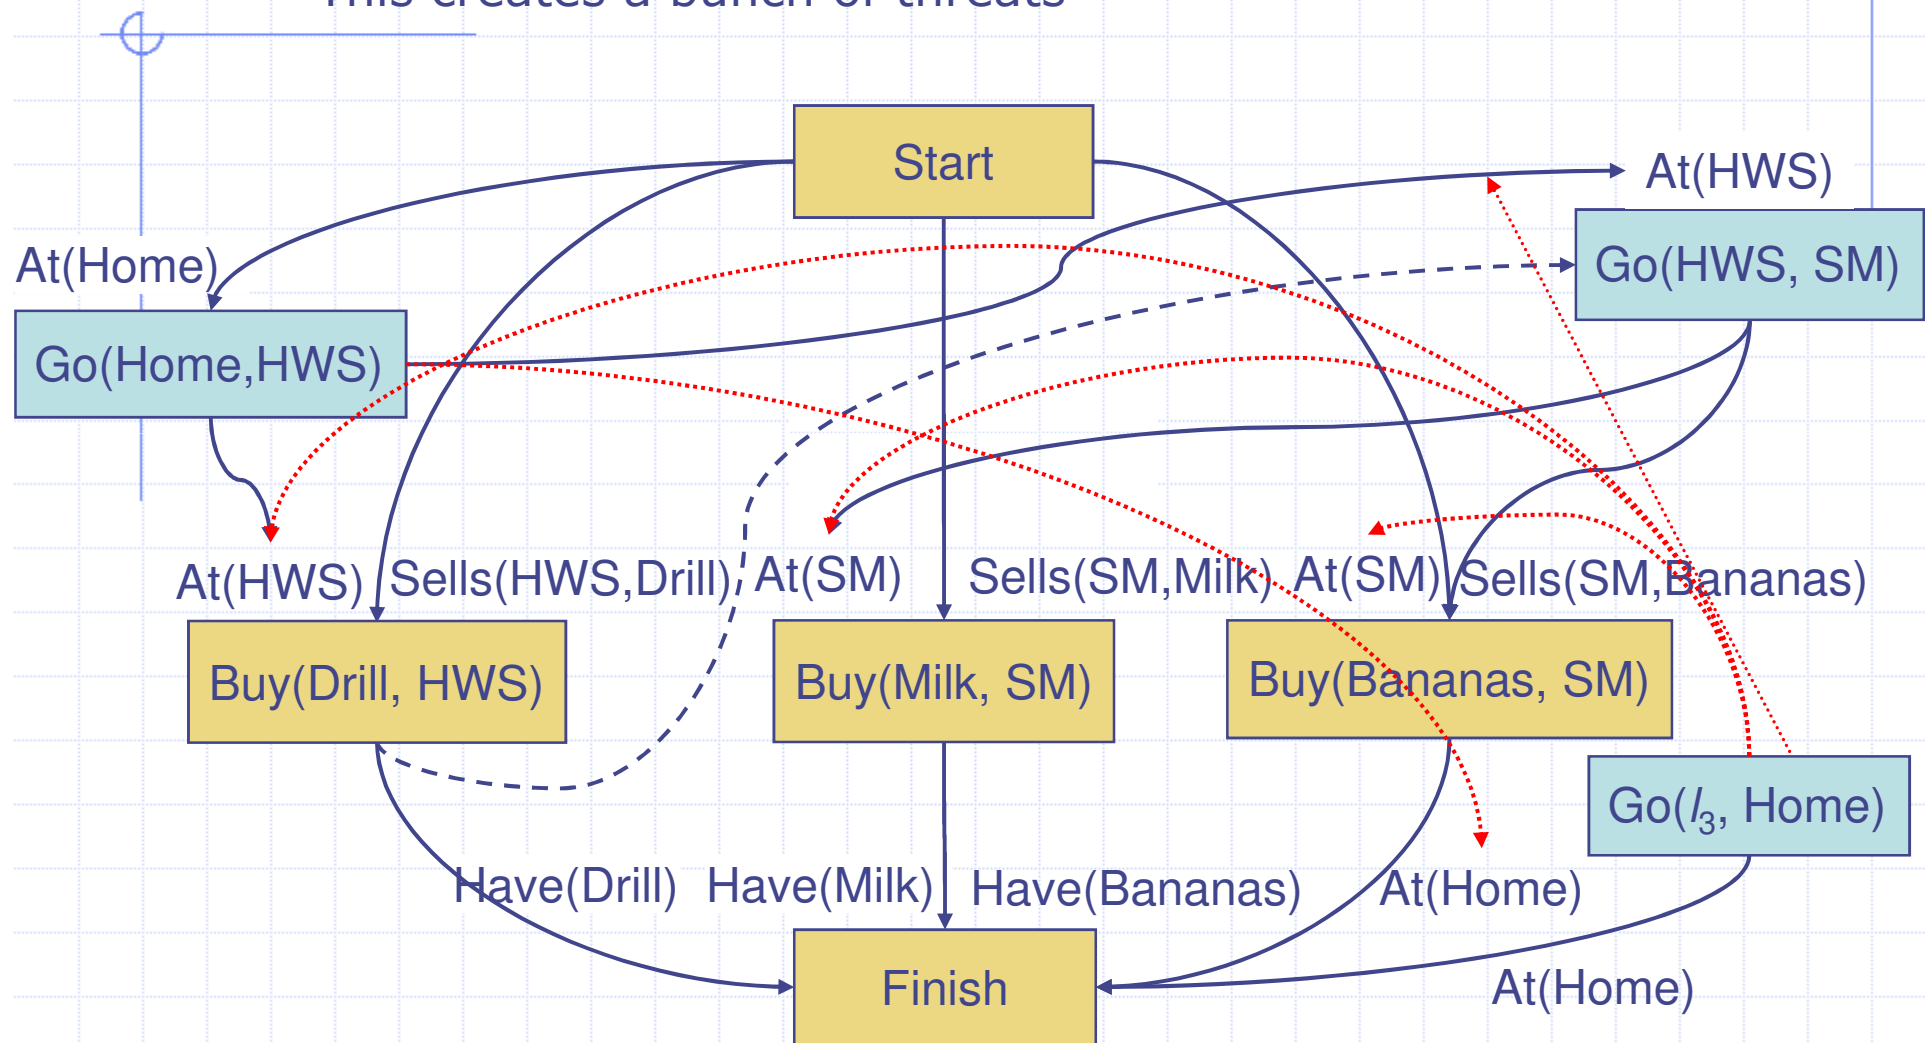


- Nondeterministic choice: how to establish $At(I_2)$?
 - We'll do it from $Go(Home, HWS)$, with $I_2 = HWS$

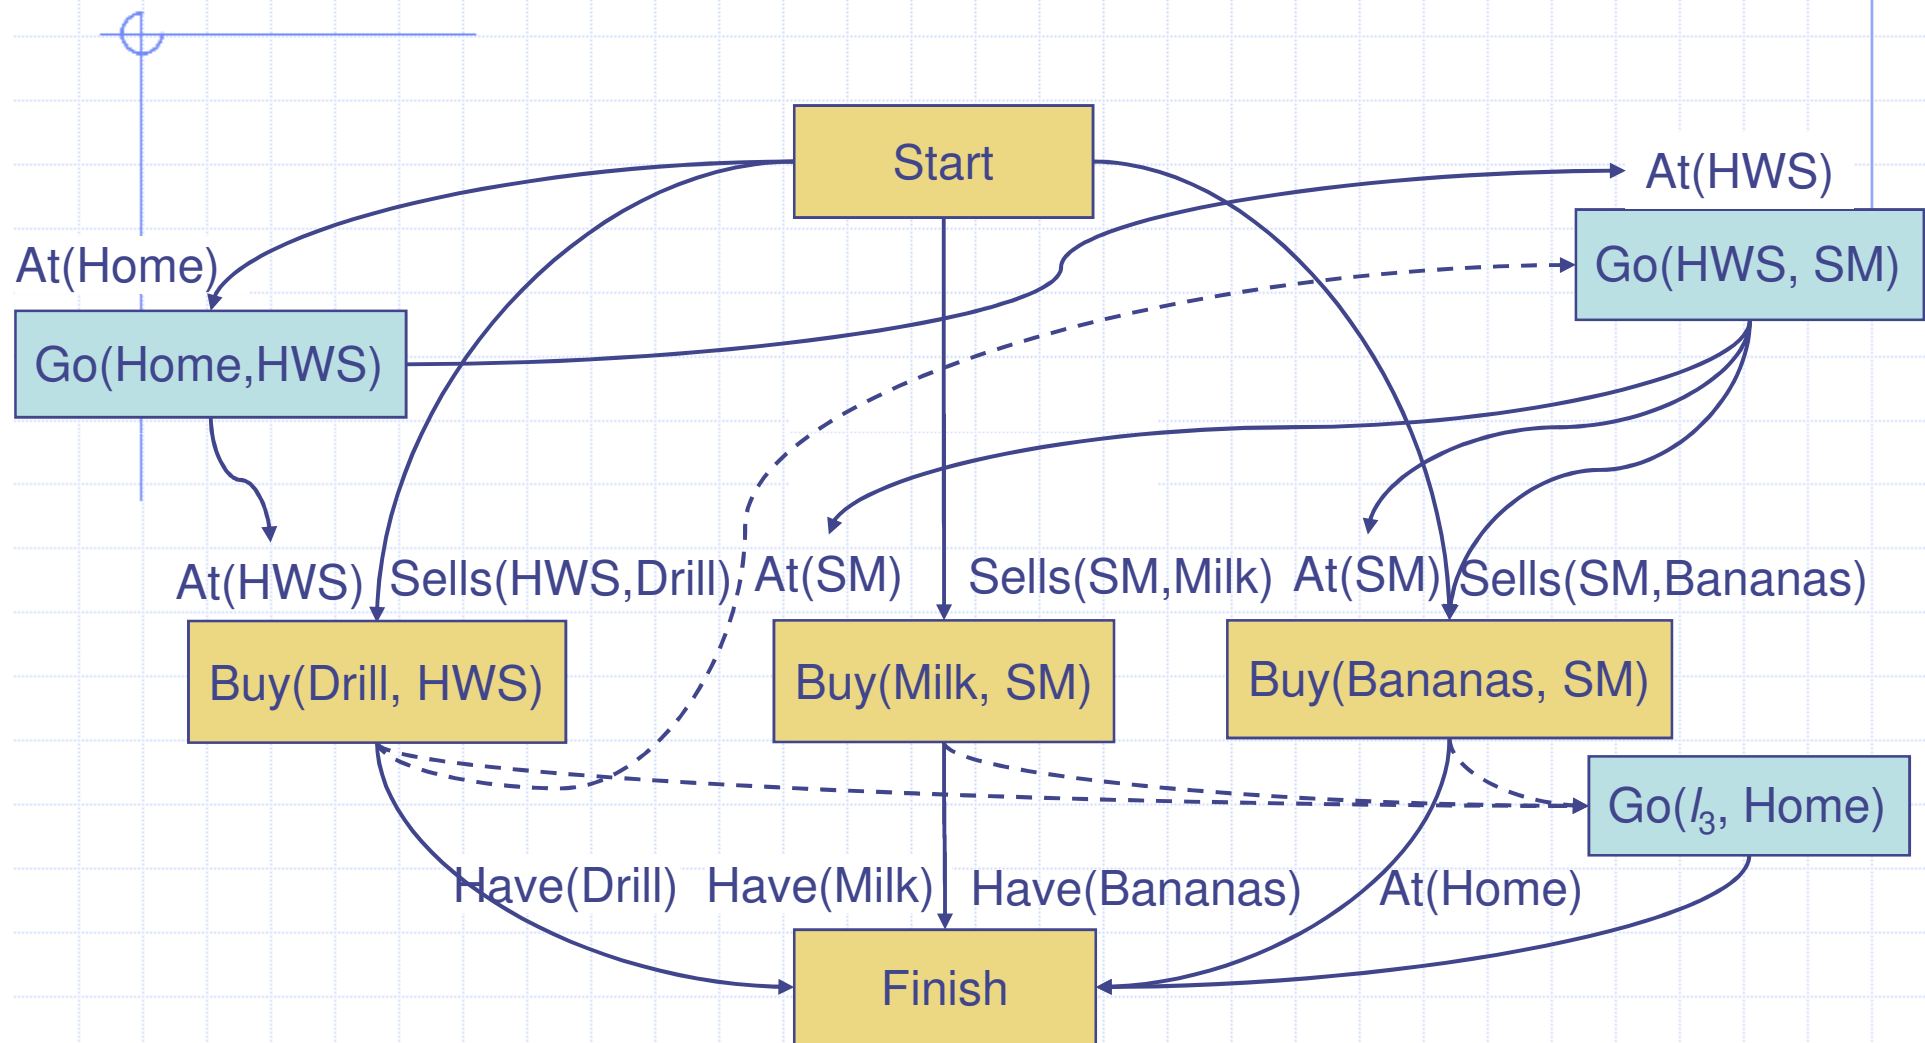


Example (continued)

- The only possible way to establish $\text{At}(\text{Home})$ for Finish
 - This creates a bunch of threats

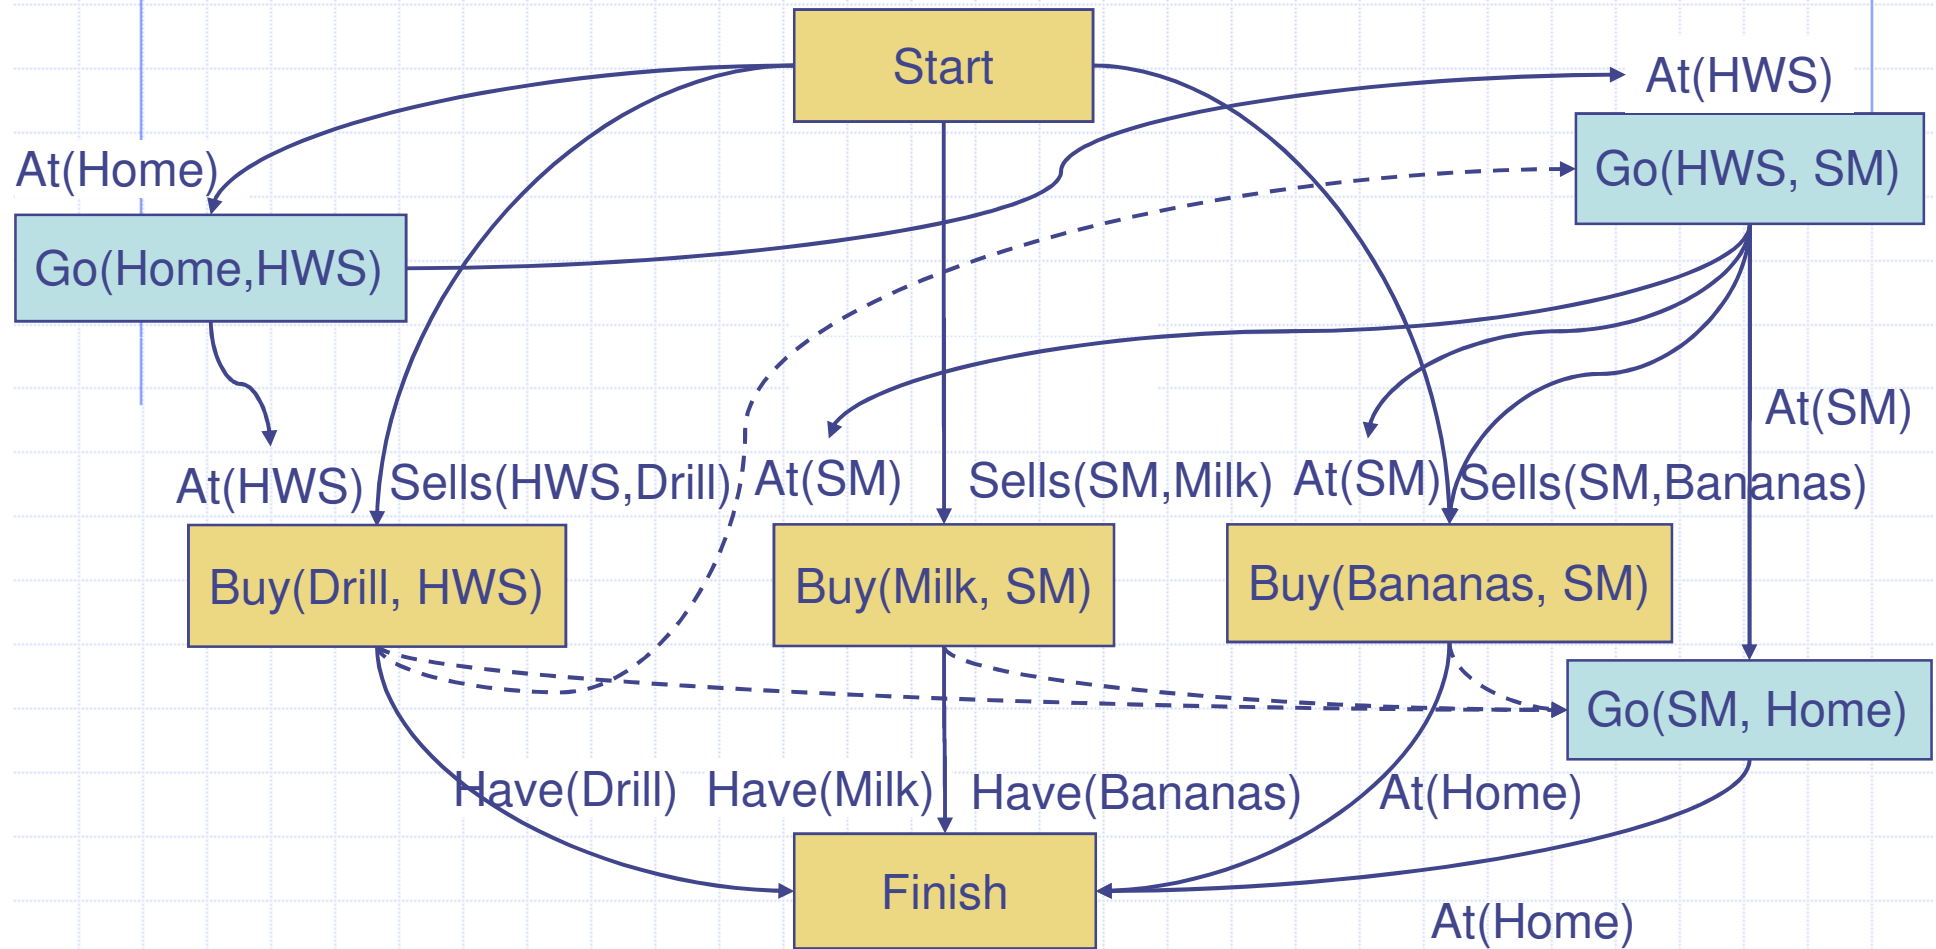


- To remove the threats to $At(SM)$ and $At(HWS)$, make them precede $Go(I_3, Home)$
 - This also removes the other threats

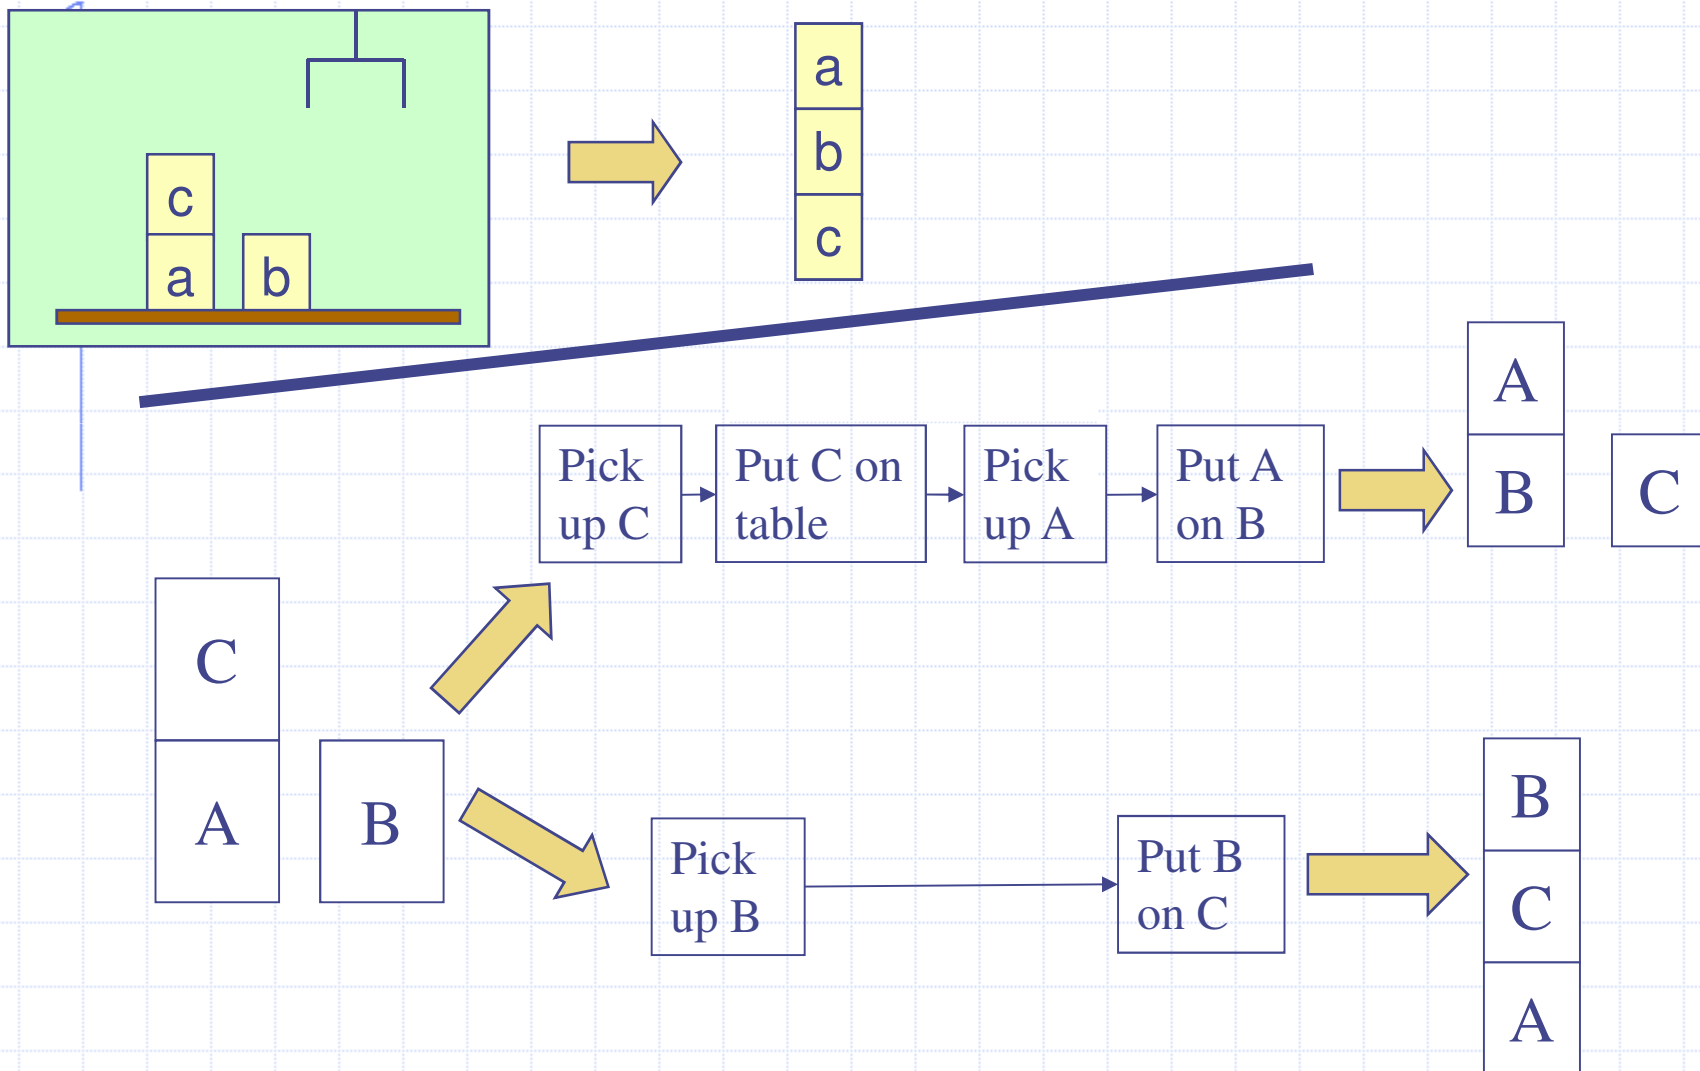


Final Plan




- Establish $At(I_3)$ with $I_3=SM$



Reminder: STRIPS and the Sussman Anomaly



POP and the Sussman Anomaly



Pick
up C

Put C on
table

Pick
up A


Put A
on B



Can we sort these
steps
to achieve these
two goals?


Pick
up B

Put B
on C



Summary

- How to choose which flaw to resolve first and how to resolve it?
 - Heuristics
- PSP doesn't commit to orderings and instantiations until necessary
 - Avoids generating search trees like this one:
- Problem: how to prune infinitely long paths?
 - Loop detection is based on recognizing states we've seen before
 - In a partially ordered plan, we don't know the states
- Can we prune if we see the same *action* more than once?
 - ... — go(b,a) — go(a,b) — go(b,a) — at(a)
 - No. Sometimes we might need the same action several times in different states of the world.



Heuristics in plan-space planning

PSP(π)

$flaws \leftarrow \text{OpenGoals}(\pi) \cup \text{Threats}(\pi)$

if $flaws = \emptyset$ then return(π)

select any flaw $\phi \in flaws$

$resolvers \leftarrow \text{Resolve}(\phi, \pi)$

if $resolvers = \emptyset$ then return(failure)

nondeterministically choose a resolver $\rho \in resolvers$

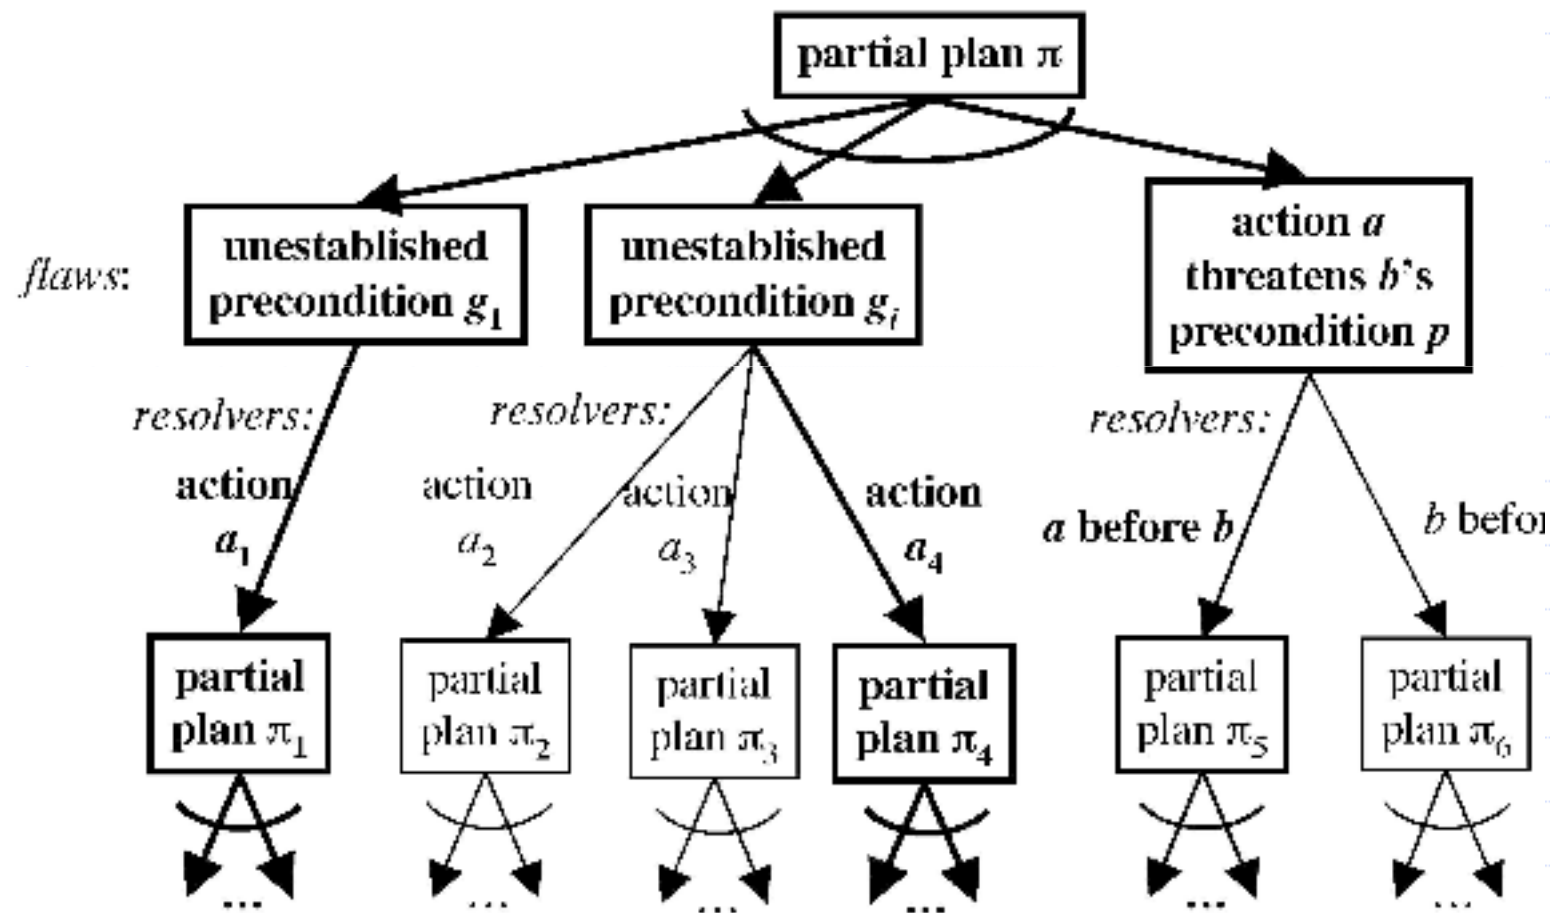
$\pi' \leftarrow \text{Refine}(\rho, \pi)$

return(PSP(π'))

end

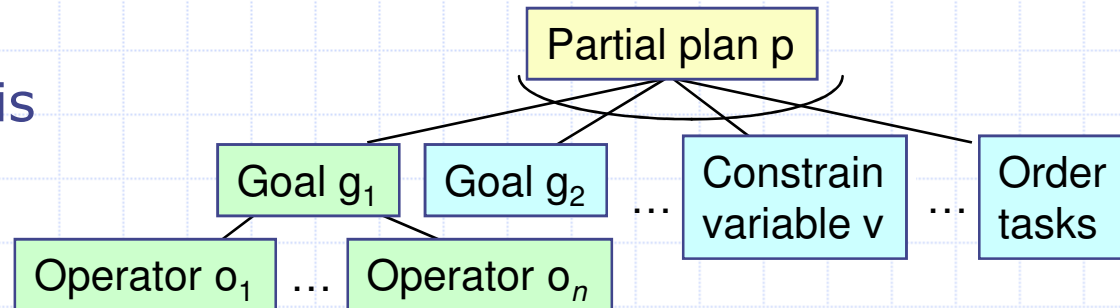
- Flaw-selection heuristic
- Resolver-selection heuristic

Heuristics in plan-space planning

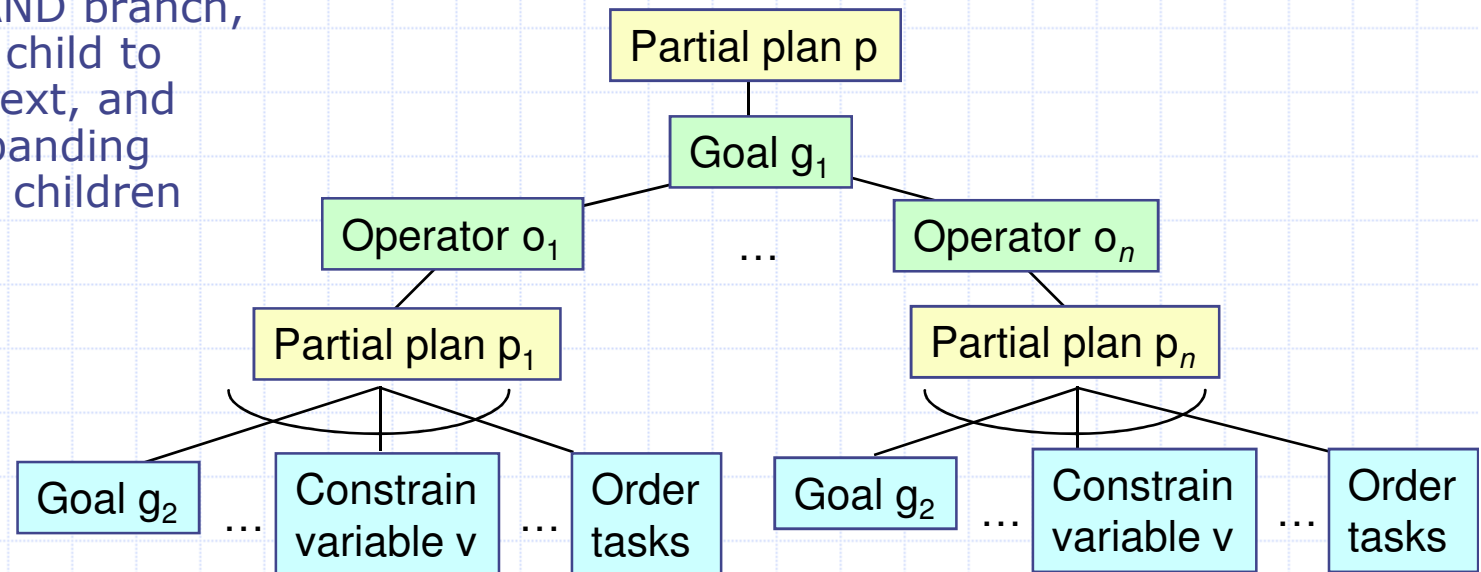


Serializing and AND/OR Tree

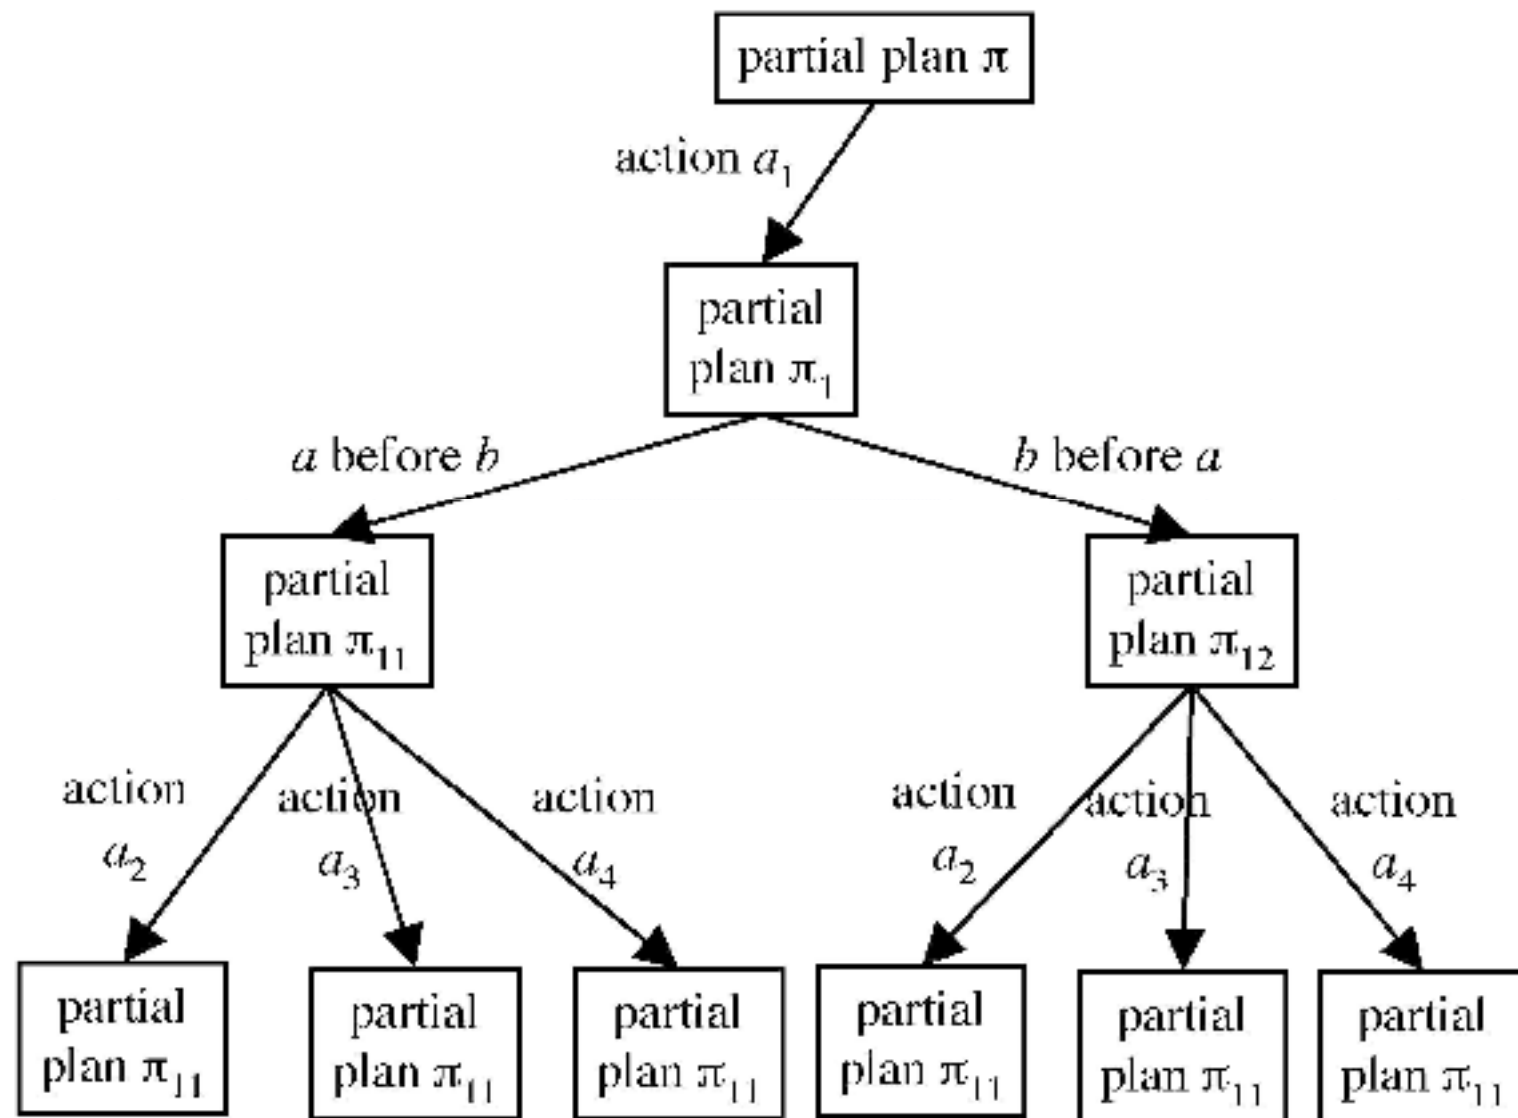
- The search space is an AND/OR tree



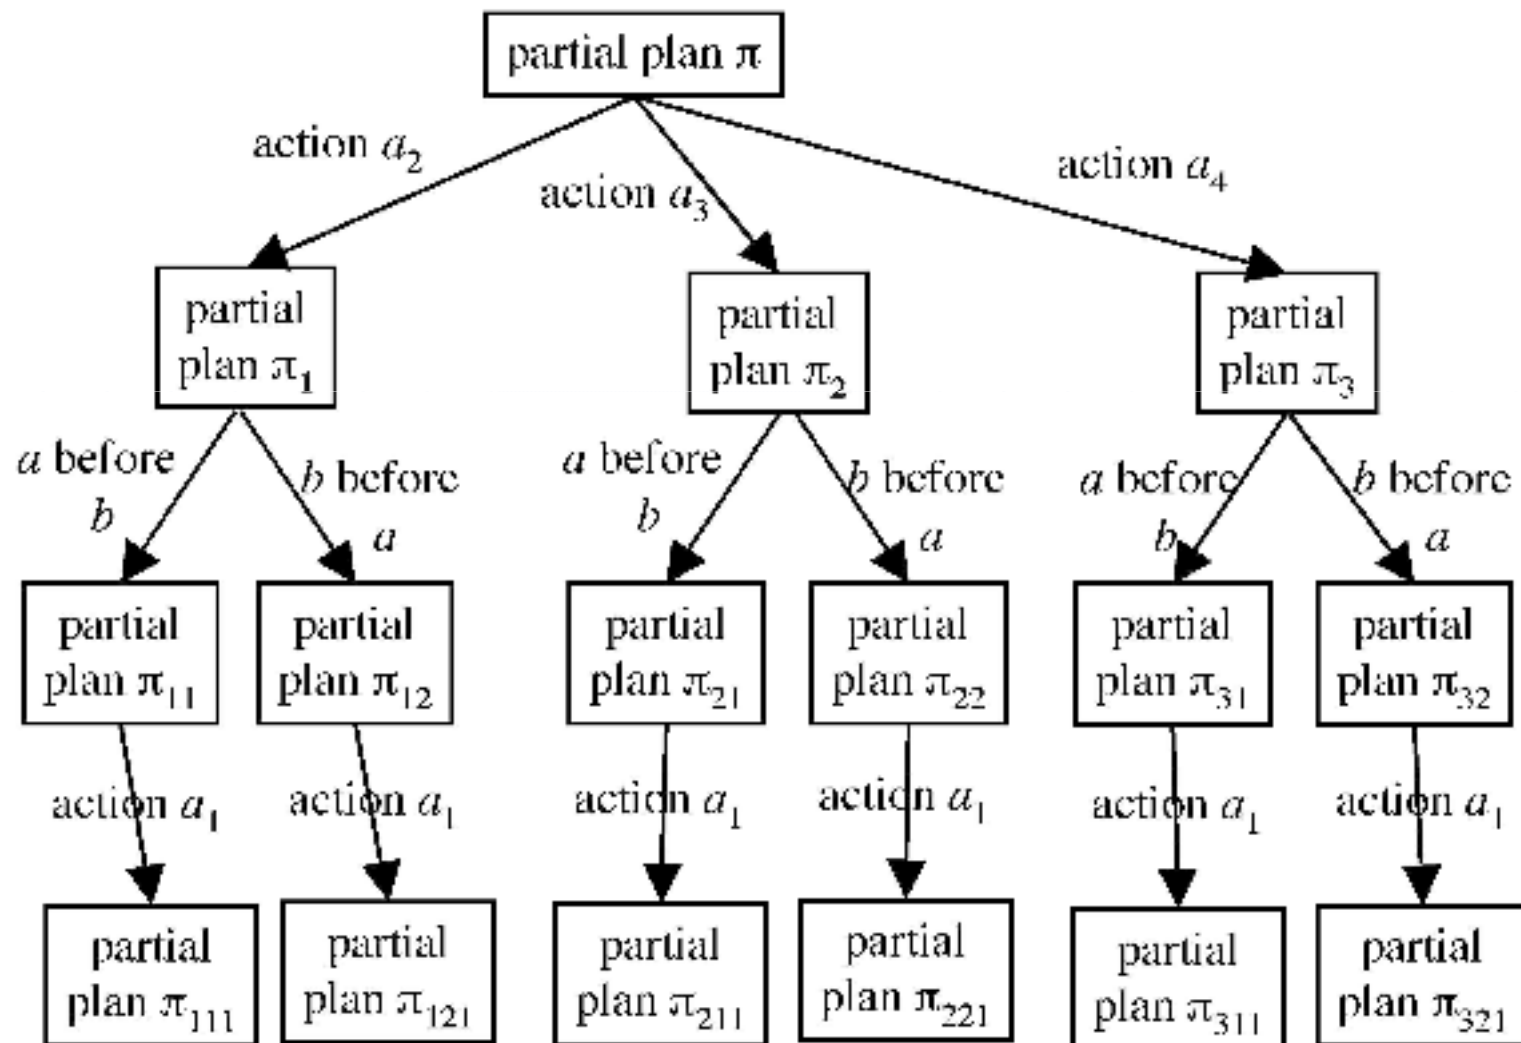
- Deciding what flaw to work on next = *serializing* this tree (turning it into a plan-space tree)
 - at each AND branch, choose a child to expand next, and delay expanding the other children



One Serialization (order of resolution: g_1, threat, g_2)



Another Serialization (order of resolution: g_2 , threat, g_1)



Why Does This Matter?

- Different refinement strategies produce different serializations
 - the search spaces have different numbers of nodes
- In the worst case, the planner will search the entire serialized search space
- The smaller the serialization, the more likely that the planner will be efficient

Why Does This Matter?

- A particular goal 'g':
 - 'B' resolvers → 'B' partial plans
 - 'a' threats in each plan, 'm' resolvers for each threat
 - Effective branching factor of search by POP: $B_g = B * m^a$
- Time complexity of a solution plan with 'n' steps/actions and 'p' preconditions/goals per step:
 - $O(B_g)^{np}$
- Exponential number of nodes, so in order for PSP to be efficient we need:
 - Good flaw-selection heuristics
 - Good resolver-selection heuristics

Flaw-selection heuristics

- Fewest Alternatives First (FAF)
 - Choose the flaw that has the smallest number of alternatives
 - In this case, choose to resolve precondition g_1
- Last-In, First-out (LIFO)
- Delay potential threats
- Zero-LIFO (ZLIFO)

Resolver-selection heuristics

- Select the least cost partial plan ($\Pi_1, \Pi_2, \dots, \Pi_n$)
- Use an A* algorithm $f(\Pi)=g(\Pi)+h(\Pi)$ where:
 - $g(\Pi)$ is the cost of the path from the root node (initial empty plan) to plan Π
 - $h(\Pi)$ is an estimate of the least cost of the path from Π to a solution plan (goal state)

Node-Selection Heuristic

- Suppose we're searching a **tree** in which each edge (s, s') has a cost $c(s, s')$


- If p is a path, let $c(p)$ = sum of the edge costs
- For classical planning, this is the length of p

- For every state s , let

- $g(s)$ = cost of the path from s_0 to s
- $h^*(s)$ = least cost of all paths from s to goal nodes
- $f^*(s) = g(s) + h^*(s)$ = least cost of all paths from s_0 to goal nodes that go through s

- Suppose $h(s)$ is an estimate of $h^*(s)$

- Let $f(s) = g(s) + h(s)$
 - $f(s)$ is an estimate of $f^*(s)$
- h is *admissible* if for every state s , $0 \leq h(s) \leq h^*(s)$
- If h is admissible then f is a lower bound on f^*



The A* Algorithm

- A^* on trees:

loop


choose the leaf node S such that $f(S)$ is smallest
if s is a solution then return it and exit
expand it (generate its children)

- On graphs, A^* is more complicated

- additional machinery to deal with multiple paths to the same node

- If a solution exists (and certain other conditions are satisfied), then:

- If $h(s)$ is admissible, then A^* is guaranteed to find an optimal solution
- The more "informed" the heuristic is (i.e., the closer it is to h^*), the smaller the number of nodes A^* expands
- If $h(s)$ is within c of being admissible, then A^* is guaranteed to find a solution that's within c of optimal



Resolver-selection heuristics

- Elements of a plan Π that can take part of the function $f(\Pi)=g(\Pi)+h(\Pi)$:
 - $N(\Pi)$: number of steps in plan Π
 - $OP(\Pi)$: open goals/preconditions in plan Π
 - $T(\Pi)$: threats in plan Π
- Usually $f(\Pi)=N(\Pi)+OP(\Pi)$:
 - Does this 'f' function guarantee to find an optimal solution, i.e., is it an A* algorithm?