

# AI Planning.

## Planning-graph techniques.

## Heuristic planning.



Eva Onaindia

Universitat Politècnica de València

# Acknowledgements

Most of the slides used in this course are taken or are modifications from Dana Nau's lecture slides for the textbook *Automated Planning*, licensed under the Creative Commons Attribution-NonCommercial-ShareAlike License:

<http://creativecommons.org/licenses/by-nc-sa/2.0/>

<http://creativecommons.org/licenses/by/3.0/es/>

I would like to gratefully acknowledge Dana Nau's contributions and thank him for generously permitting me to use aspects of his presentation material.

# Outline




- History
- Motivation
- Reachability analysis
- The planning graph
  - Example
  - Mutual exclusion
  - Example (continued)
- The Graphplan algorithm
- Solution extraction
- Heuristics derived from planning graphs
- Discussion

# History

- Before Graphplan came out, most planning researchers were working on PSP-like planners
  - POP, SNLP, UCPOP, etc.
- Graphplan caused a sensation because it was so much faster
- Many subsequent planning systems have used ideas from it
  - IPP, STAN, GraphHTN, SGP, Blackbox, Medic, TGP, LPG
  - Many of them are much faster than the original Graphplan

# Motivation

- A big source of inefficiency in search algorithms is the *branching factor*
  - the number of children of each node
- e.g., a backward search may try lots of actions that can't be reached from the initial state
- and forward search may try lots of actions that do not reach the goal state



# Motivation

- One way to reduce branching factor:
- First create a *relaxed problem*
  - Remove some restrictions of the original problem
    - Want the relaxed problem to be easy to solve (polynomial time)
  - The solutions to the relaxed problem will include all solutions to the original problem
    - actions that occur in the solutions to the relaxed problem will also occur in all solutions to the original problem
- Then do a modified version of the original search
  - Restrict its search space to include only those actions that occur in solutions to the relaxed problem

# Motivation

- State-space planners provide a plan as a sequence of actions
- Plan-space planning synthesize a plan as a partially ordered set of actions; any sequence that meets the constraints of the partial order is a valid plan
- Planning-graph approaches take a middle ground. Their output is a sequence of sets of actions, e.g.  $\langle \{a_1, a_2\}, \{a_3, a_4\}, \{a_5, a_6, a_7\} \rangle$
- Plan-space planning: *least-commitment principle*  
Planning-graph approach: *strong commitment* while planning, actions are considered fully instantiated and at specific time steps; *least-commitment* with respect to the order of the actions at a time step
- Planning-graph approaches rely on the idea of ***relaxation of the reachability analysis***

# Reachability analysis

- Analysis performed by planners to make sure the goals of the problem are reachable
- State reachability:
  - given a set **A** of actions, a state **s** is reachable from some initial state **s0** if there is a sequence of actions in **A** that defines a path from **s0** to **s**
  - Reachability analysis consists of analyzing which states can be reached from **s0** in some number of steps and how to reach them
  - Reachability can be computed *exactly* through a *reachability tree* => it cannot be computed in a tractable way
  - Reachability can be *approximated* through a *planning graph* => relaxation of the reachability analysis

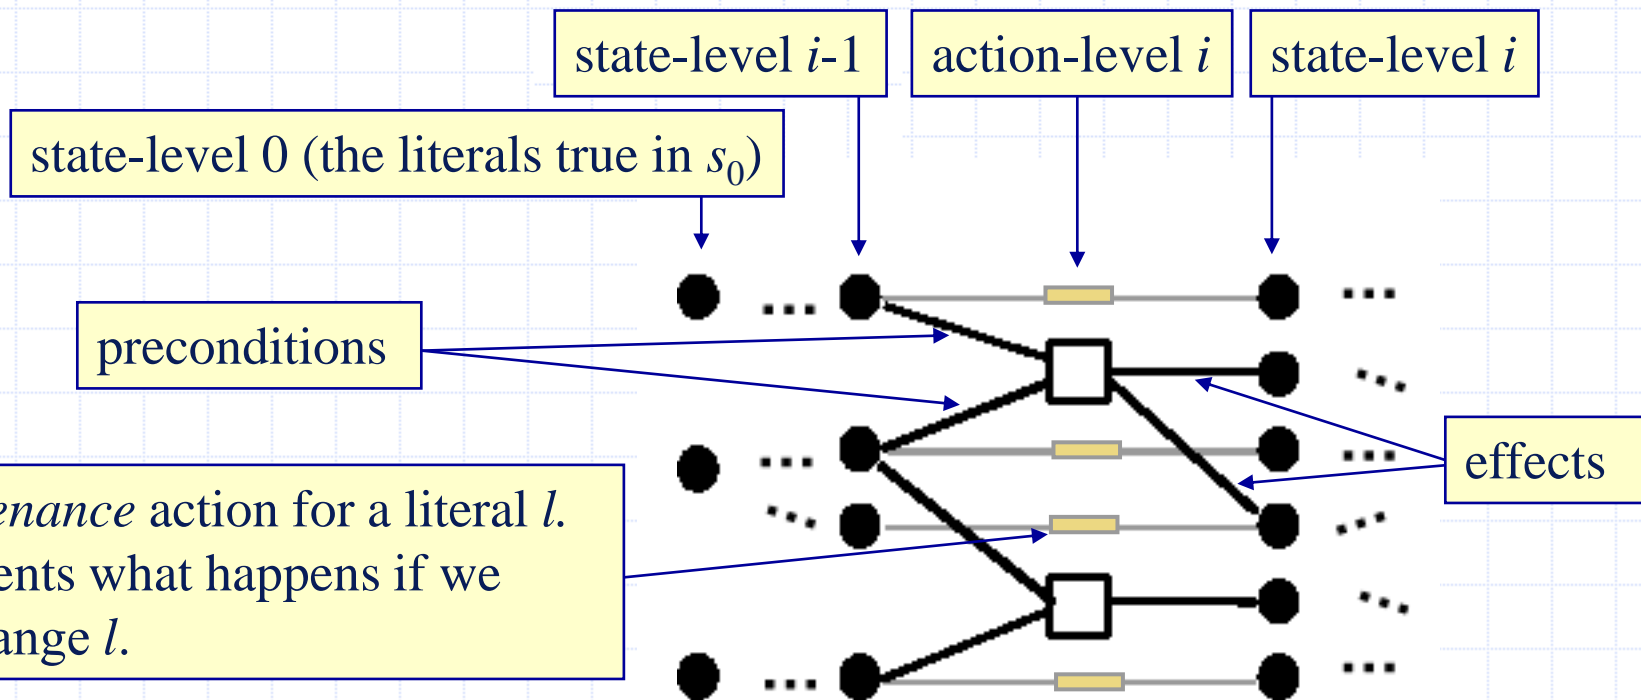


# The planning graph (I)

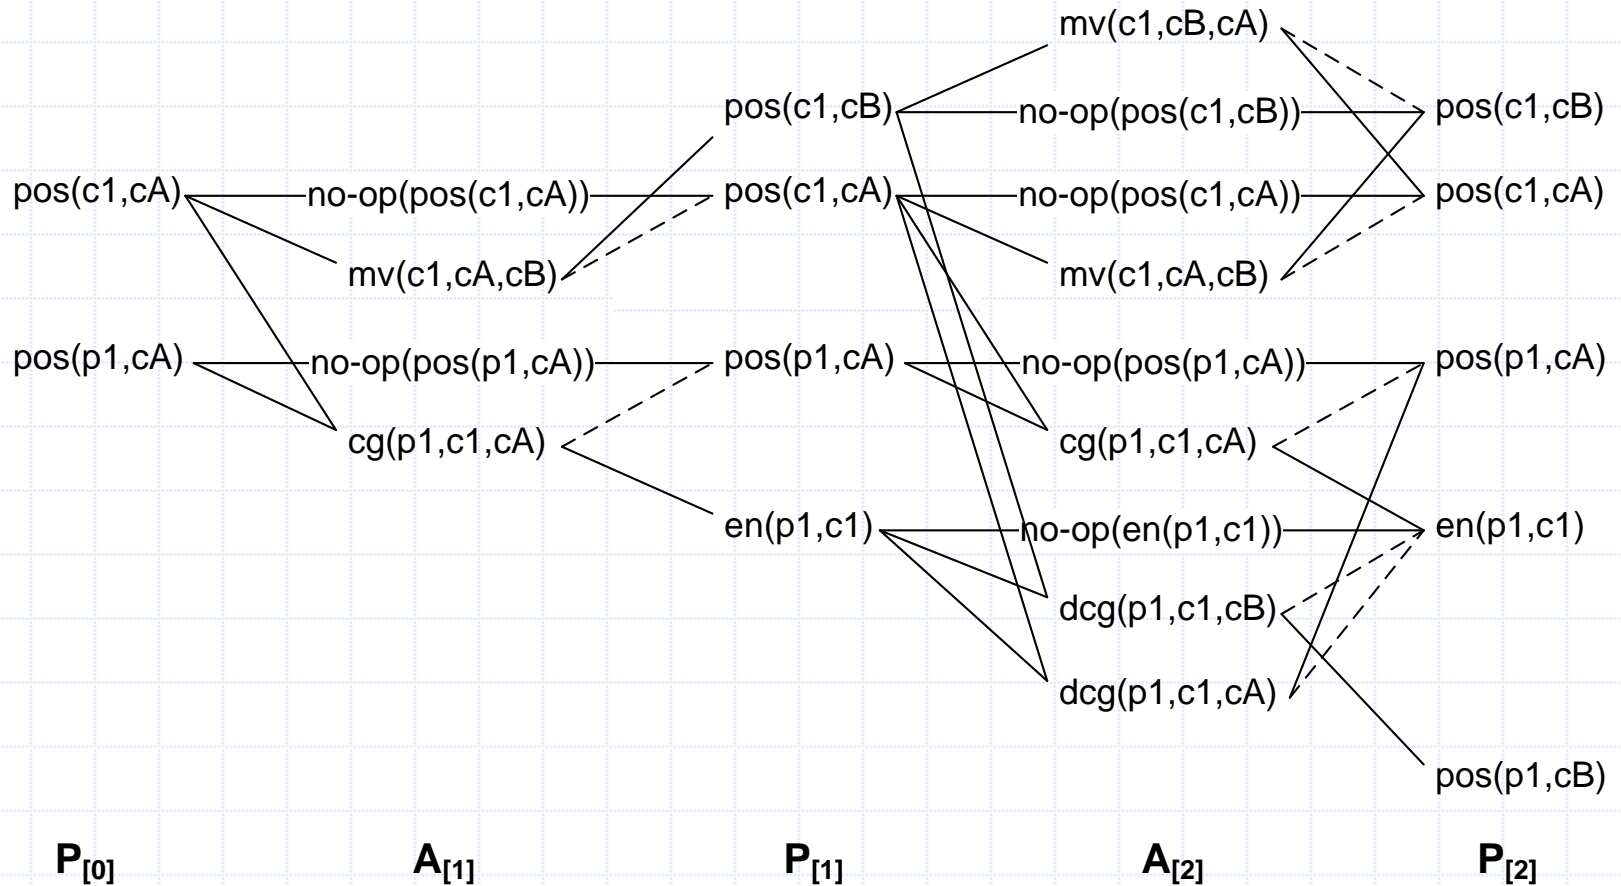
- A layered graph
- Two kinds of layers alternate:
  - Literals (propositions) (shown with circles)
  - Actions (shown with squares)
- Every two layers correspond to a discrete time (time step)
- No variables as in action schemas
- The first layer is a literal layer which shows all the literals that are true in the initial state
- Every action has a link from each of its preconditions and a link to each of its effects.
- Straight lines between two literals at consecutive literal levels denote *NoOp*

# The planning graph (II)


- Search space for a relaxed version of the planning problem
- Alternating layers of ground literals and actions
  - Nodes at action-level  $i$ : actions that might be possible to execute at time  $i$
  - Nodes at state-level  $i$ : literals that might possibly be true at time  $i$
  - Edges: preconditions and effects



# Example



# Mutual Exclusion




- Two actions at the same action-level are mutex if
  - *Inconsistent effects*: an effect of one negates an effect of the other
  - *Interference*: one deletes a precondition of the other
  - *Competing needs*: **they have mutually exclusive preconditions**
- Otherwise they don't interfere with each other
  - Both may appear in a solution plan
- Two literals at the same state-level are mutex if
  - *Inconsistent support*: one is the negation of the other, **or all ways of achieving them are pairwise mutex**

Recursive  
propagation  
of mutexes

## Example (level A[1])

**A[1]**

$mv(c1, cA, CB) \times \{cg(p1, c1, cA), no-op(pos(c1, cA))\}$   
 $cg(p1, c1, cA) \times \{no-op(pos(p1, cA)), mv(c1, cA, CB)\}$



$P_{[0]}$

$A_{[1]}$

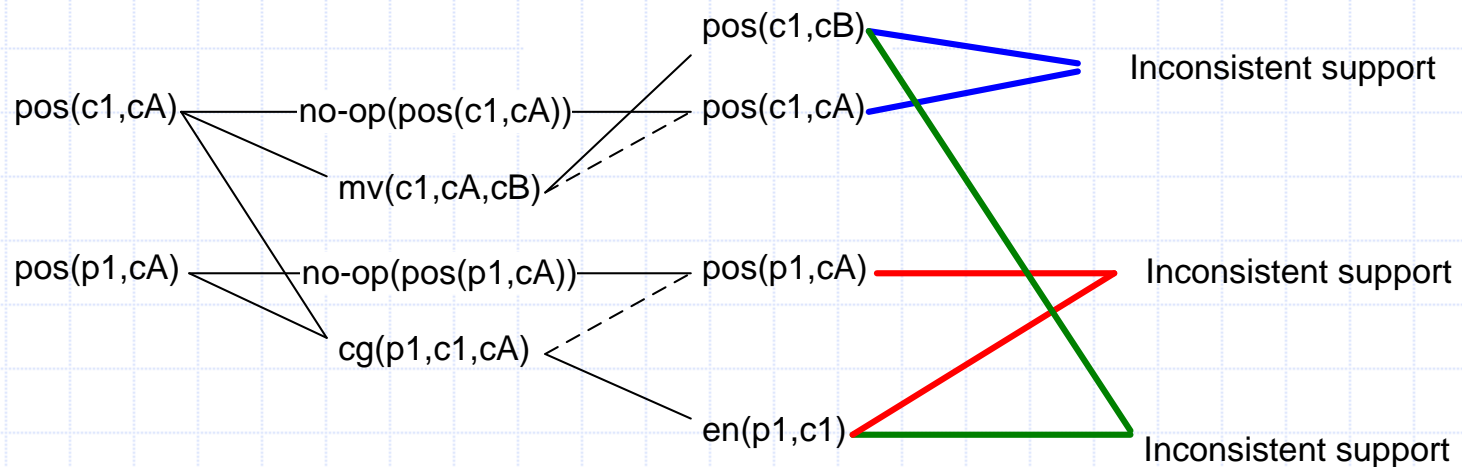
## Example (level P[1])

**P[1]**

$\text{en}(p1, c1) \times \{\text{pos}(p1, cA), \text{pos}(c1, cB)\}$

$\text{pos}(c1, cA) \times \{\text{pos}(c1, cB)\}$

$\text{pos}(c1, cB) \times \{\text{en}(p1, c1), \text{pos}(c1, cA)\}$

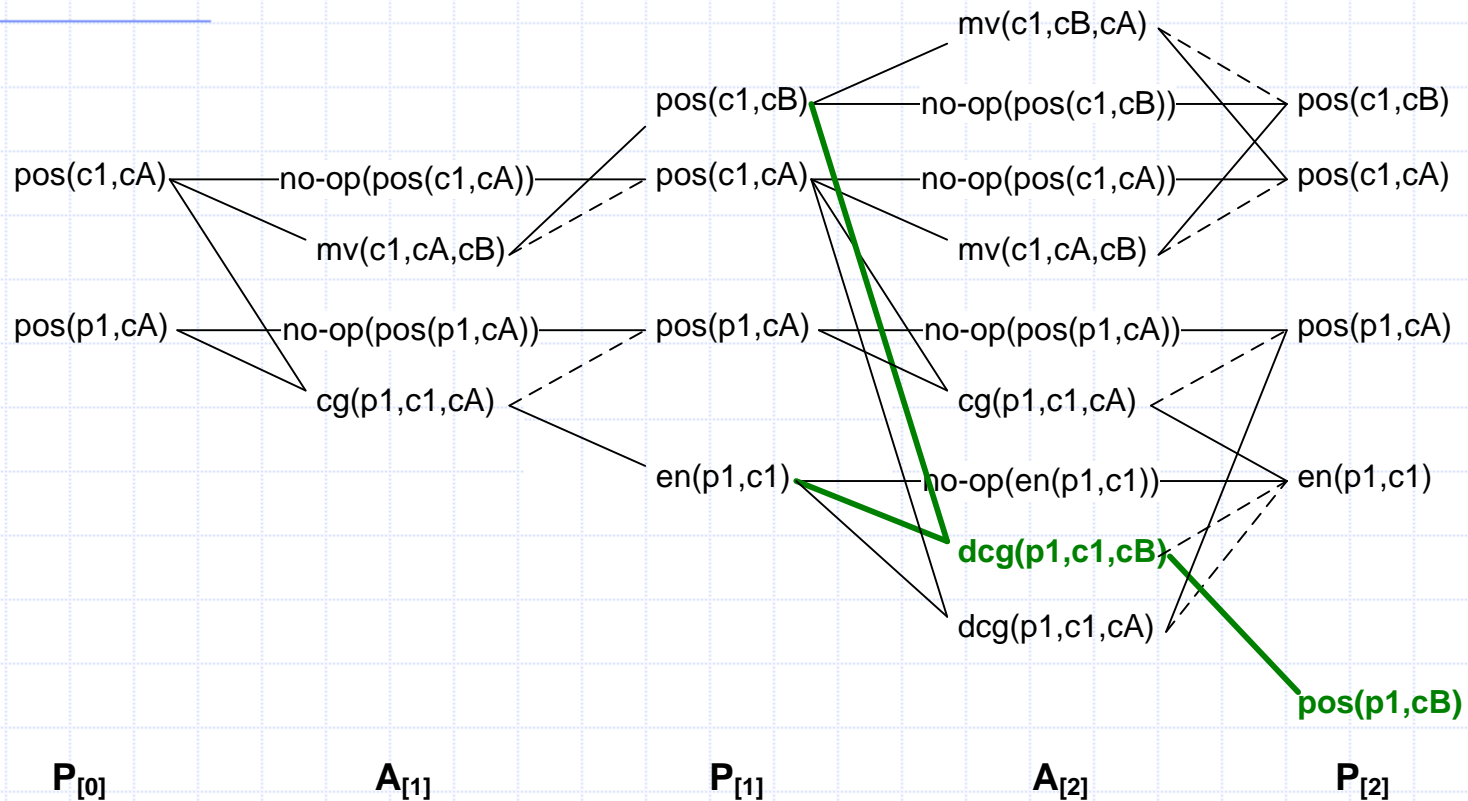


$P_{[0]}$

$A_{[1]}$

$P_{[1]}$

## Example (level A[2])



Action  **$dcg(p1,c1,cB)$**  would not appear in  **$A_{[2]}$**  because its preconditions are inconsistent. Consequently, literal  $pos(p1,cB)$  would not appear in  **$P_{[2]}$** , which is an indication that the goal cannot be achieved at level  **$P_{[2]}$**  and so the graph must be extended one more level ( **$A_{[3]}$ ,  $P_{[3]}$** )

## Example (level A[2])

**A[2]**

**mv(c1,cA,cB)** x

{mv(c1,cB,cA), cg(p1,c1,cA), dcg(p1,c1,cA), no-op(pos(c1,cA)),  
no-op(pos(c1,cB))}

**mv(c1,cB,cA)** x

{no-op(pos(c1,cB)), no-op(en(p1,c1)), no-op(pos(c1,cA)),  
cg(p1,c1,cA), dcg(p1,c1,cA), mv(c1,cA,cB)}

**cg(p1,c1,cA)** x

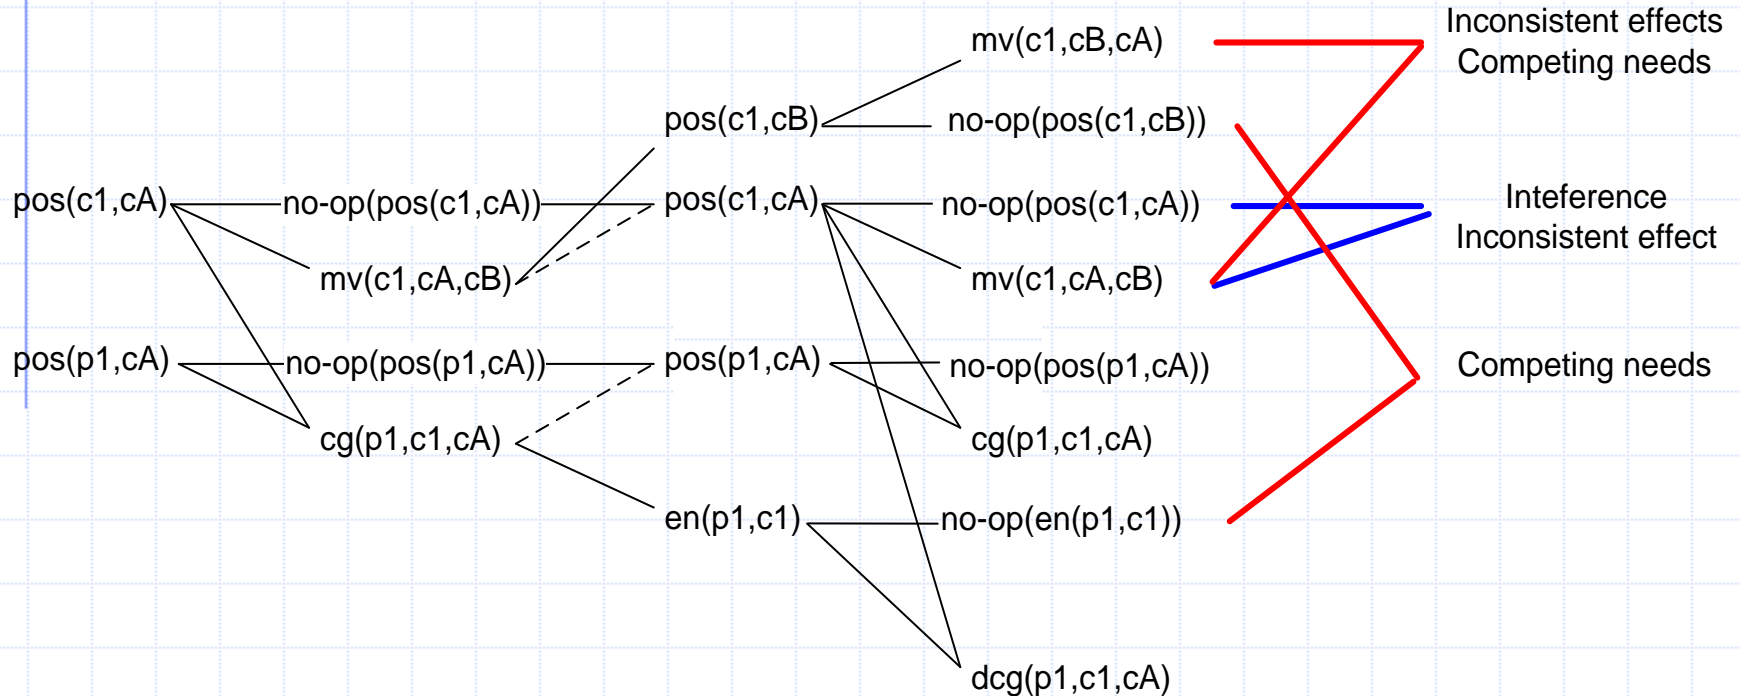
{no-op(pos(p1,cA)), mv(c1,cA,cB), mv(c1,cB,cA), no-op(pos(c1,cB)),  
dcg(p1,c1,cA), no/op(en(p1,c1))}

**dcg(p1,c1,cA)** x

{cg(p1,c1,cA), no-op(en(p1,c1)), mv(c1,cB,cA), no-op(pos(c1,cB)),  
mv(c1,cA,cB), no-op(pos(p1,cA))}



## Example (level A[2])



P<sub>[0]</sub>

A<sub>[1]</sub>

P<sub>[1]</sub>

A<sub>[2]</sub>

# Example (level P[2])

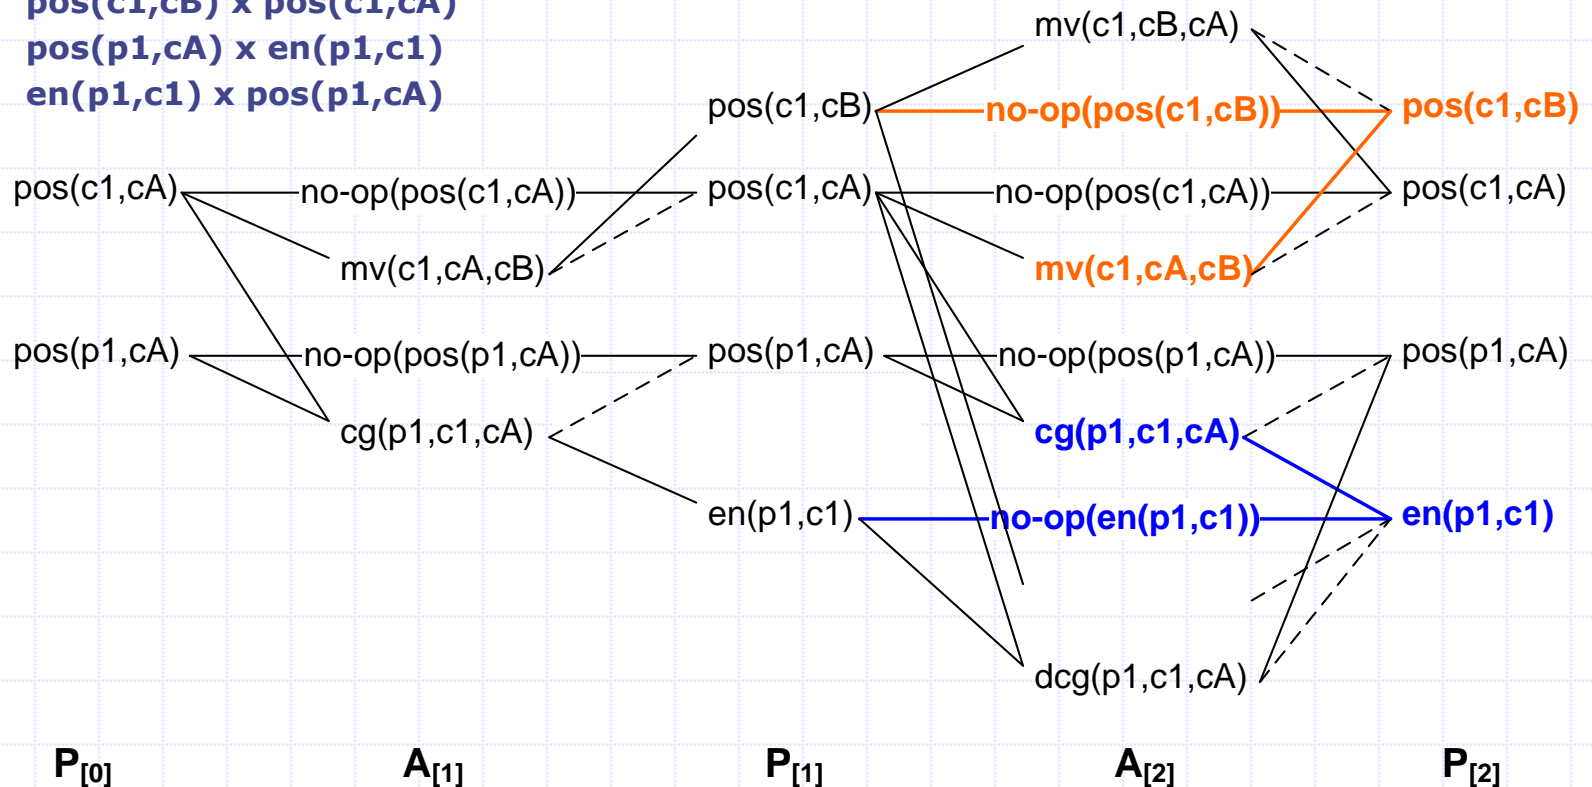
**P[2]**

**pos(c1,cA) x pos(c1,cB)**

**pos(c1,cB) x pos(c1,cA)**

**pos(p1,cA) x en(p1,c1)**

**en(p1,c1) x pos(p1,cA)**



**pos(c1,cB)** and **en(p1,c1)** are no longer mutex because now not all ways of achieving **pos(p1,cB)** are pairwise mutex with all ways of achieving **en(p1,c1)**

## Example (level P[2])

Literals in P[2]

pos(c1,cB)

en(p1,c1)

Actions in A[2] achieving  
the literals

no-op(pos(c1,cB))

mv(c1,cA,cB)

no-op(en(p1,c1))

cg(p1,c1,cA)

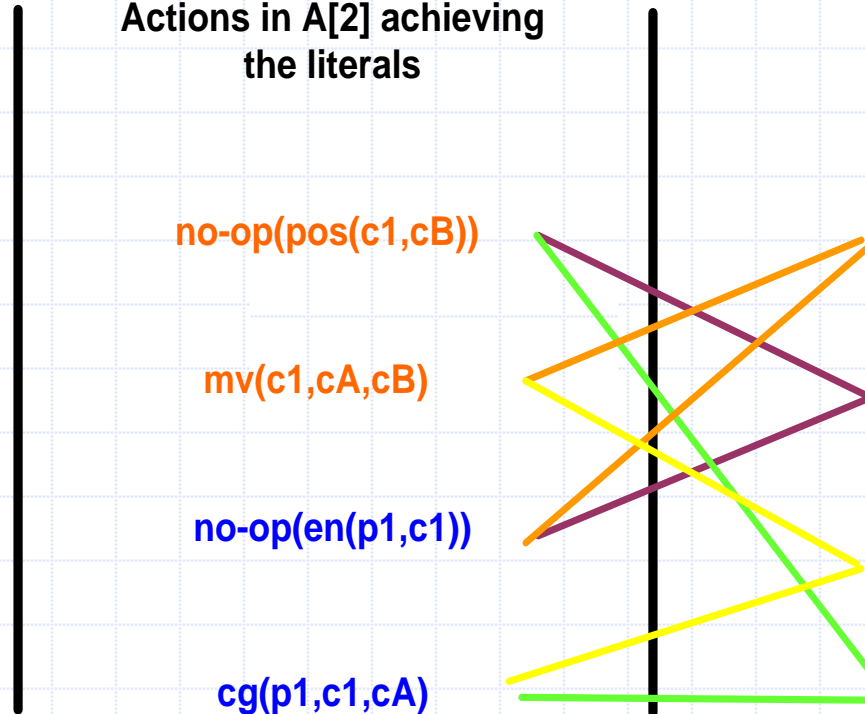
Mutex relations

Non-mutex

Competing needs in A[2]

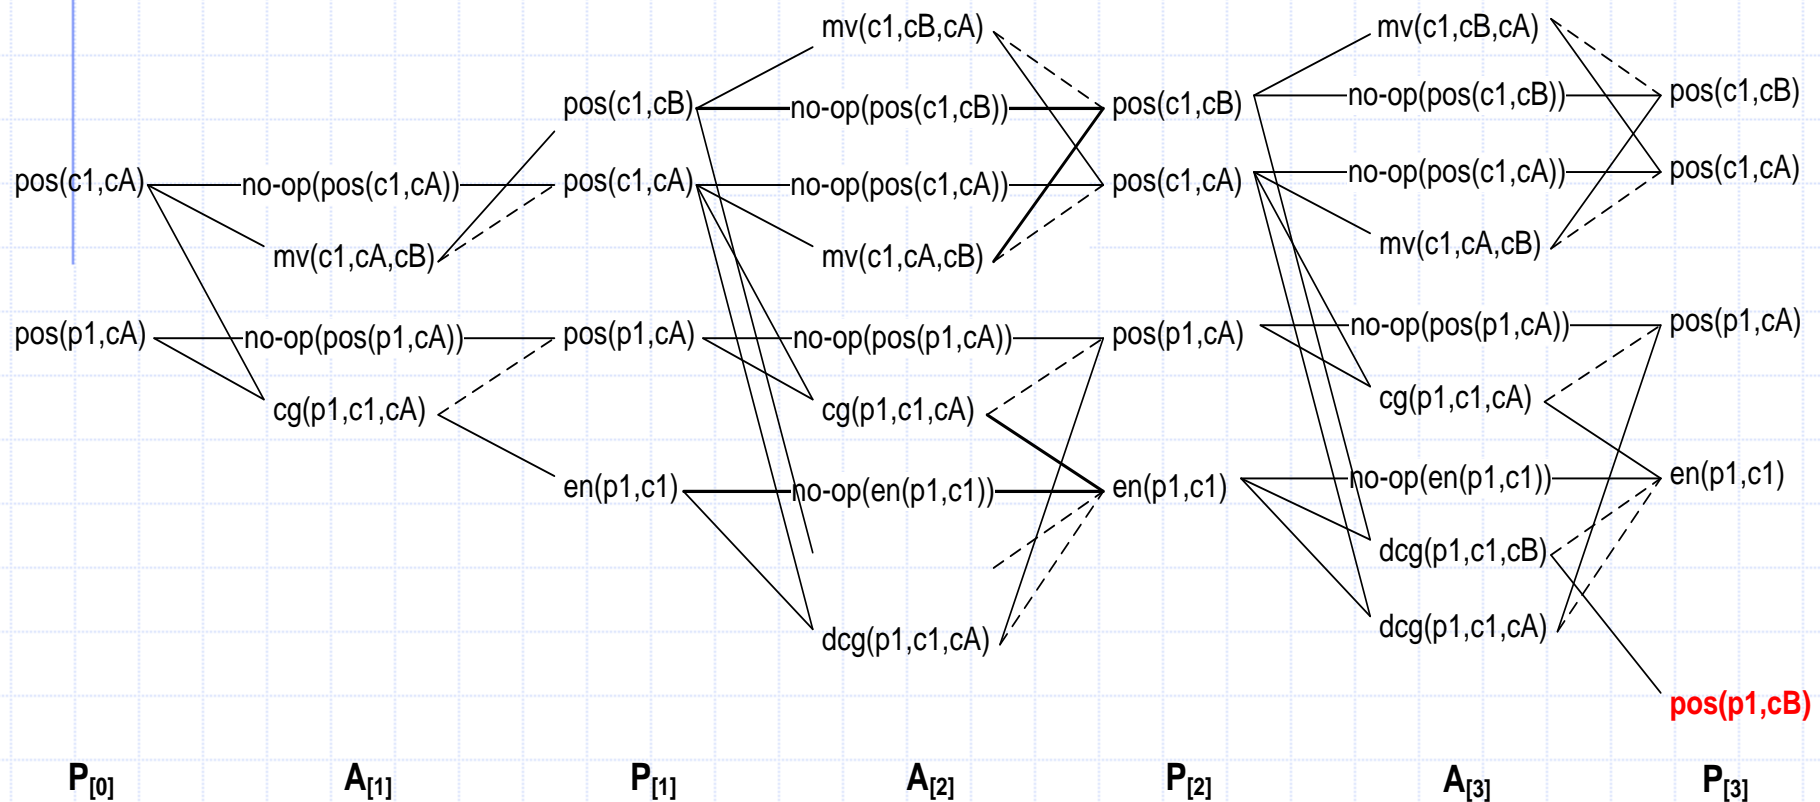
Interferences in A[2]

Competing needs in A[2]



## Example (layers A[3], P[3])

- The planning graph is extended level by level until the goals are achieved (necessary but not sufficient condition)

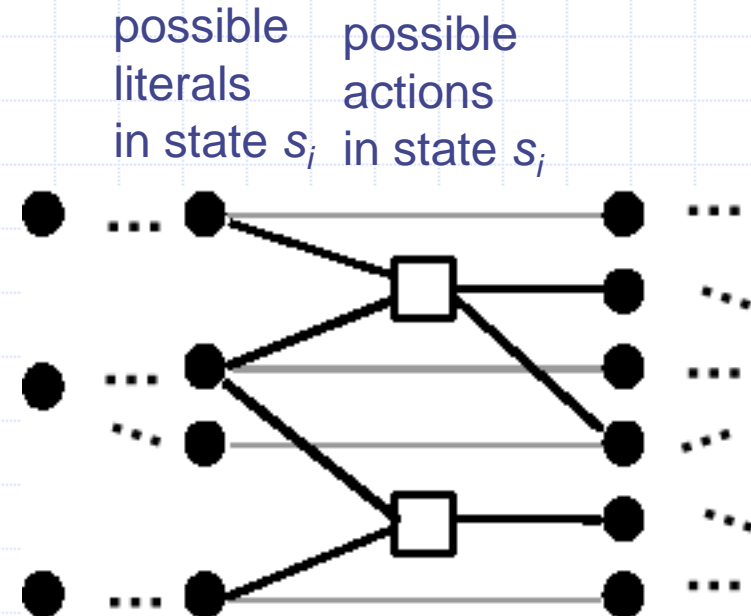


# The Graphplan algorithm

procedure Graphplan:

- for  $k = 0, 1, 2, \dots$ 
  - *Graph expansion:*
    - create a “planning graph” that contains  $k$  “levels”
  - Check whether the planning graph satisfies a necessary (but insufficient) condition for plan existence
  - If it does, then
    - do *solution extraction*:
      - backward search, modified to consider only the actions in the planning graph
      - if we find a solution, then return it

relaxed  
problem



## Solution extraction

- Check to see whether there's a possible solution
  - All of the goals appear at a proposition level
  - None are mutex with each other
- Thus, there's a chance that a plan exists

# Solution Extraction

The set of goals we are trying to achieve

The level of the state  $s_j$

procedure Solution-extraction( $g, j$ )

if  $j=0$  then return the solution

for each literal  $l$  in  $g$

nondeterministically choose an action  
to use in state  $s_{j-1}$  to achieve  $l$


if any pair of chosen actions are mutex  
then backtrack

$g' := \{\text{the preconditions of the chosen actions}\}$

Solution-extraction( $g', j-1$ )

end Solution-extraction

A real action or a maintenance action



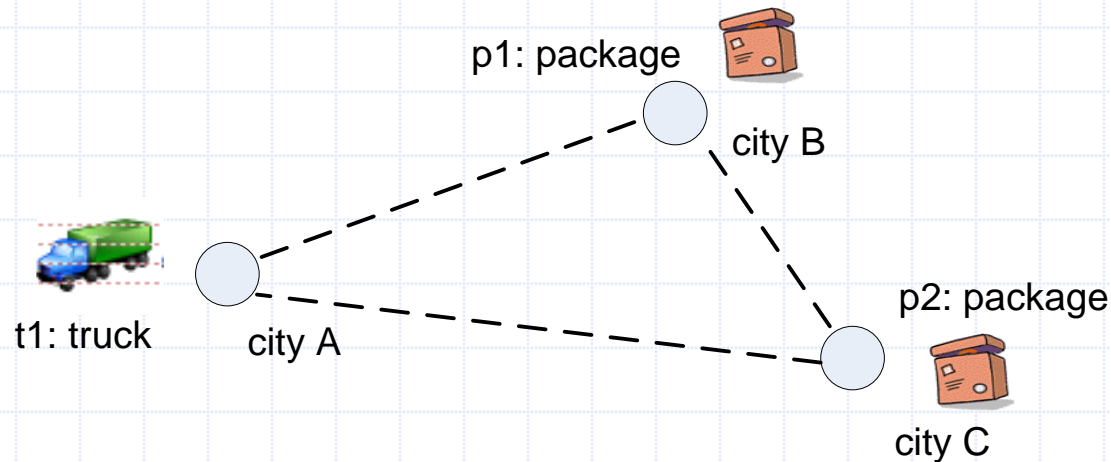
## Solution extraction

- If any pair of chosen actions are mutex then backtrack:
  - Choose another action for a literal so that the pair is not mutex
- If all sets of actions (combinations of actions) in level  $A_i$  for the literals at level  $P_i$  contain mutex actions then go back and do more graph expansion => generate another action level and another proposition level
- Termination of Graphplan:
  - Literals increase monotonically
  - Actions increase monotonically
  - Mutexes decrease monotonically
  - The graph levels off in a finite number of steps (two subsequent proposition levels are identical)



# Solution extraction and mutex relations

- The mutual exclusion rules do not guarantee to find *all* mutual exclusion relationships, but usually find a large number of them (in fact, determining *all* mutual exclusion relationships can be as hard as finding a plan).
- Mutual exclusion rules only find binary mutex but there also exists other higher-order mutex, eg. *ternary* mutex ...
- Example:



Goals: (at p1 cityA) and (at p2 cityA)


## Solution extraction and mutex relations

- Shortest solution plan involves six steps:
  - {move (t1,cityA,cityB)}
  - {load (t1,p1,cityB)}
  - {move (t1,cityB,cityC)}
  - {load (t1,p2,cityC)}
  - {move (t1,cityC,cityA)}
  - {unload (t1,p1,cityA), unload (t1,p2,cityA)}
- However, Graphplan infers that the goals can be achieved together by level 5 of the graph. This is because **(in p1 t1) (in p2 t1)** and **(at t1 cityA)** are achieved by level 4 of the graph and Graphplan will not find any mutex relation between any pair of literals:


## Solution extraction and mutex relations

- the pairs  $\{(\text{in } p1 \ t1), (\text{at } t1 \ \text{cityA})\}$  and  $\{(\text{in } p2 \ t1), (\text{at } t1 \ \text{cityA})\}$  are not mutex at level 4 because there is a 3-step plan which will get the truck loaded with either package to the destination; and when two literals are found non-mutex, they can never subsequently become mutex since if they can occur in the same state then they will always be able to appear in the same state
- the pair  $\{(\text{in } p1 \ t1), (\text{in } p2 \ t1)\}$  is also not mutex at level 5 because there is a 4-step plan which will have the truck loaded with both packages although not at the destination
- then, no *binary* mutex relation is found with the three literals by level 4 although there is a mutex among the three literals
- Consequently, Graphplan begins searching for a plan from level 5 when it is impossible for one to be found until level 6 has been constructed.

# Solution extraction and mutex relations



# Solution extraction and mutex relations



## Discussion

- Advantage:
  - The backward-search part of Graphplan—which is the hard part—will only look at the actions in the planning graph
  - smaller search space than PSP; thus faster
- Disadvantage:
  - To generate the planning graph, Graphplan creates a huge number of ground atoms
  - Many of them may be irrelevant
- Can alleviate (but not eliminate) this problem by assigning data types to the variables and constants
  - Only instantiate variables to terms of the same data type
- For classical planning, the advantage outweighs the disadvantage
  - GraphPlan solves classical planning problems much faster than PSP

# Heuristics derived from planning graphs

- Recall how GraphPlan works:

loop

*Graph expansion:*

this takes polynomial time

extend a “planning graph” forward from the initial state  
until we have achieved a necessary (but insufficient) condition for plan existence


*Solution extraction:*

this takes exponential time

search backward from the goal, looking for a correct plan  
if we find one, then return it

repeat

# Heuristics derived from planning graphs



- In the graph, there are alternating layers of ground literals and actions
- The number of “action” layers is a lower bound on the number of actions in the plan
- Construct a planning graph, starting at  $s$
- $h(g_i)$ : estimate to achieve  $g_i$  from  $s$
- $h(g_i)$  = level of the first layer that “possibly achieves”  $g_i$



# Heuristics derived from planning graphs

- Heuristics for a set of goal  $\mathbf{G} = \{\text{pos}(p1, cB), \text{pos}(c1, cA)\}$ :
  - The **max level heuristic** takes the maximum level cost of any of the goals (admissible, not very accurate).  
$$h_{\max}(G) = \max_{g \in G} h(g) \quad // \quad h_{\max}(G) = \max(0, 3) = 3$$
  - The sum **level heuristic** returns the sum of the level costs of the goals (inadmissible, works well in practice)  
$$h_{\text{sum}}(G) = \sum_{g \in G} h(g) \quad // \quad h_{\text{sum}}(G) = 0 + 3 = 3$$
  - The **max<sub>2</sub> level heuristic** takes the maximum level at which coexist any pair of goals.  
$$h_{\max_2}(G) = \max_{\{g1, g2\} \in G} h(g1 \wedge g2) \quad // \quad h_{\max_2}(G) = \max(4) = 4$$
  - The **max<sub>k</sub> level heuristic** takes the maximum level at which coexist any  $k$  goals. 
$$h_{\max_k}(G) = \max_{\{g1, g2, \dots, gk\} \in G} h(g1 \wedge g2 \wedge \dots \wedge gk)$$


# Heuristics derived from planning graphs

- Heuristics used by forward state-space planners (like FF)
- Application of the heuristic over each state in the tree
- Planning graph + mutex calculation over each state => very costly process
- Computing heuristics on a **relaxed planning graph (RPG)**:
  - Ignoring negated effects
  - No mutex calculation
  - Slide 11 shows an example on how the RPG will look like
- $\text{hmax}(G) = \max(0, 2) = 2$
- $\text{hsum}(G) = 0 + 2 = 2$
- $\text{hmax}_2(G) = \max(2) = 2$

# Heuristics derived from planning graphs

- **Relaxed Planning Graph:**
  - No delete effects
  - No mutex
  - No no-op actions because we are only interested in the first appearance of any action/proposition
- Relaxed plan heuristic: extract a plan from a **RPG**
- See examples of the blocks-world domain

# The FastForward Planner



- Use a heuristic function similar to  $h(s) = \Delta^g(s, g)$ 
  - Some ways to improve it (I'll skip the details)
- Don't want an A\*-style search (takes too much memory)
- Instead, use a greedy procedure:

until we have a solution, do

expand the current state  $s$

$s :=$  the child of  $s$  for which  $h(s)$  is smallest

(i.e., the child we think is closest to a solution)

# The FastForward Planner


- Use a heuristic function  $h(s) = \text{relaxed plan}$
- Don't want an A\*-style search (takes too much memory)
- Instead, use a greedy procedure:

until we have a solution, do

expand the current state  $s$

$s := \text{the child of } s \text{ for which } h(s) \text{ is smallest}$

(i.e., the child we think is closest to a solution)



# The FastForward Planner


- Use a heuristic function  $h(s) = \text{relaxed plan}$
- Don't want an A\*-style search (takes too much memory)
- Instead, use a greedy procedure:

until we have a solution, do

expand the current state  $s$

$s := \text{the child of } s \text{ for which } h(s) \text{ is smallest}$

(i.e., the child we think is closest to a solution)



# The FastForward Planner


- Use a heuristic function  $h(s) = \text{relaxed plan}$
- Don't want an A\*-style search (takes too much memory)
- Instead, use a greedy procedure:

until we have a solution, do

expand the current state  $s$

$s := \text{the child of } s \text{ for which } h(s) \text{ is smallest}$

(i.e., the child we think is closest to a solution)



# The FastForward Planner

- Use a heuristic function similar to  $h(s) = \text{relaxed plan}$
- Don't want an A\*-style search (takes too much memory)
- Instead, use a greedy procedure:

until we have a solution, do

expand the current state  $s$

$s :=$  the child of  $s$  for which  $h(s)$  is smallest

(i.e., the child we think is closest to a solution)

- There are some ways FF improves on this
  - e.g. a way to escape from local minima
    - breadth-first search, stopping when a node with lower cost is found
- Can't guarantee how fast it will find a solution, or how good a solution it will find
  - However, it works pretty well on many problems

