

Chapter 2. Multilayer Perceptron

Neural Networks

2023/2024

Máster Universitario en Inteligencia Artificial, Reconocimiento
de Formas e Imagen Digital

Departamento de Sistemas Informáticos y Computación

Index

- 1 Introduction ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 Bibliography ▷ 66

Index

- 1 *Introduction* ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 Bibliography ▷ 66

Introduction

- From Perceptron to Multi-Layer Perceptron (MLP)
- From linear separability to complex decision boundaries (generalized linear discriminant functions)
 - Add layers to the original Perceptron architecture
 - Activation units to overcome the linear limitations
- Back-propagation:
 - Non-convex optimization
 - Initialization
 - Local-minima

Linear Discriminant Functions

- Graphical representation of a LDF.
Given an input $\mathbf{x} \in \mathbb{R}^D$

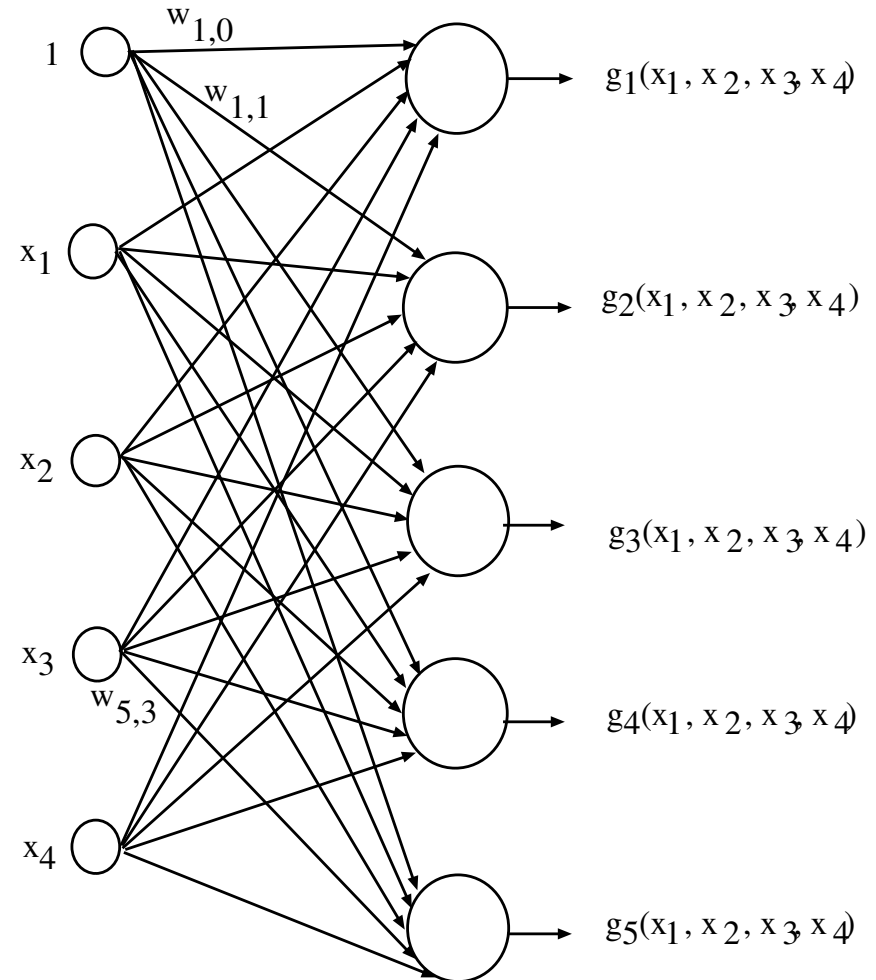
$$\hat{c} = G(\mathbf{x}) \equiv \operatorname{argmax}_{1 \leq c \leq C} g_c(\mathbf{x})$$

$$g_c(\mathbf{x}) = \mathbf{w}_c^t \mathbf{x} + w_{c0}$$

$$\mathbf{g}(\mathbf{x}) = \mathbf{W}\mathbf{x}$$

\mathbf{W} is a matrix of dimension $C \times (D + 1)$ (including the thresholds w_{c0}) and \mathbf{x} is of dimension $D + 1$ ($x_0 = 1$)

- $E = \mathbb{R}^4$,
- $\mathbf{w}_c \in \mathbb{R}^4$, $w_{c0} \in \mathbb{R}$ for $1 \leq c \leq 5$
- Classes = $\{1, 2, 3, 4, 5\}$



Generalized Linear Discriminant Functions

- Graphical representation of a GLDF.
Given an input $\mathbf{x} \in \mathbb{R}^D$ and a function $\Phi : \mathbb{R}^D \rightarrow \mathbb{R}^{D'}$

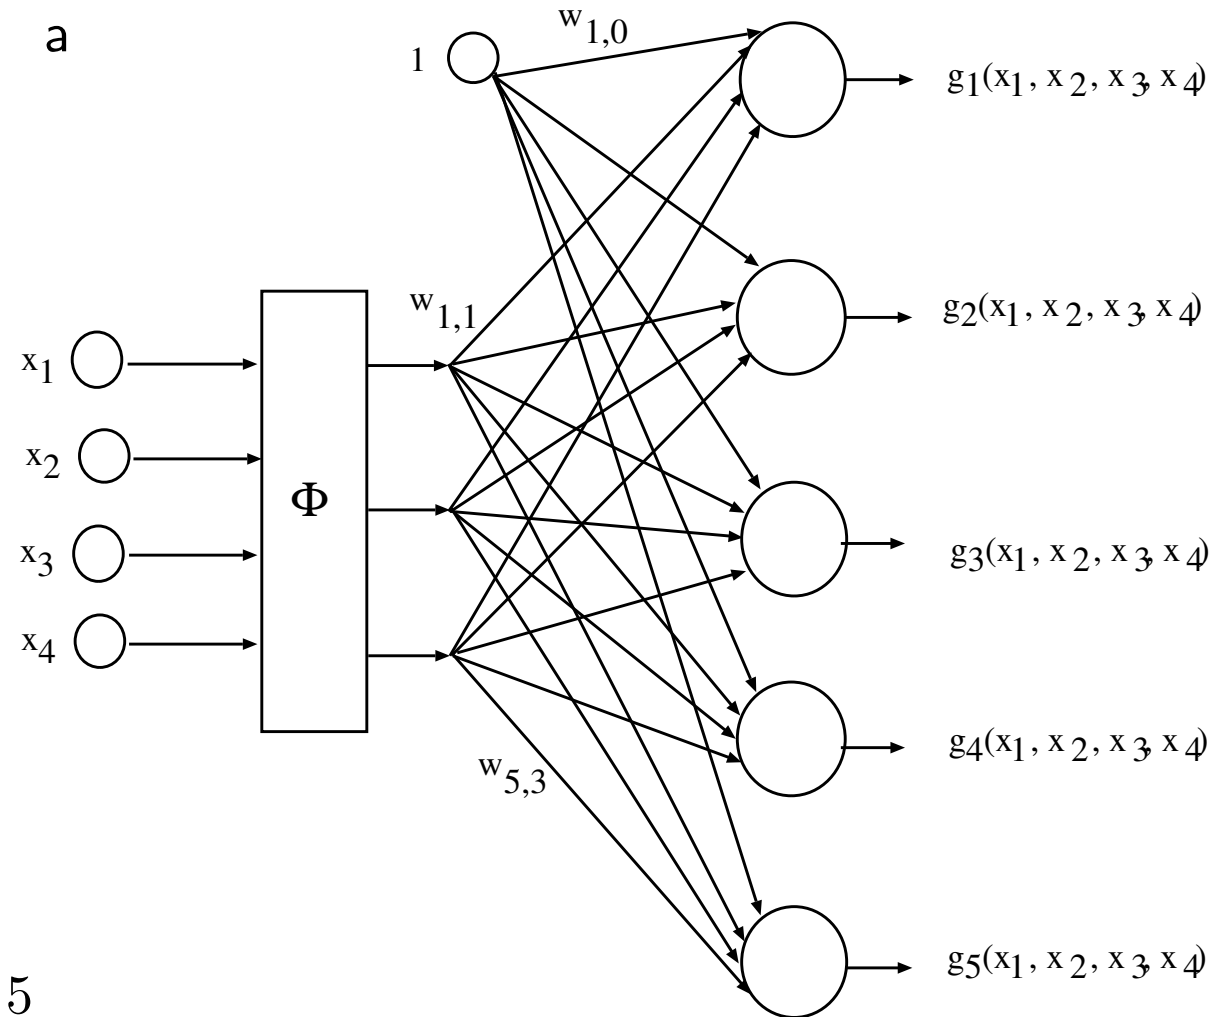
$$\hat{c} = G(\mathbf{x}) \equiv \operatorname{argmax}_{1 \leq c \leq C} g_c(\mathbf{x})$$

$$g_c(\mathbf{x}) = \mathbf{w}_c^t \Phi(\mathbf{x}) + w_{c0}$$

$$\mathbf{g}(\mathbf{x}) = \mathbf{W}\Phi(\mathbf{x})$$

In this case, \mathbf{W} is a matrix of dimension $C \times (D' + 1)$.

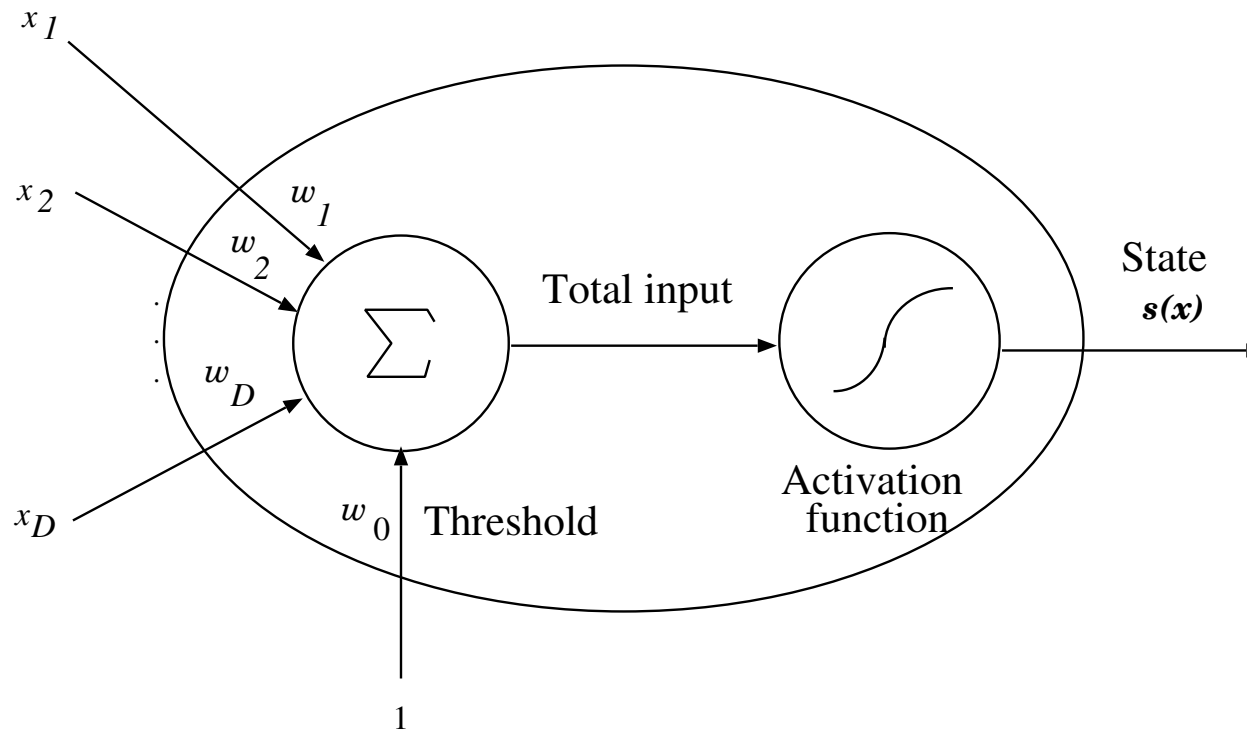
- $E = \mathbb{R}^4$,
- $\Phi : E \rightarrow \mathbb{R}^3$
- Classes = $\{1, 2, 3, 4, 5\}$
- $\mathbf{w}_c \in \mathbb{R}^3$, $w_{c0} \in \mathbb{R}$ for $1 \leq c \leq 5$



Index

- 1 Introduction ▷ 3
- 2 *Multilayer Perceptron* ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 Bibliography ▷ 66

Logistic linear discriminant functions



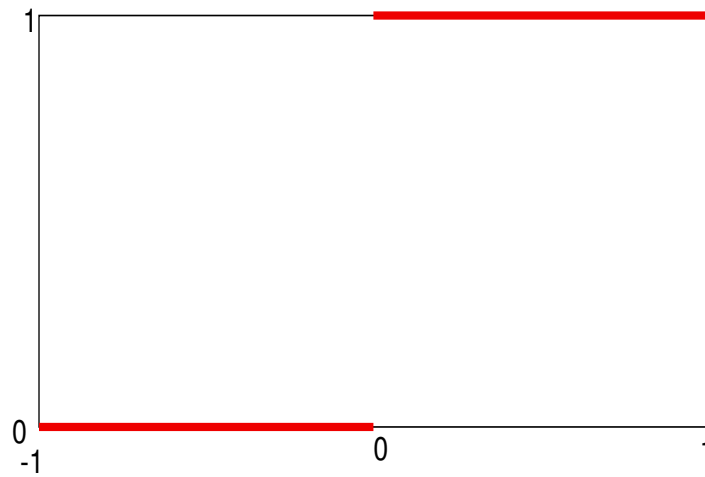
Compact notation: from $\mathbf{x} = (x_1, \dots, x_D)$ to $\mathbf{x} = (1, x_1, \dots, x_D)$, and from $\mathbf{w} = (w_1, \dots, w_D)$ to $\mathbf{w} = (w_0, w_1, \dots, w_D)$:

$$s(\mathbf{x}) = f\left(\sum_{k=1}^D w_k x_k + w_0\right) = f(\mathbf{w}^t \mathbf{x}), \text{ where } f \text{ is an activation function.}$$

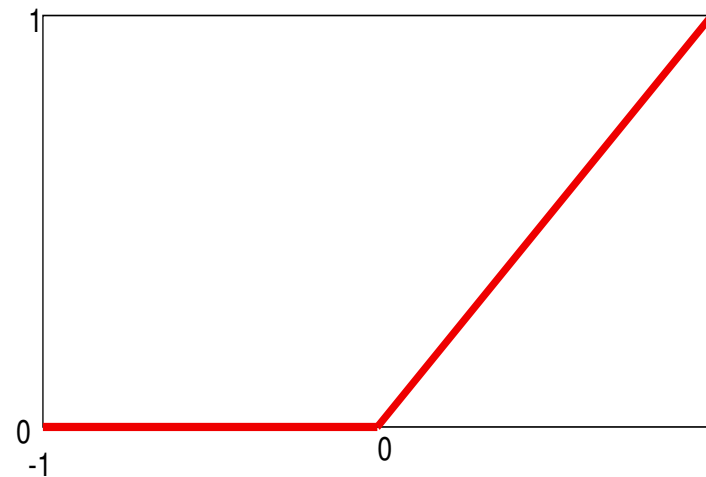
Activation functions

- *Linear*: $f_L(z) = z, z \in \mathbb{R}$
- *Step*: $f_E(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z < 0 \end{cases}, z \in \mathbb{R}$
- *ReLU* (rectified linear unit): $f_R(z) = \max(0, z), z \in \mathbb{R}$
- *PReLU* (parametric rectified linear unit): $f_{PR}(z) = \begin{cases} z & \text{if } z > 0 \\ a z & \text{if } z \leq 0 \end{cases}, z \in \mathbb{R}$
- *Sigmoid*: $f_S(z) = \frac{1}{1 + \exp(-z)}, z \in \mathbb{R}$
- *Hyperbolic tangent*: $f_T(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}, z \in \mathbb{R} \quad (f_T(z) = 2f_S(2z) - 1)$
- *Softmax*: $f_{SM}(z_j) = \frac{\exp(z_j)}{\sum_{j'=1}^M \exp(z_{j'})}; (f_{SM}(z_j) = f_S(z_j - \ln(\sum_{j' \neq j} \exp(z_{j'}))))$, $z_1, \dots, z_M \in \mathbb{R}$
- Others: Exponential Linear Units (ELU), Maxout, ...

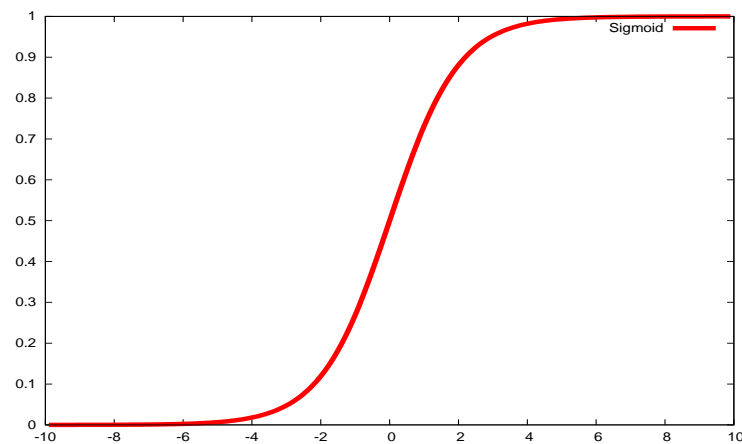
Activation functions



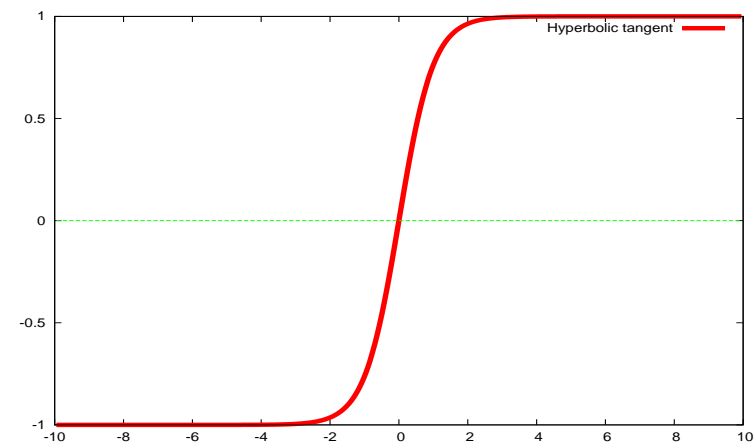
Step



ReLU

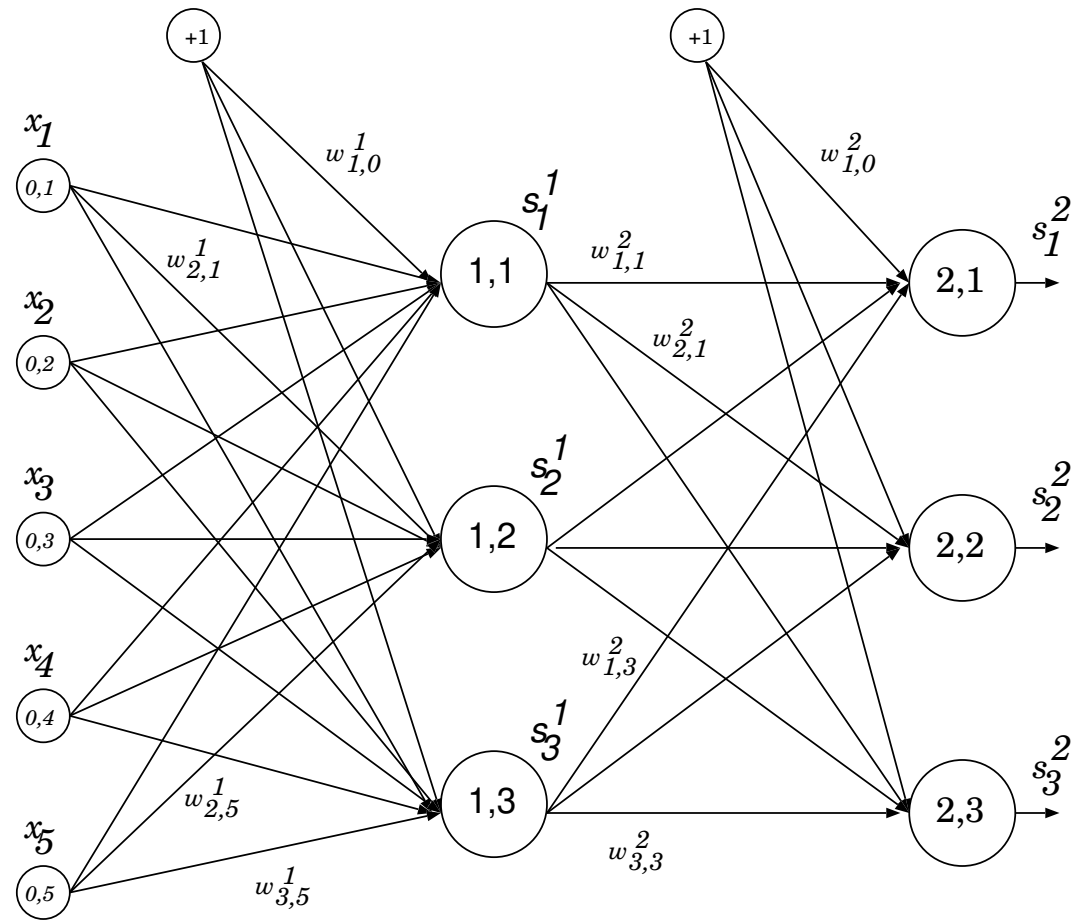


Sigmoid



Hyperbolic tangent

A two-layer perceptron



Hidden layer

Output layer

$$s_j^1 = f(\sum_i w_{j,i}^1 x_i), 1 \leq j \leq M_1$$

$$s_j^2 = f(\sum_i w_{j,i}^2 s_i^1), 1 \leq j \leq M_2$$

A two-layer perceptron

A two-layer perceptron consists of a combination of logistic linear discriminant functions grouped in layers and defines a set of M_2 discriminant functions:

$$g_k(\mathbf{x}; \theta) \equiv s_k^2(\mathbf{x}) = f\left(\sum_{j=0}^{M_1} w_{k,j}^2 s_j^1(\mathbf{x})\right) = f\left(\sum_{j=0}^{M_1} w_{k,j}^2 f\left(\sum_{j'=0}^{M_0} w_{j,j'}^1 x_{j'}\right)\right)$$

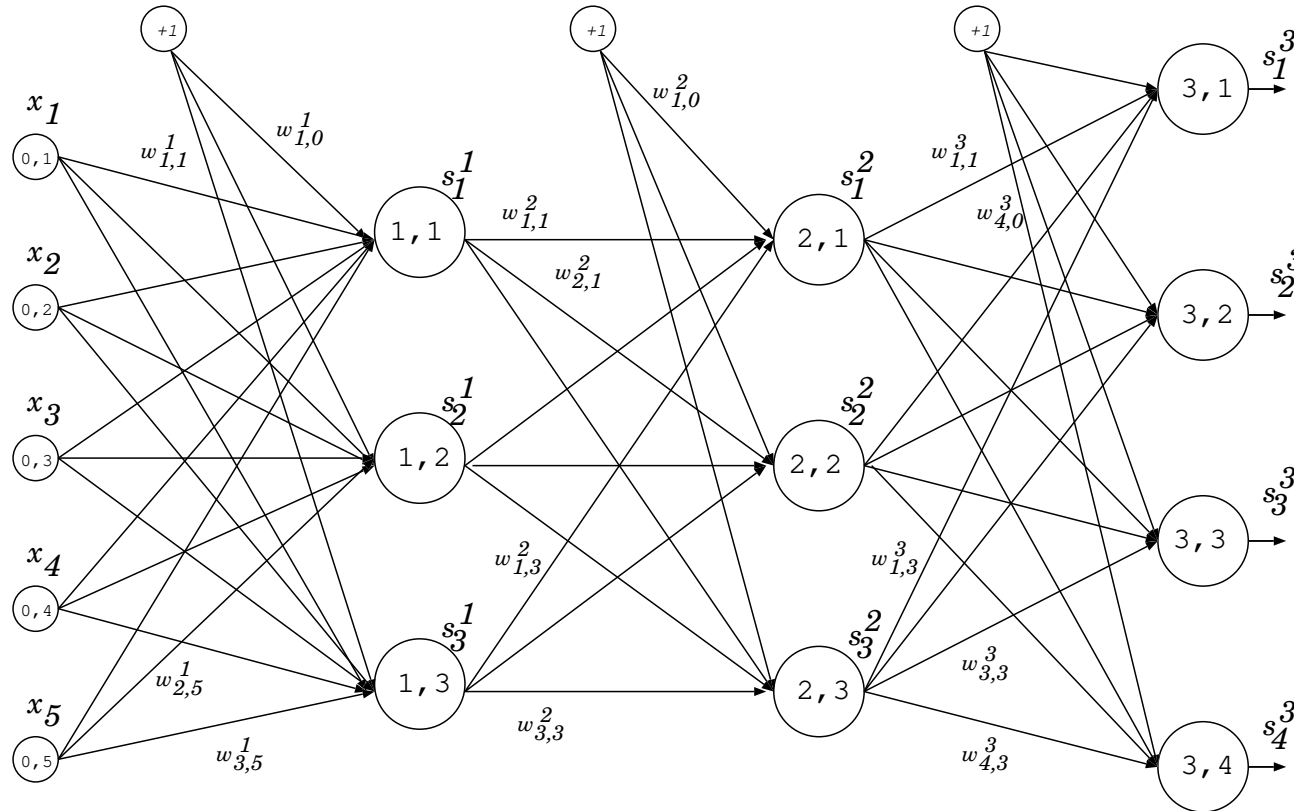
for $1 \leq k \leq M_2$, $M_0 \equiv D$. Therefore,

$$\theta \equiv \mathbf{w} = (w_{1,0}^1, \dots, w_{M_1,M_0}^1, w_{1,0}^2, \dots, w_{M_2,M_1}^2)$$

In compact notation: $\mathbf{g}(\mathbf{x}; \theta) \equiv \mathbf{s}^2(\mathbf{x}) = \mathbf{f}(\mathbf{W}^2 \mathbf{f}(\mathbf{W}^1 \mathbf{x}))$

Regression problem: Let A be a training set $\{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^{M_0}$, $\mathbf{t}_n \in \mathbb{R}^{M_2}$ search for \mathbf{w} such that $\mathbf{s}^2(\mathbf{x}_n) = \mathbf{t}_n$ with $1 \leq n \leq N$.

A three-layer perceptron



Two hidden layers

An output layer

$$1 \leq j \leq M_1 \quad 1 \leq j \leq M_2 \quad 1 \leq j \leq M_3$$

$$s_j^1 = f(\sum_i w_{j,i}^1 x_i) \quad s_j^2 = f(\sum_i w_{j,i}^2 s_i^1) \quad s_j^3 = f(\sum_i w_{j,i}^3 s_i^2)$$

A three-layer perceptron

A **three-layer perceptron** defines a set of M_3 discriminant functions:

$$\begin{aligned} g_k(\mathbf{x}; \theta) &\equiv s_k^3(\mathbf{x}) = f\left(\sum_{j=0}^{M_2} w_{k,j}^3 s_j^2(\mathbf{x})\right) \\ &= f\left(\sum_{j=0}^{M_2} w_{k,j}^3 f\left(\sum_{j'=0}^{M_1} w_{j,j'}^2 s_{j'}^1(\mathbf{x})\right)\right) = f\left(\sum_{j=0}^{M_2} w_{k,j}^3 f\left(\sum_{j'=0}^{M_1} w_{j,j'}^2 f\left(\sum_{j''=0}^{M_0} w_{j',j''}^1 x_{j''}\right)\right)\right) \end{aligned}$$

for $1 \leq k \leq M_3$, $M_0 \equiv D$. Therefore,

$$\theta \equiv \mathbf{w} = (w_{10}^1, \dots, w_{M_1, M_0}^1, w_{1,0}^2, \dots, w_{M_2, M_1}^2, w_{1,0}^3, \dots, w_{M_3, M_2}^3)$$

In compact notation: $\mathbf{g}(\mathbf{x}; \theta) \equiv \mathbf{s}^3(\mathbf{x}) = \mathbf{f}(\mathbf{W}^3 \mathbf{f}(\mathbf{W}^2 \mathbf{f}(\mathbf{W}^1 \mathbf{x})))$

Regression problem: Let A be a training set $\{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^{M_0}$, $\mathbf{t}_n \in \mathbb{R}^{M_3}$: search for \mathbf{w} such that $\mathbf{s}^3(\mathbf{x}_n) = \mathbf{t}_n$ with $1 \leq n \leq N$.

Multilayer perceptrons and activation functions

- Given a two-layer perceptron,

$$g_k(\mathbf{x}) \equiv s_k^2(\mathbf{x}) = f\left(\sum_{j=0}^{M_1} w_{k,j}^2 s_j^1(\mathbf{x})\right) = f\left(\sum_{j=0}^{M_1} w_{k,j}^2 f\left(\sum_{j'=0}^{M_0} w_{j,j'}^1 x_{j'}\right)\right)$$

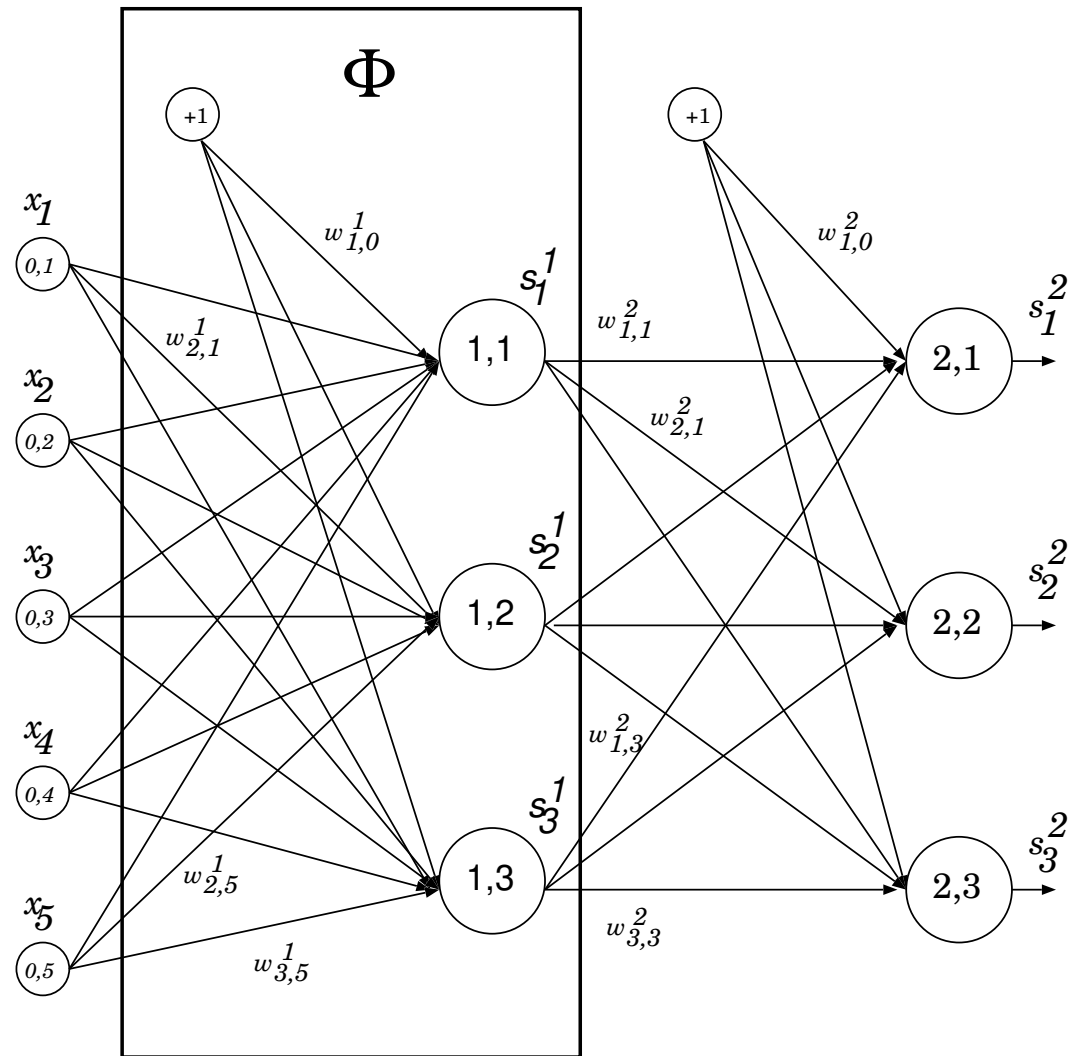
- If all the activation functions are linear, a multilayer perceptron defines **a linear discriminant function**:

$$g_k(\mathbf{x}) \equiv s_k^2(\mathbf{x}) = \sum_{j'=0}^{M_0} \left(\sum_{j=0}^{M_1} w_{k,j}^2 w_{j,j'}^1 \right) x_{j'} = \sum_{j'=0}^{M_0} w_{k,j'} x_{j'}$$

- If at least one activation function is not linear (and all activation functions in the output layer are linear), a multilayer perceptron defines a **generalized linear discriminant function**:

$$g_k(\mathbf{x}) \equiv s_k^2(\mathbf{x}) = \sum_{j=0}^{M_1} w_{k,j}^2 f\left(\sum_{j'=0}^{M_0} w_{j,j'}^1 x_{j'}\right) = \sum_{j=0}^{M_1} w_{k,j}^2 \Phi_j(\mathbf{x})$$

A two-layer perceptron



The multilayer perceptron as a classifier

For a problem of C classes, the multilayer perceptron has C units and the training target vectors \mathbf{t}_n have the form for $1 \leq n \leq N$:

$$t_{nc} = \begin{cases} 1 (+1) & \text{if } \mathbf{x}_n \text{ is of the class } c \\ 0 (-1) & \text{otherwise} \end{cases}$$

The classifier is:

$$G(\mathbf{x}) = \operatorname{argmax}_{1 \leq k \leq C} g_k(\mathbf{x}; \theta) \equiv \operatorname{argmax}_{1 \leq k \leq M_2 \equiv C} s_k^2(\mathbf{x})$$

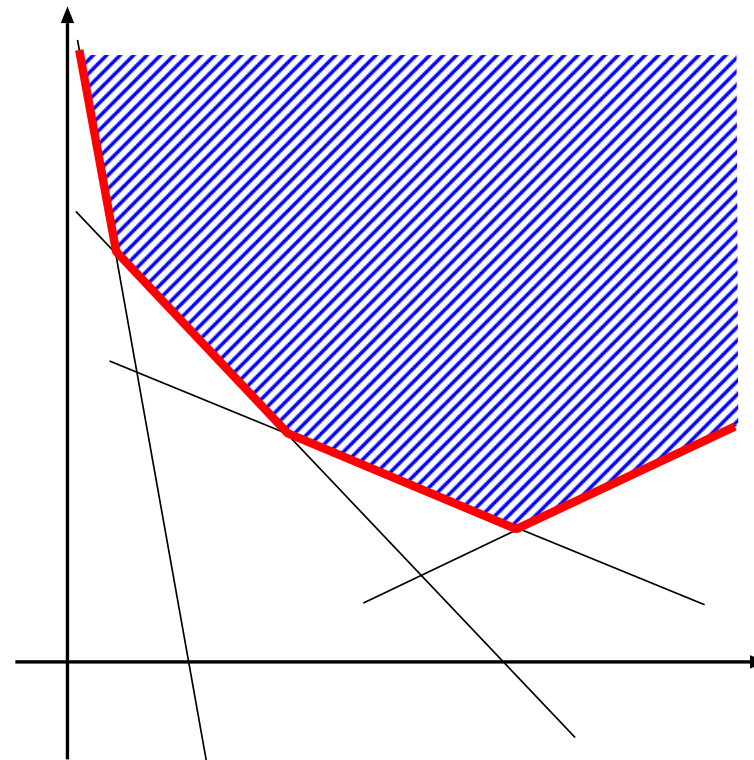
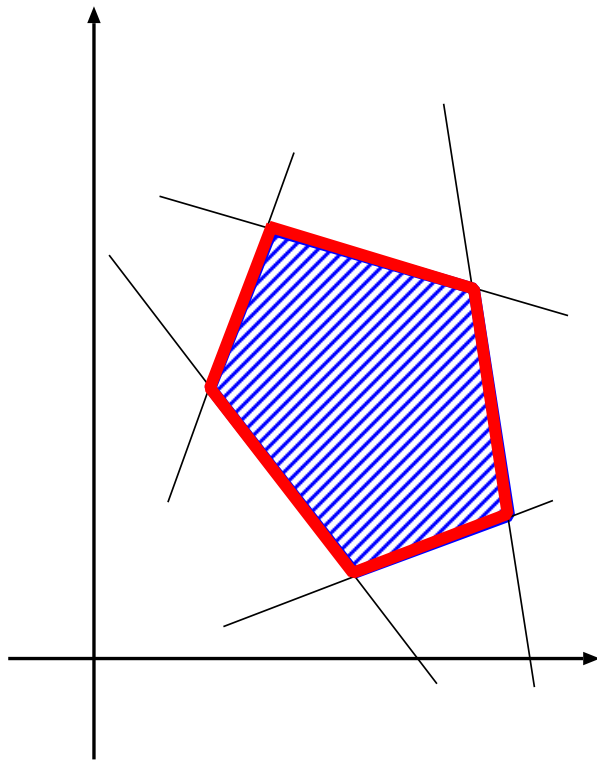
Softmax activation function is used in the output layer.

Given a training sample with N patterns, is there a multilayer perceptron that classifies correctly the training sample?

- If the training sample is linearly separable: a perceptron without hidden layers.
- A multilayer perceptron of 1 hidden layer with $N - 1$ nodes and step activation functions can classify correctly the sample.
- About generalization?

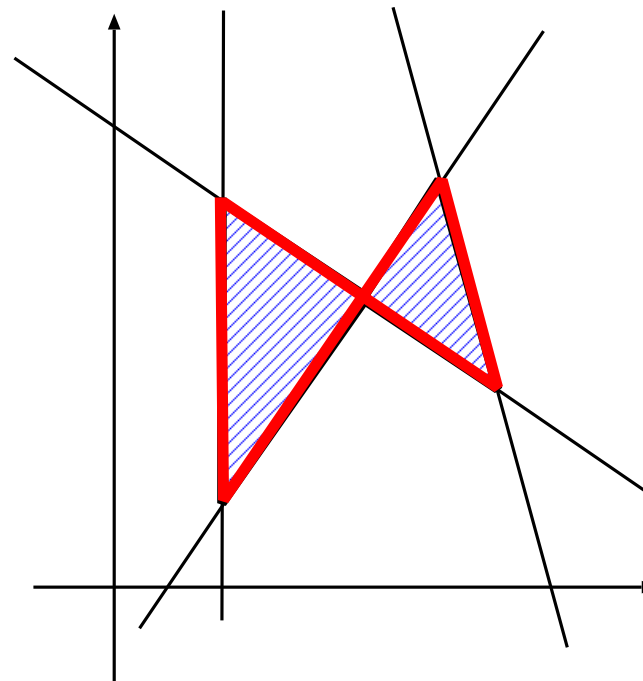
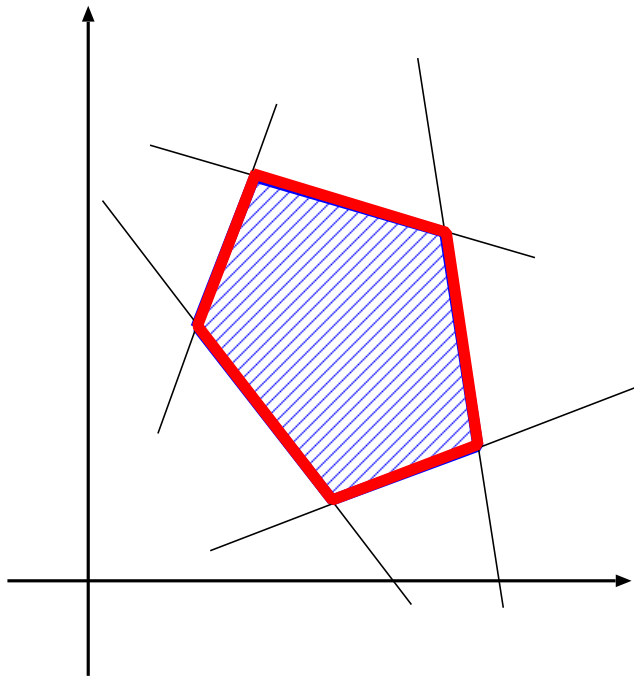
Properties of the multilayer perceptron

- A multilayer perceptron can implement decision borders that are linear at intervals.



Properties of the multilayer perceptron

- A multilayer perceptron with a hidden layer and step activation functions can implement convex decision borders.
- Any decision border based on hyperplanes can be implemented by means of a multilayer perceptron with two hidden layers and step activation functions.



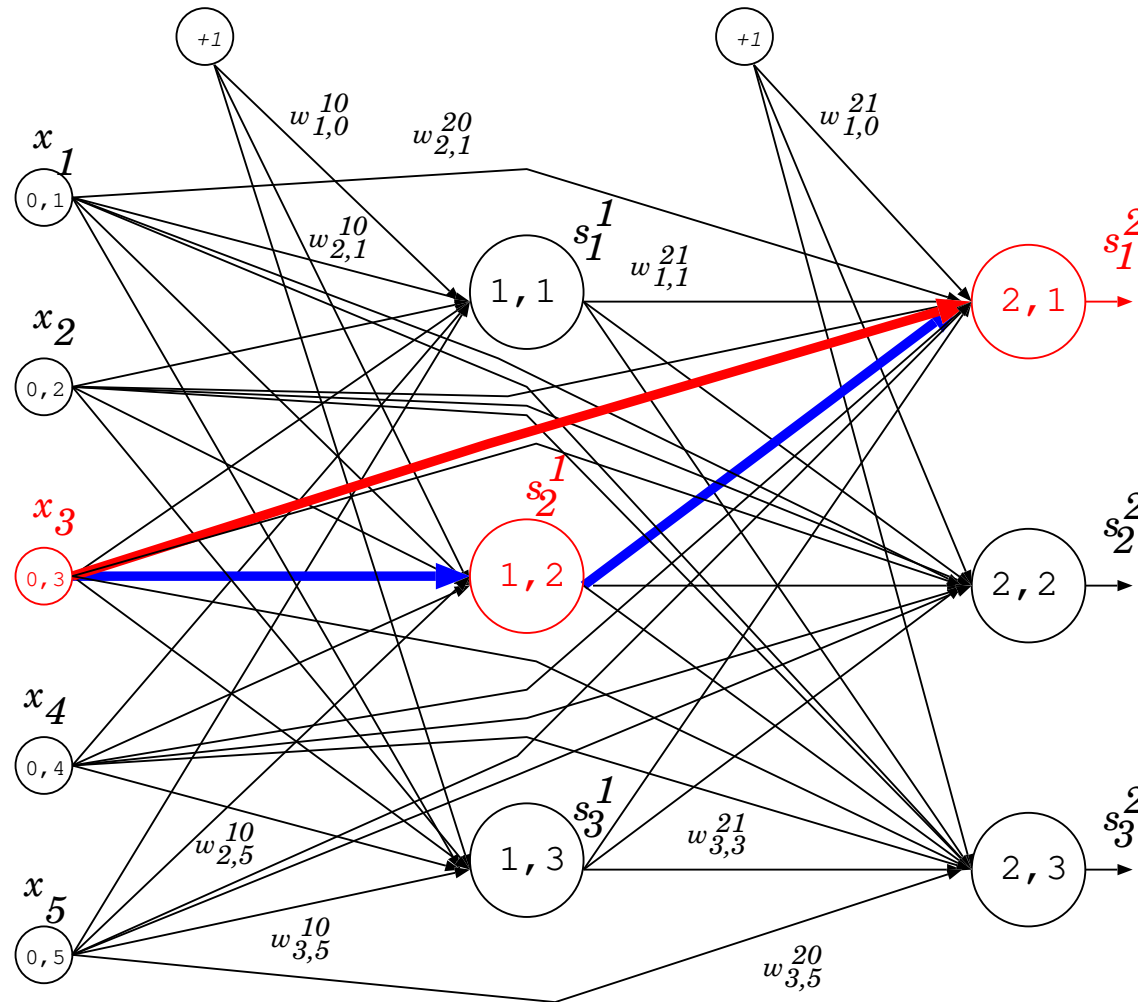
Regression with a multilayer perceptron (function approximation)

$$F(\mathbf{x}) : \mathbb{R}^{M_0} \rightarrow \mathbb{R}^{M_2} : s_k^2(\mathbf{x}) = \sum_{j=0}^{M_1} w_{k,j}^2 f\left(\sum_{j'=0}^{M_0} w_{j,j'}^1 x_{j'}\right) \text{ for } 1 \leq k \leq M_2$$

Linear activation function is usually adopted in the output layer.

- Any function can be arbitrary approached by means of a multilayer perceptron of two hidden layers and step activation functions and therefore with sigmoid functions too.
- Any function can be arbitrary approached by means of a multilayer perceptron of a hidden layer and step activation functions and therefore with sigmoid functions if the number of hidden nodes is high enough.

Feed-forward networks



$$s^1(\mathbf{x}) = \mathbf{f}(\mathbf{W}^{1,0}\mathbf{x})$$

$$s^2(\mathbf{x}) = \mathbf{f}(\mathbf{W}^{2,1}s^1(\mathbf{x}) + \mathbf{W}^{2,0}\mathbf{x})$$

Index

- 1 Introduction ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 *Error Back-propagation algorithm* ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 Bibliography ▷ 66

Error backpropagation algorithm for the multilayer perceptron

REGRESSION PROBLEM: Given a topology of a multilayer perceptron and $A = \{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^{N_0}$, $\mathbf{t}_n \in \mathbb{R}^{N_2}$, search for \mathbf{w} such that minimizes the objective function (mean squared error):

$$E_A(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N \sum_{k=1}^{M_2} (t_{n,k} - s_k^2(\mathbf{x}_n; \mathbf{w}))^2$$

that is,

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} E_A(\mathbf{w})$$

SOLUTION, gradient descent ($1 \leq k \leq 2, 1 \leq i \leq M_k, 0 \leq j \leq M_{k-1}$):

$$\Delta w_{i,j}^k = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{i,j}^k} \quad (\Delta \mathbf{W}^k = -\rho \nabla_{\mathbf{W}^k} E_A(\mathbf{w}))$$

Gradient descent

$$\begin{aligned} \mathbf{w}(1) &= \text{arbitrary} \\ \mathbf{w}(k+1) &= \mathbf{w}(k) - \rho_k \nabla J|_{\mathbf{w}=\mathbf{w}(k)} \end{aligned}$$

Where $\rho_k \in \mathbb{R}^{>0}$ is a *learning rate* and $\nabla J|_{\mathbf{w}=\mathbf{w}(k)} \equiv \left(\frac{\partial J}{\partial w_0}, \frac{\partial J}{\partial w_1}, \dots, \frac{\partial J}{\partial w_D} \right) \Big|_{\mathbf{w}=\mathbf{w}(k)}$

1-dimension

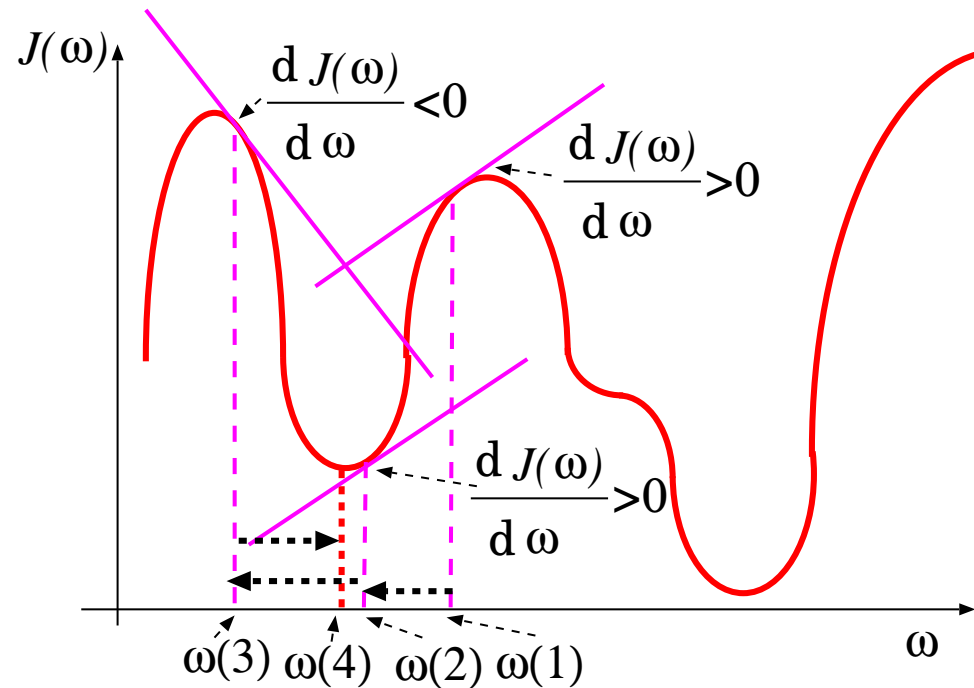
$$w(1) = \text{arbitrary}$$

$$w(2) = w(1) - \rho \frac{dJ}{dw} \Big|_{w(1)}$$

$$w(3) = w(2) - \rho \frac{dJ}{dw} \Big|_{w(2)}$$

$$w(4) = w(3) - \rho \frac{dJ}{dw} \Big|_{w(3)}$$

$$\frac{dJ}{dw} \Big|_{w(4)} = 0$$



Derivation of backpropagation algorithm (1)

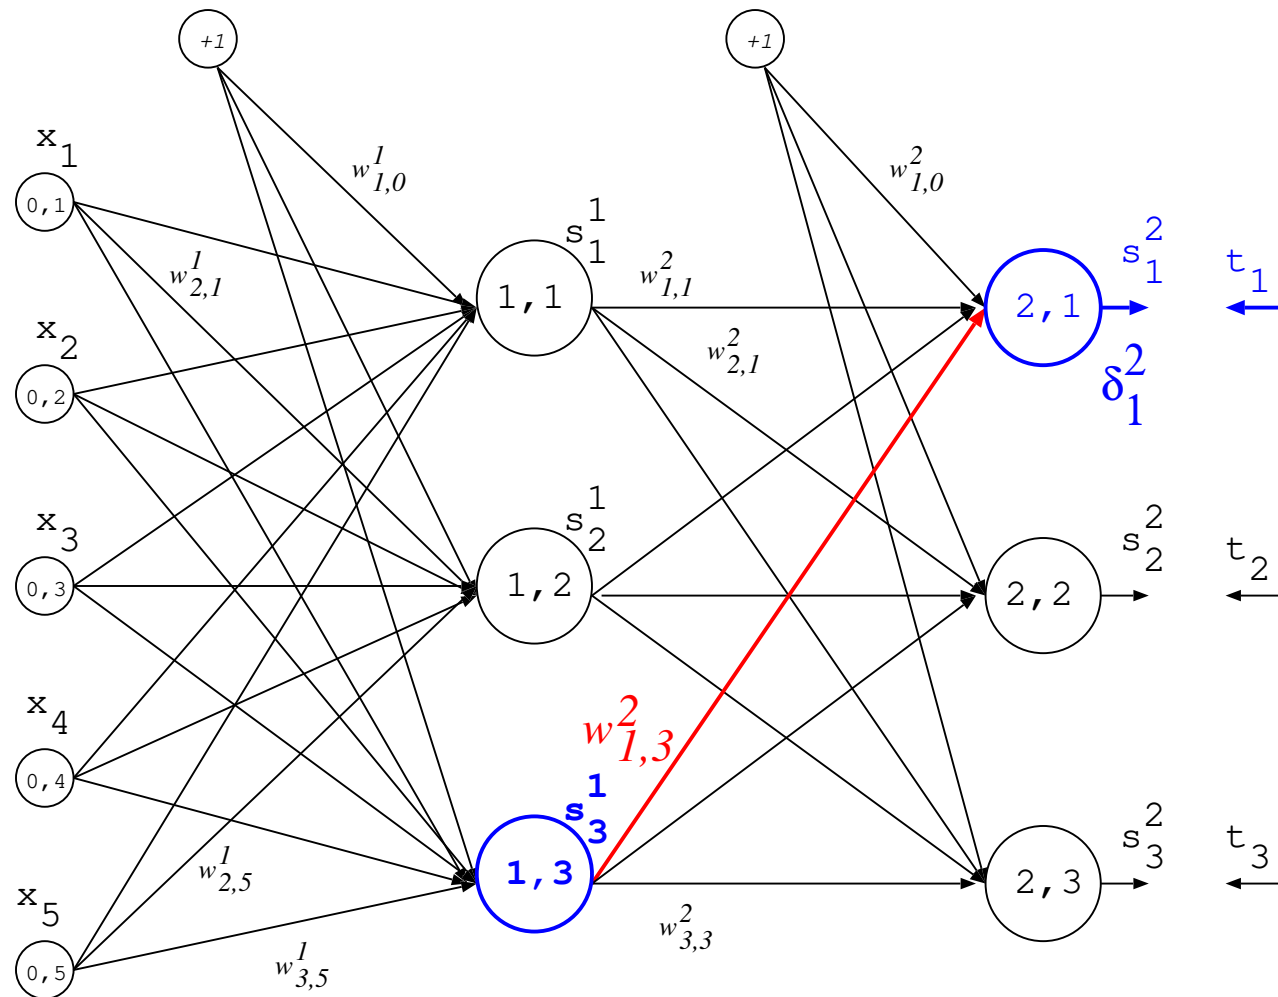
- Update weight of the output layer $w_{i,j}^2$ (if $N = 1$)

$$E_A(\mathbf{w}) = \frac{1}{2} \sum_{m=1}^{M_2} (t_m - s_m^2)^2; \quad s_m^2 = f(z_m^2); \quad z_m^2 = \left(\sum_{l=1}^{M_1} w_{m,l}^2 s_l^1 \right)$$

$$\frac{\partial E_A}{\partial w_{i,j}^2} = \frac{\partial E_A}{\partial s_i^2} \frac{\partial s_i^2}{\partial z_i^2} \frac{\partial z_i^2}{\partial w_{i,j}^2} = (-1)(t_i - s_i^2) f'(z_i^2) s_j^1 = -\delta_i^2 s_j^1$$

$$\Delta w_{i,j}^2 = \rho \delta_i^2 s_j^1$$

Derivation of backpropagation algorithm (2)



$$\Delta w_{1,3}^2 = \rho \delta_1^2 s_3^1 = \rho (t_1 - s_1^2) f'(z_1^2) s_3^1$$

Derivation of backpropagation algorithm (3)

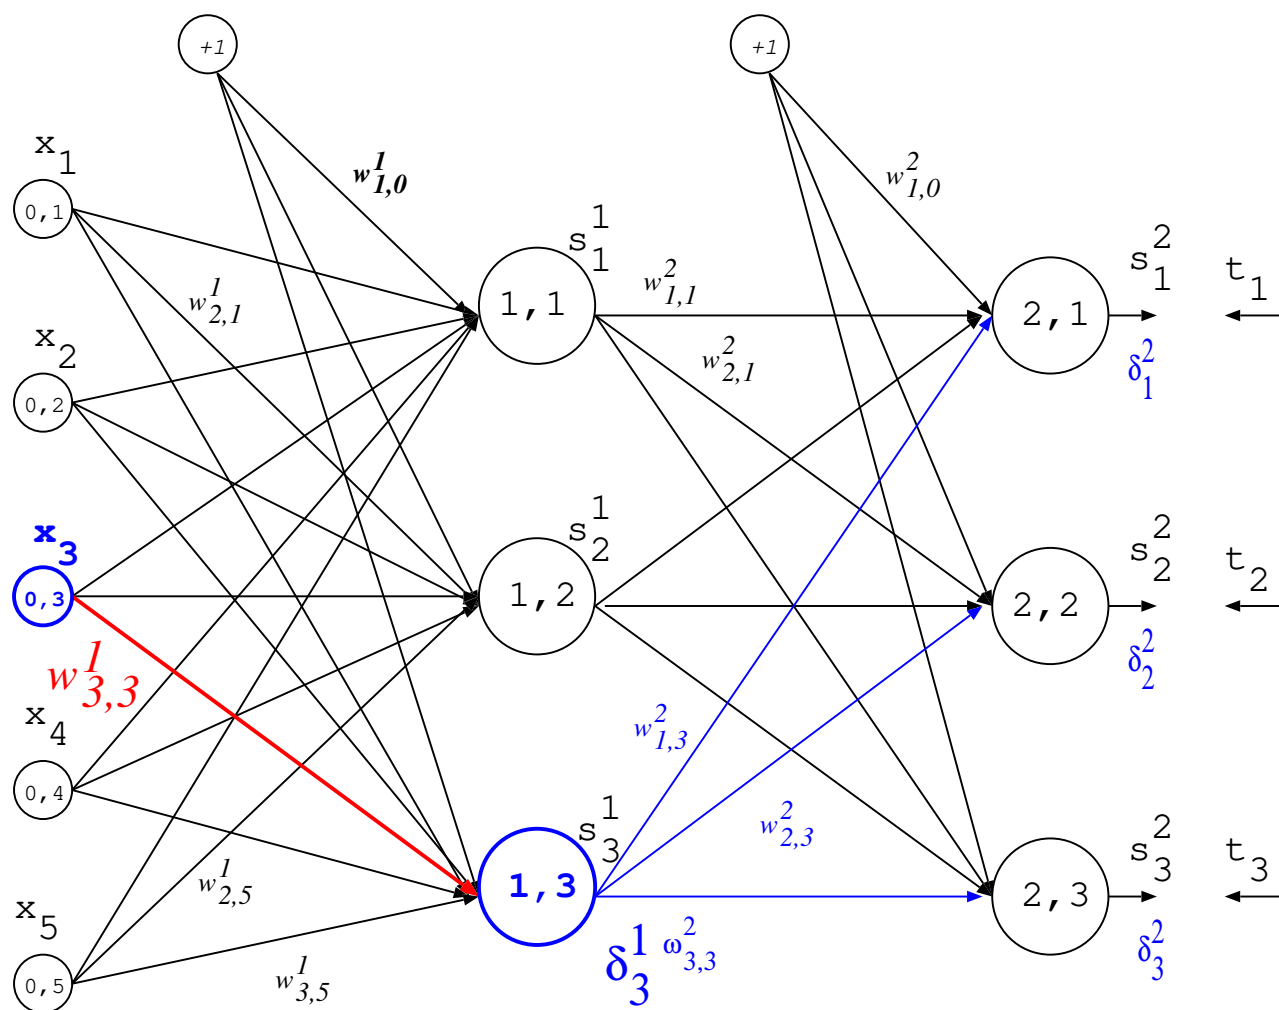
- Update the weight of the hidden layer $w_{i,j}^1$ (for $N = 1$)

$$E_A(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^{M_2} \left(t_k - s_k^2 \right)^2; \quad s_k^2 = f(z_k^2); \quad z_k^2 = \sum_{m=1}^{M_1} w_{k,m}^2 s_m^1; \quad s_m^1 = f(z_m^1); \quad z_m^1 = \sum_{l=1}^{M_0} w_{m,l}^1 x_l$$

$$\begin{aligned} \frac{\partial E_A}{\partial w_{i,j}^1} &= \sum_{k=1}^{M_2} \frac{\partial E_A}{\partial s_k^2} \frac{\partial s_k^2}{\partial z_k^2} \frac{\partial z_k^2}{\partial s_i^1} \frac{\partial s_i^1}{\partial z_i^1} \frac{\partial z_i^1}{\partial w_{i,j}^1} = \sum_{k=1}^{M_2} -\delta_k^2 w_{k,i}^2 f'(z_i^2) x_j = \\ &\quad - \left(f'(z_i^2) \sum_{k=1}^{M_2} \delta_k^2 w_{k,i}^2 \right) x_j = -\delta_i^1 x_j \end{aligned}$$

$$\Delta w_{i,j}^1 = -\rho \frac{\partial E_A}{\partial w_{i,j}^1} = \rho \delta_i^1 x_j$$

Derivation of backpropagation algorithm (4)



$$\Delta w^1_{3,3} = \rho \delta^1_3 x_3 = \rho \left(\sum_r \delta^2_r w^2_{r,3} \right) f'(z^1_3) x_3$$

BackProp for N training samples

- Updating the weights of the output layer: ($1 \leq i \leq M_2$, $0 \leq j \leq M_1$):

$$\Delta w_{ij}^2 = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{ij}^2} = \frac{\rho}{N} \sum_{n=1}^N \delta_i^2(\mathbf{x}_n) s_j^1(\mathbf{x}_n)$$

$$\delta_i^2(\mathbf{x}_n) = (t_{ni} - s_i^2(\mathbf{x}_n)) f'(z_i^2(\mathbf{x}_n)) \text{ with } z_i^2(\mathbf{x}_n) = \sum_{j=0}^{M_1} w_{ij}^2 s_j^1(\mathbf{x}_n)$$

- Updating the weights of the hidden layer: ($1 \leq i \leq M_1$, $0 \leq j \leq M_0$):

$$\Delta w_{ij}^1 = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{ij}^1} = \frac{\rho}{N} \sum_{n=1}^N \delta_i^1(\mathbf{x}_n) x_{nj}$$

$$\delta_i^1(\mathbf{x}_n) = \left(\sum_{r=1}^{M_2} \delta_r^2(\mathbf{x}_n) w_{ri}^2 \right) f'(z_i^1(\mathbf{x}_n)) \text{ with } z_i^1(\mathbf{x}_n) = \sum_{j=0}^{M_0} w_{ij}^1 x_{nj}$$

BackProp for N training samples for a three-layer perceptron

- Updating the weights of the output layer: ($1 \leq i \leq M_3$, $0 \leq j \leq M_2$)

$$\Delta w_{ij}^3 = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{ij}^3} = \frac{\rho}{N} \sum_{n=1}^N \delta_i^3(\mathbf{x}_n) s_j^2(\mathbf{x}_n) \quad \delta_i^3(\mathbf{x}_n) = \left(t_{ni} - s_i^3(\mathbf{x}_n) \right) f'(z_i^3(\mathbf{x}_n))$$

- Updating the weights of the second hidden layer: ($1 \leq i \leq M_2$, $0 \leq j \leq M_1$)

$$\Delta w_{ij}^2 = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{ij}^2} = \frac{\rho}{N} \sum_{n=1}^N \delta_i^2(\mathbf{x}_n) s_j^1(\mathbf{x}_n) \quad \delta_i^2(\mathbf{x}_n) = \left(\sum_{r=1}^{M_3} \delta_r^3(\mathbf{x}_n) w_{ri}^3 \right) f'(z_i^2(\mathbf{x}_n))$$

- Updating the weights of the first hidden layer: ($1 \leq i \leq M_1$, $0 \leq j \leq M_0$)

$$\Delta w_{ij}^1 = -\rho \frac{\partial E_A(\mathbf{w})}{\partial w_{ij}^1} = \frac{\rho}{N} \sum_{n=1}^N \delta_i^1(\mathbf{x}_n) x_{nj} \quad \delta_i^1(\mathbf{x}_n) = \left(\sum_{r=1}^{M_2} \delta_r^2(\mathbf{x}_n) w_{ri}^2 \right) f'(z_i^1(\mathbf{x}_n))$$

BackProp algorithm

Input: Topology, Initial weights w_{ij}^l , $1 \leq l \leq L$, $1 \leq i \leq M_l$, $0 \leq j \leq M_{l-1}$,
learning rate ρ , Convergence conditions, N training samples A

Output: Weights of connections that minimize the mean squared error of A

While no convergence

For $1 \leq l \leq L$, $1 \leq i \leq M_l$, $0 \leq j \leq M_{l-1}$, initialize $\Delta w_{ij}^l = 0$

For each training sample $(\mathbf{x}, \mathbf{t}) \in A$

From the input to the output layers ($l = 0, \dots, L$):

For $1 \leq i \leq M_l$ if $l = 0$ then $s_i^0 = x_i$ else compute z_i^l y $s_i^l = f(z_i^l)$

From the output to the input layers ($l = L, \dots, 1$),

For each node ($1 \leq i \leq M_l$)

Compute $\delta_i^l = \begin{cases} f'(z_i^l) (t_{ni} - s_i^L) & \text{if } l == L \\ f'(z_i^l) (\sum_r \delta_r^{l+1} w_{ri}^{l+1}) & \text{otherwise} \end{cases}$

For each weight w_{ij}^l ($0 \leq j \leq M_{l-1}$) compute: $\Delta w_{ij}^l = \Delta w_{ij}^l + \rho \delta_i^l s_j^{l-1}$

For $1 \leq l \leq L$, $1 \leq i \leq M_l$, $0 \leq j \leq M_{l-1}$, update the weights: $w_{ij}^l = w_{ij}^l + \frac{1}{N} \Delta w_{ij}^l$

Computational cost for each iteration: $O(N D)$, $N = |A|$, $D = \text{number of weights}$

Incremental BackProp algorithm

Input: Topology, initial weights w_{ij}^l , $1 \leq l \leq L$, $1 \leq i \leq M_l$, $0 \leq j \leq M_{l-1}$,
learning rate ρ , convergence conditions, N training samples A

Salidas: Weights that minimize the mean squared error of A

While no convergence

For each training sample $(\mathbf{x}, \mathbf{t}) \in A$ (in random order)

From the input to the output layers ($l = 0, \dots, L$):

For $1 \leq i \leq M_l$ if $l = 0$ then $s_i^0 = x_i$ else compute z_i^l y $s_i^l = f(z_i^l)$

From the output to the input layer ($l = L, \dots, 1$),

For each node ($1 \leq i \leq M_l$)

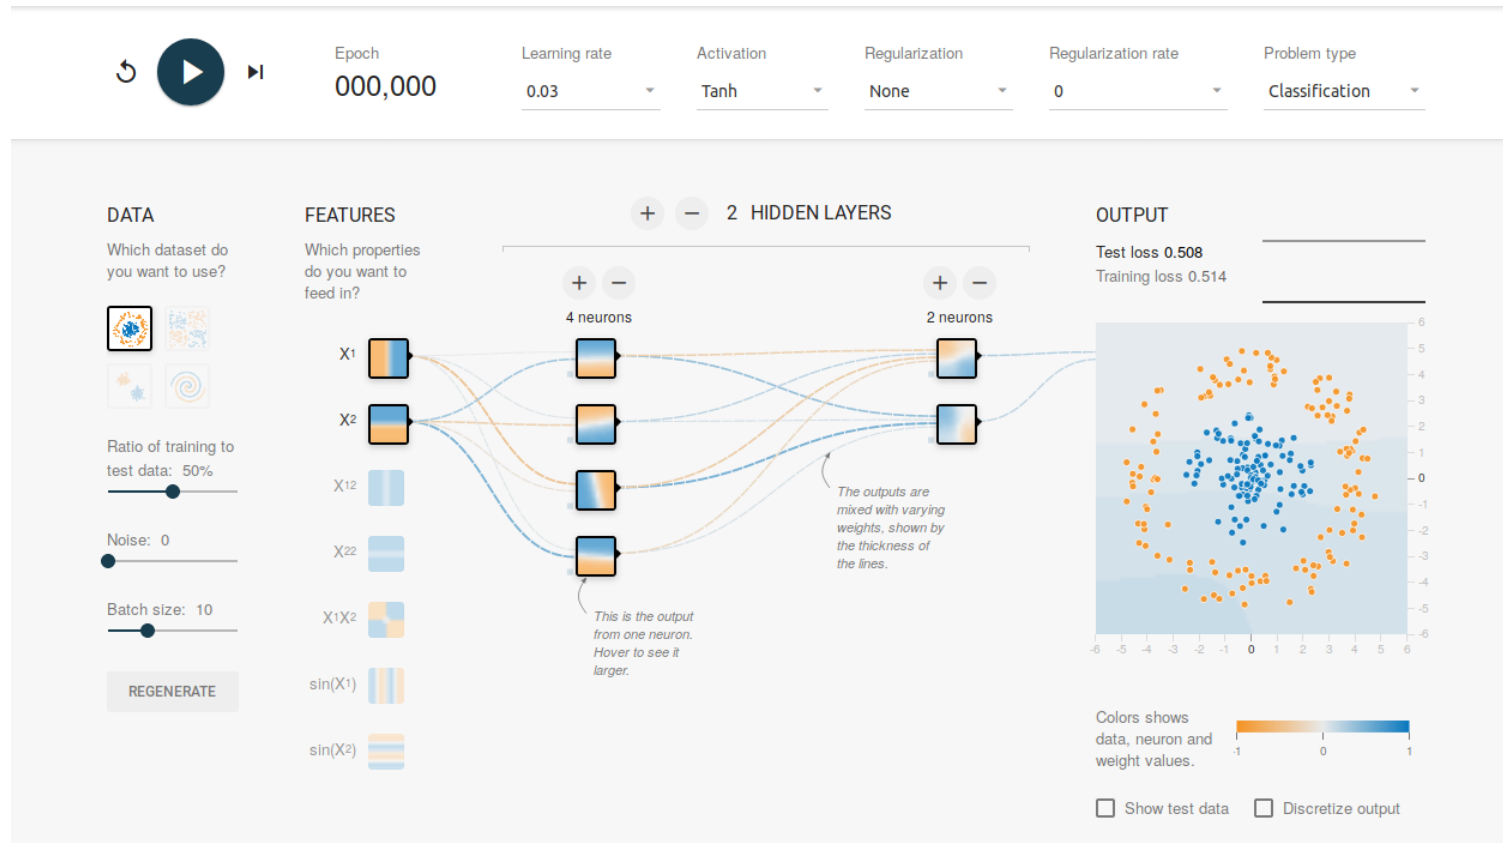
Compute $\delta_i^l = \begin{cases} f'(z_i^l) (t_{ni} - s_i^L) & \text{if } l == L \\ f'(z_i^l) (\sum_r \delta_r^{l+1} w_{ri}^{l+1}) & \text{otherwise} \end{cases}$

For each weight w_{ij}^l ($0 \leq j \leq M_{l-1}$) compute: $\Delta w_{ij}^l = \rho \delta_i^l s_j^{l-1}$

For $1 \leq l \leq L$, $1 \leq i \leq M_l$, $0 \leq j \leq M_{l-1}$, update weights: $w_{ij}^l = w_{ij}^l + \frac{1}{N} \Delta w_{ij}^l$

Computational cost for each iteration *mientras*: $O(N D)$, $N = |A|$, $D =$ number of weights

A demo of BackProp



<http://playground.tensorflow.org/>

A particular case: The Widrow-Hoff algorithm (Adaline)

- Given a training set $A = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^{D+1}$, $t_n \in \mathbb{R}$, search for $\mathbf{w} \in \mathbb{R}^{D+1}$: $\mathbf{w}^t \mathbf{x}_n = t_n$ (or $\mathbf{w}^t \mathbf{x}_n \approx t_n$) $1 \leq n \leq N$ ($\mathbf{x}_n = (1, x_{n1}, x_{n2}, \dots, x_{nD})$ and $\mathbf{w} = (w_0, w_1, \dots, w_D)$)

- Minimize the Widrow-Hoff function:
$$J_A(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^t \mathbf{x}_n - t_n)^2$$

- Solution by gradient descent:

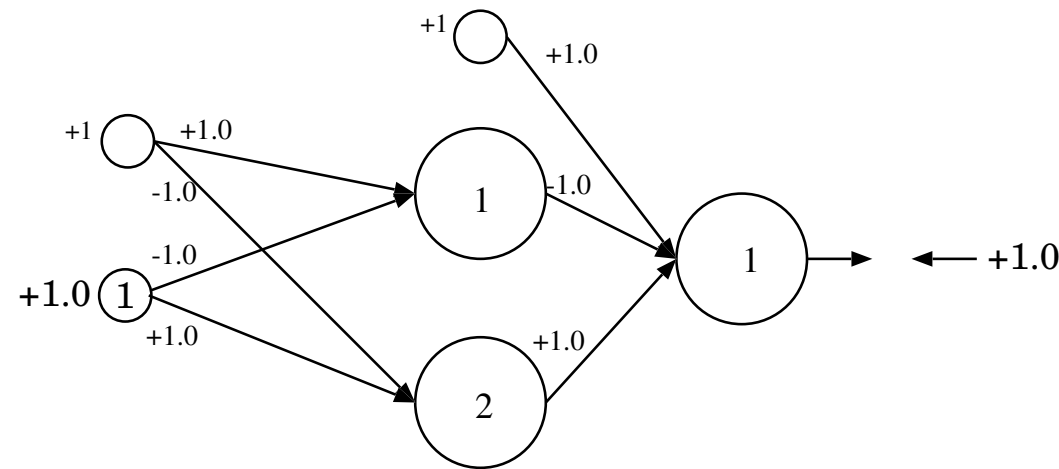
$$\begin{aligned} \mathbf{w}(1) &= \text{arbitrary} \\ \mathbf{w}(l+1) &= \mathbf{w}(l) - \rho_l \sum_{n=1}^N (\mathbf{w}(l)^t \mathbf{x}_n - t_n) \mathbf{x}_n \end{aligned}$$

- Correction sample by sample or on-line:

$$\begin{aligned} \mathbf{w}(1) &= \text{arbitrary} \\ \mathbf{w}(l+1) &= \mathbf{w}(l) + \rho_l (t(l) - \mathbf{w}(l)^t \mathbf{x}(l)) \mathbf{x}(l) \end{aligned}$$

Exercise

- The multiplayer perceptron of the figure is used for a regression problem with the hyperbolic tangent function as the activation function for all the nodes and a learning rate of $\rho = 0.5$.



Given an input sample $x = 1$ and the corresponding target $t = +1$, calculate:

- The outputs of all the nodes
- The corresponding errors in the output and hidden node
- The new weights

BackProp for classification

- Softmax is used as the activation function in the output layer.
- Training

Problem: Given a topology of a multilayer perceptron and $A = \{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$, with $\mathbf{x}_n \in \mathbb{R}^{M_0}$, $\mathbf{t}_n \in \{0, 1\}^{M_2 \equiv C}$, search for \mathbf{w} such that minimizes the objective function (cross-entropy):

$$C_A(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^{M_2} t_{n,k} \log s_k^2(\mathbf{x}_n; \mathbf{w})$$

Solution, gradient descent: $\Delta w_{i,j}^k = -\rho \frac{\partial C_A(\mathbf{w})}{\partial w_{i,j}^k}$

- The cross-entropy leads to faster training and improved generalization (Bishop 2006)

Derivation of BackProp for classification

- For $N = 1$, $C_A(\mathbf{w}) = - \sum_{k=1}^{M_2} t_k \log s_k^2(\mathbf{x}; \mathbf{w})$
- Solution, gradient descent: $\Delta w_{i,j}^k = -\rho \frac{\partial C_A(\mathbf{w})}{\partial w_{i,j}^k}$
- Update weight of the output layer $w_{i,j}^2$ (if $N = 1$) ($z_k^2 = \sum_{m=1}^{M_1} w_{k,m}^2 s_m^1$)

$$C_A(\mathbf{w}) = \sum_{m=1}^{M_2} t_m \log s_m^2; \quad s_m^2 = f(z_m^2); \quad z_m^2 = f\left(\sum_{l=1}^{M_1} w_{m,l}^2 s_l^1\right)$$

$$\frac{\partial C_A}{\partial w_{i,j}^2} = \frac{\partial C_A}{\partial s_i^2} \frac{\partial s_i^2}{\partial z_i^2} \frac{\partial z_i^2}{\partial w_{i,j}^2} = \frac{t_i}{s_i^2} f'(z_i^2) s_j^1 = -\delta_i^2 s_j^1$$

$$\Delta w_{i,j}^2 = -\rho \frac{\partial E_A}{\partial w_{i,j}^2} = -\rho \frac{t_i}{s_i^2} f'(z_i^2) s_j^1 = \rho \delta_i^2 s_j^1$$

Convergence of the error backpropagation algorithm

GENERAL THEOREM OF CONVERGENCE: *Let λ_k be the eigenvalues of the matrix $\frac{\partial^2 E_A(\mathbf{w})}{\partial \omega_i \partial \omega_j}$ for a given \mathbf{w} . If $|1 - \lambda_k \rho| < 1 \ \forall k$, when the number of iterations tends to ∞ , \mathbf{w} tends to a local minimum of $E_A(\mathbf{w})$*

Learning factor

- $\rho < 2/\lambda_{max}$ (Bishop, 95)
- Large $\rho \Rightarrow$ fast convergence and tendency to oscillate.
- Small $\rho \Rightarrow$ slow convergence.

Probabilistic interpretation of the output of a multilayer perceptron

1. A multilayer perceptron of L layers as a classifier in C classes
2. $A = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$ with $\mathbf{x}_n \in \mathbb{R}^d, \mathbf{t}_n \in \mathbb{R}^C$ $1 \leq n \leq N$ and $t_{n,k} = 1$ if $c(\mathbf{x}_n) = k$ and $t_{n,k} = 0$ otherwise for $1 \leq k \leq C$
3. The squared error:

$$E_A(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N \sum_{k=1}^C (t_{n,k} - s_k^L(\mathbf{x}_n))^2$$

ASSUMPTIONS: A has been generated following a distribution $\Pr(\mathbf{x}, \mathbf{t})$ and A is sufficiently huge and representative of \Pr :

If a global minimum of the mean squared error is reached and the a-posteriori probability is implementable by means of a multilayer perceptron, then, the outputs of the multilayer perceptron implement the a-posteriori probability underlying in the training samples: $s_k^L(\mathbf{x}) = \Pr(k | \mathbf{x})$.

Likelihood and squared error

- Assume that: $t_{n,k} = s_k^L(\mathbf{x}_n; \mathbf{w}) + \epsilon$, where ϵ is a Gaussian noise:

$$p(t_{n,k} | \mathbf{x}_n; \mathbf{w}) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left(-\frac{(s_k^L(\mathbf{x}_n; \mathbf{w}) - t_{nk})^2}{2\sigma^2} \right)$$

- As $\mathbf{t}_n \in \mathbb{R}^C$, let us assume that $p(\mathbf{t}_n | \mathbf{x}_n; \mathbf{w}) = \prod_{k=1}^C p(t_{n,k} | \mathbf{x}_n; \mathbf{w})$
- Given a training sample $A = \{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ the **maximum likelihood estimation** of \mathbf{w} is:

$$\begin{aligned} \operatorname{argmax}_{\mathbf{w}} \prod_{n=1}^N p(\mathbf{t}_n | \mathbf{x}_n; \mathbf{w}) &= \operatorname{argmin}_{\mathbf{w}} \left(-\sum_{n=1}^N \log p(\mathbf{t}_n | \mathbf{x}_n; \mathbf{w}) \right) \\ &= \operatorname{argmin}_{\mathbf{w}} \left(\sum_{n=1}^N \sum_{k=1}^C (s_k^L(\mathbf{x}_n; \mathbf{w}) - t_{n,k})^2 \right) \end{aligned}$$

In this case, the maximum likelihood estimation conveys to a **mean squared error minimization**.

Index

- 1 Introduction ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 *Some issues about BackProp* ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 Bibliography ▷ 66

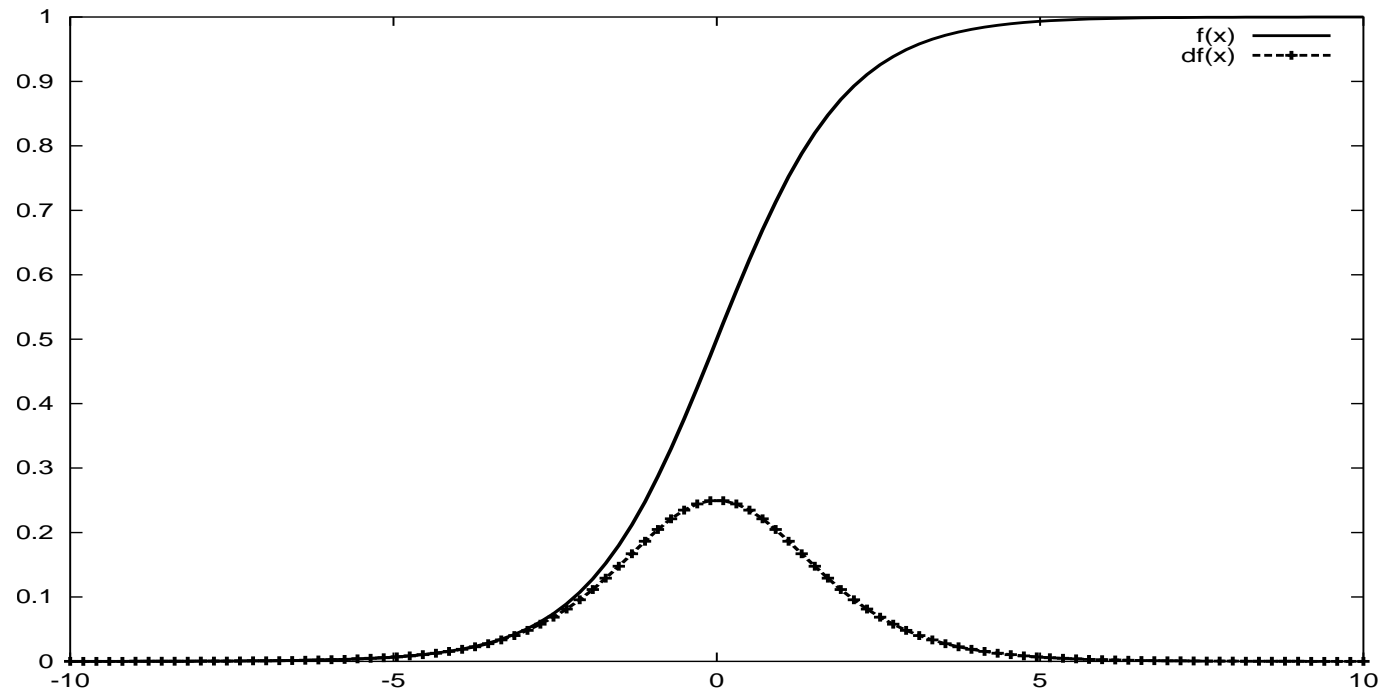
Derivatives of activation functions

- *Linear*: $f_L(z) = z \Rightarrow f'_L(z) = 1, z \in \mathbb{R}$.
- *Step*: $f_E(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z < 0 \end{cases} \Rightarrow f'_E(z) = \begin{cases} 0 & \text{if } z > 0 \text{ or } z < 0 \\ \text{not deriv.} & z = 0 \end{cases}$
- *ReLU*: $f_R(z) = \max(0, z) \Rightarrow f'_R(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z < 0 \\ \text{not derivable} & z = 0 \end{cases}$
- *PReLU*: $f_{PR}(z) = \begin{cases} z & \text{if } z > 0 \\ a \cdot z & \text{if } z < 0 \end{cases} \Rightarrow f'_R(z) = \begin{cases} 1 & \text{if } z > 0 \\ a & \text{if } z < 0 \\ \text{not derivable} & z = 0 \end{cases}$
- *Sigmoid*: $f_S(z) = \frac{1}{1+\exp(-z)} \Rightarrow f'_S(z) = f_S(z) (1 - f_S(z))$
- *Hyperbolic tangent*: $f_T(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)} \Rightarrow f'_T(z) = 1 - (f_T(z))^2$
- *Softmax*: $f_{SM}(z_k) = \frac{\exp(z_k)}{\sum_{k'} \exp(z_{k'})} \Rightarrow f'_{SM}(z_k) = f_{SM}(z_k) (1 - f_{SM}(z_k))$

Network paralysis

$$f(x) = \frac{1}{1 + \exp(-x)}$$

$$\frac{df(x)}{dx} = f(x) (1 - f(x))$$



PROBLEM: **Vanishing and exploding gradients**

Batch and online BackProp

- **Off-line or batch algorithm:** 1 update of the weights by epoch.
- **Mini-batch training:** B a subset of A . Batch algorithm applied to B .
- **Incremental algorithm:** sample $\mathbf{x}(l)$ at iteration l . $|A|$ updates of the weights by epoch.
- **Online algorithm:** each sample \mathbf{x} is used only once.

Gradient descent optimization algorithms (Ruder 2016)

- Stochastic gradient descent.
- Stochastic gradient descent with momentum.
- Adagrad (Adaptive Gradient)
- Adadelta (an extension of Adagrad)
- Adam (Adaptive Moment Estimation):
- Nesterov accelerated gradient, RMSProp, AdaMax, Nadam, ...

Gradient descent optimization algorithms (Ruder 2016) (1)

- **Stochastic gradient descent**: gradient descent in the l mini-batch.

$$\Delta \mathbf{w}(l) = \rho \nabla_{\mathbf{w}} E_B(\mathbf{w}(l))$$

- Stochastic gradient descent with **momentum**:

$$\Delta \mathbf{w}(l) = \rho \nabla_{\mathbf{w}} E_B(\mathbf{w}(l)) + \gamma \Delta \mathbf{w}(l - 1)$$

Gradient descent optimization algorithms (Ruder 2016) (2)

- **Adagrad** (Adaptive Gradient): ($\times \equiv$ element-wise product)

$$\mathbf{m}(l) = \mathbf{m}(l-1) + (\nabla_{\mathbf{w}} E_B(\mathbf{w}(l)) \times \nabla_{\mathbf{w}} E_B(\mathbf{w}(l)))$$

$$\mathbf{q}(l) : \forall i, \quad q_i(l) = \frac{\rho}{\sqrt{m_i(l) + \epsilon}}$$

$$\Delta \mathbf{w}(l) = \mathbf{q}(l) \times \nabla_{\mathbf{w}} E_B(\mathbf{w}(l))$$

Gradient descent optimization algorithms (Ruder 2016) (3)

- **Adadelta** (an extension of Adagrad);

$$\mathbf{m}(l) = \gamma_1 \mathbf{m}(l-1) + (1 - \gamma_1) (\nabla_{\mathbf{w}} E_B(\mathbf{w}(l)) \times \nabla_{\mathbf{w}} E_B(\mathbf{w}(l)))$$

$$\mathbf{v}(l) = \gamma_2 \mathbf{v}(l-1) + (1 - \gamma_2) (\Delta \mathbf{w}(l-1) \times \Delta \mathbf{w}(l-1))$$

$$\mathbf{q}(l) : \forall i, \quad q_i(l) = \rho \sqrt{\frac{v_i(l-1)}{m_i(l) + \epsilon}}$$

$$\Delta \mathbf{w}(l) = \mathbf{q}(l) \times \nabla_{\mathbf{w}} E_B(\mathbf{w}(l))$$

Gradient descent optimization algorithms (Ruder 2016) (4)

- **Adam** (Adaptive Moment Estimation):

$$\mathbf{m}(l) = \gamma_1 \mathbf{m}(l-1) + (1 - \gamma_1) (\nabla_{\mathbf{w}} E_B(\mathbf{w}(l)) \times \nabla_{\mathbf{w}} E_B(\mathbf{w}(l)))$$

$$\mathbf{q}(l) : q_i(l) = \frac{1}{\sqrt{\frac{m_i(l)}{1-\gamma_1} + \epsilon}}$$

$$\mathbf{v}(l) = \frac{\rho}{1 - \gamma_2} (\gamma_2 \mathbf{v}(l-1) + (1 - \gamma_2) \nabla_{\mathbf{w}} E_B(\mathbf{w}(l)))$$

$$\Delta \mathbf{w}(l) = \mathbf{q}(l) \times \mathbf{v}(l)$$

Input normalization and weight initialization

- **Input coding:** Normalize the input range to $[0, 1]$.

$$A = \{\mathbf{x}_1, \dots, \mathbf{x}_N\} \subset \mathbb{R}^D \Rightarrow \begin{cases} \mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j} \\ \sigma_j^2 = \frac{1}{N-1} \sum_{i=1}^N (x_{i,j} - \mu_j)^2 \end{cases} \quad 1 \leq j \leq D$$

$$\forall \mathbf{x} \in \mathbb{R}^D, \mathbf{x}^N : x_j^N = \frac{x_j - \mu_j}{\sigma_j} \Rightarrow \begin{cases} \mu_j^N = 0 \\ \sigma_j^N = 1 \end{cases} \quad \text{for } 1 \leq j \leq D$$

- **Weight initialization:** (n is the size of previous layer)

$$\left[-\frac{1}{\sqrt{n}}, +\frac{1}{\sqrt{n}} \right]$$

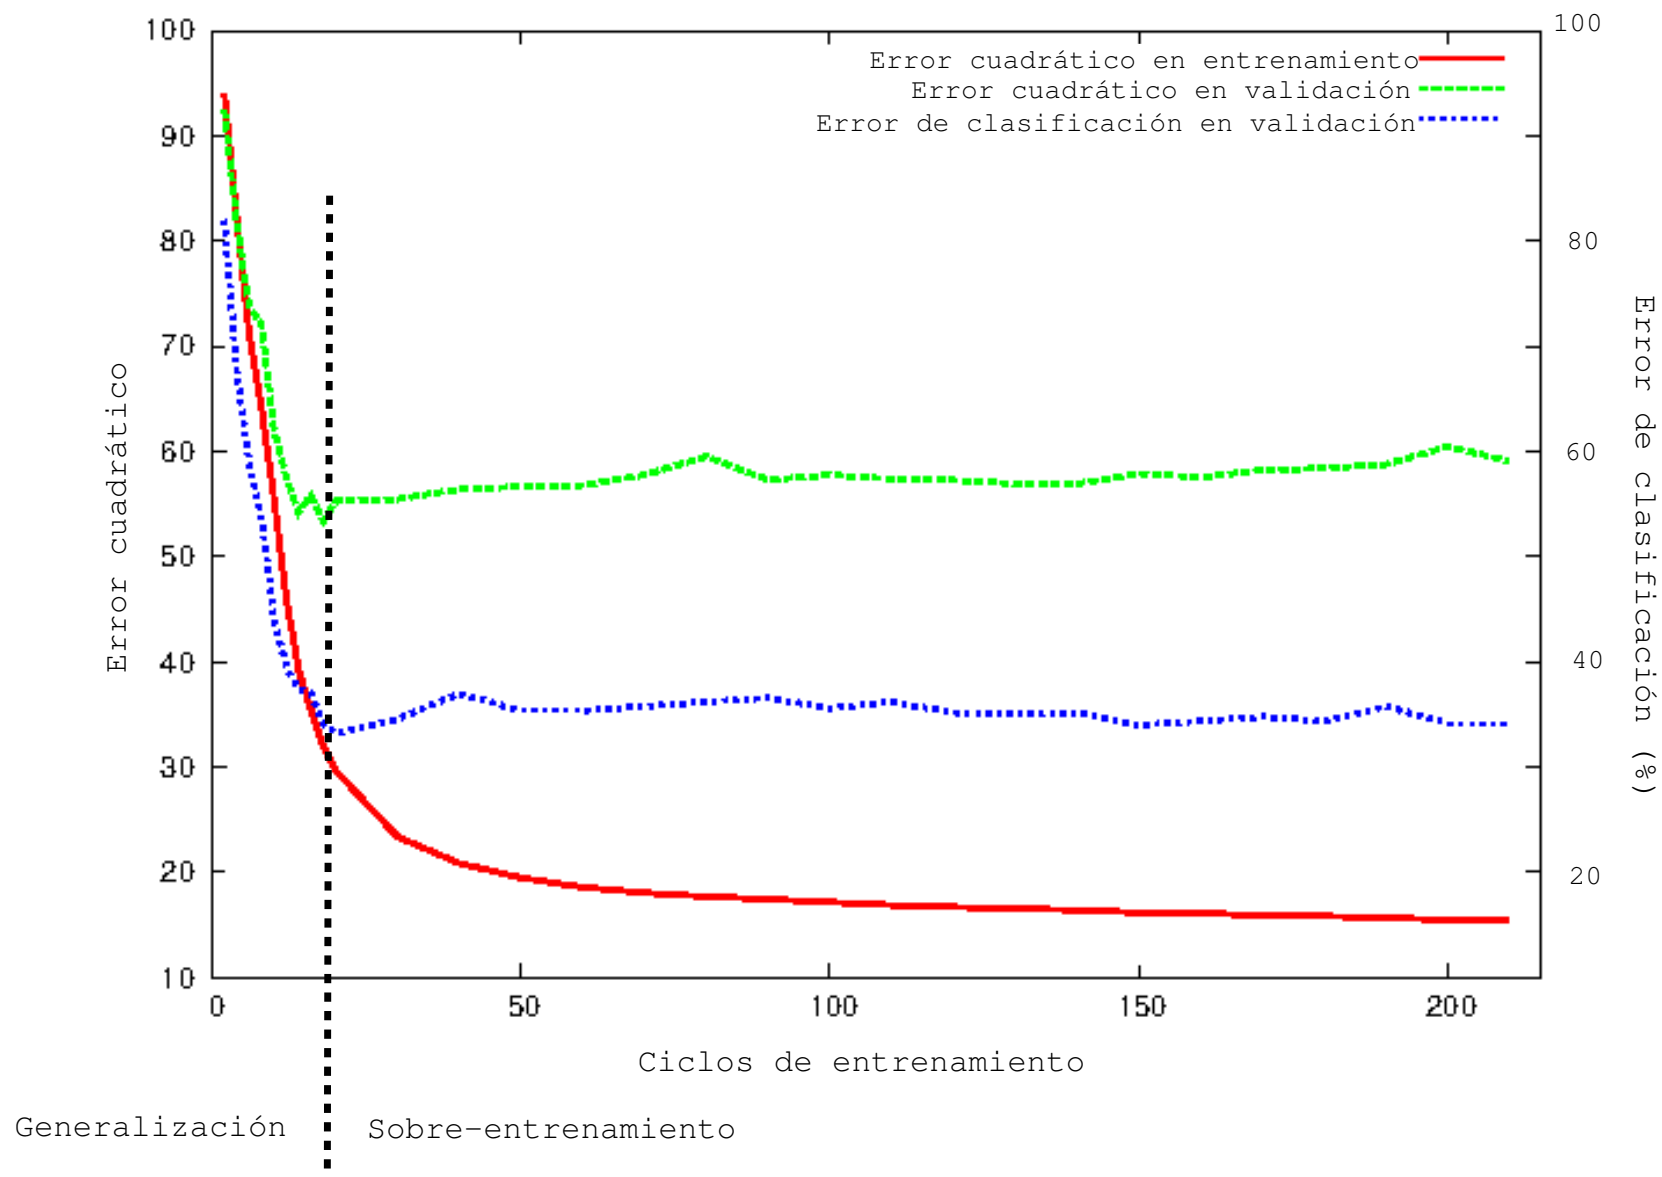
Regularization

- Problem: to prevent very big weights.
- Solution: add a regularization term to the goal function
 - Regularization $L_2 : q_S(\Theta) + \frac{\lambda}{2} \sum_{l,i,j} (\theta_{ij}^l)^2$
 - Regularization $L_1 : q_S(\Theta) + \frac{\lambda}{2} \sum_{l,i,j} \|\theta_{ij}^l\|$

Other techniques for avoiding “bad” local minimas [Koehn 2020]

- Shuffling the training data.
- Curriculum learning: From “easy” samples to “difficult” samples.
- Regularization: A new objective function to optimize, $E_A(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$.
- Adding a Gaussian noise ϵ_k : $\Delta\omega_{i,j}^k = \rho \left(\delta_i^k s_j^{k-1} + \epsilon_l \right)$

Convergence conditions



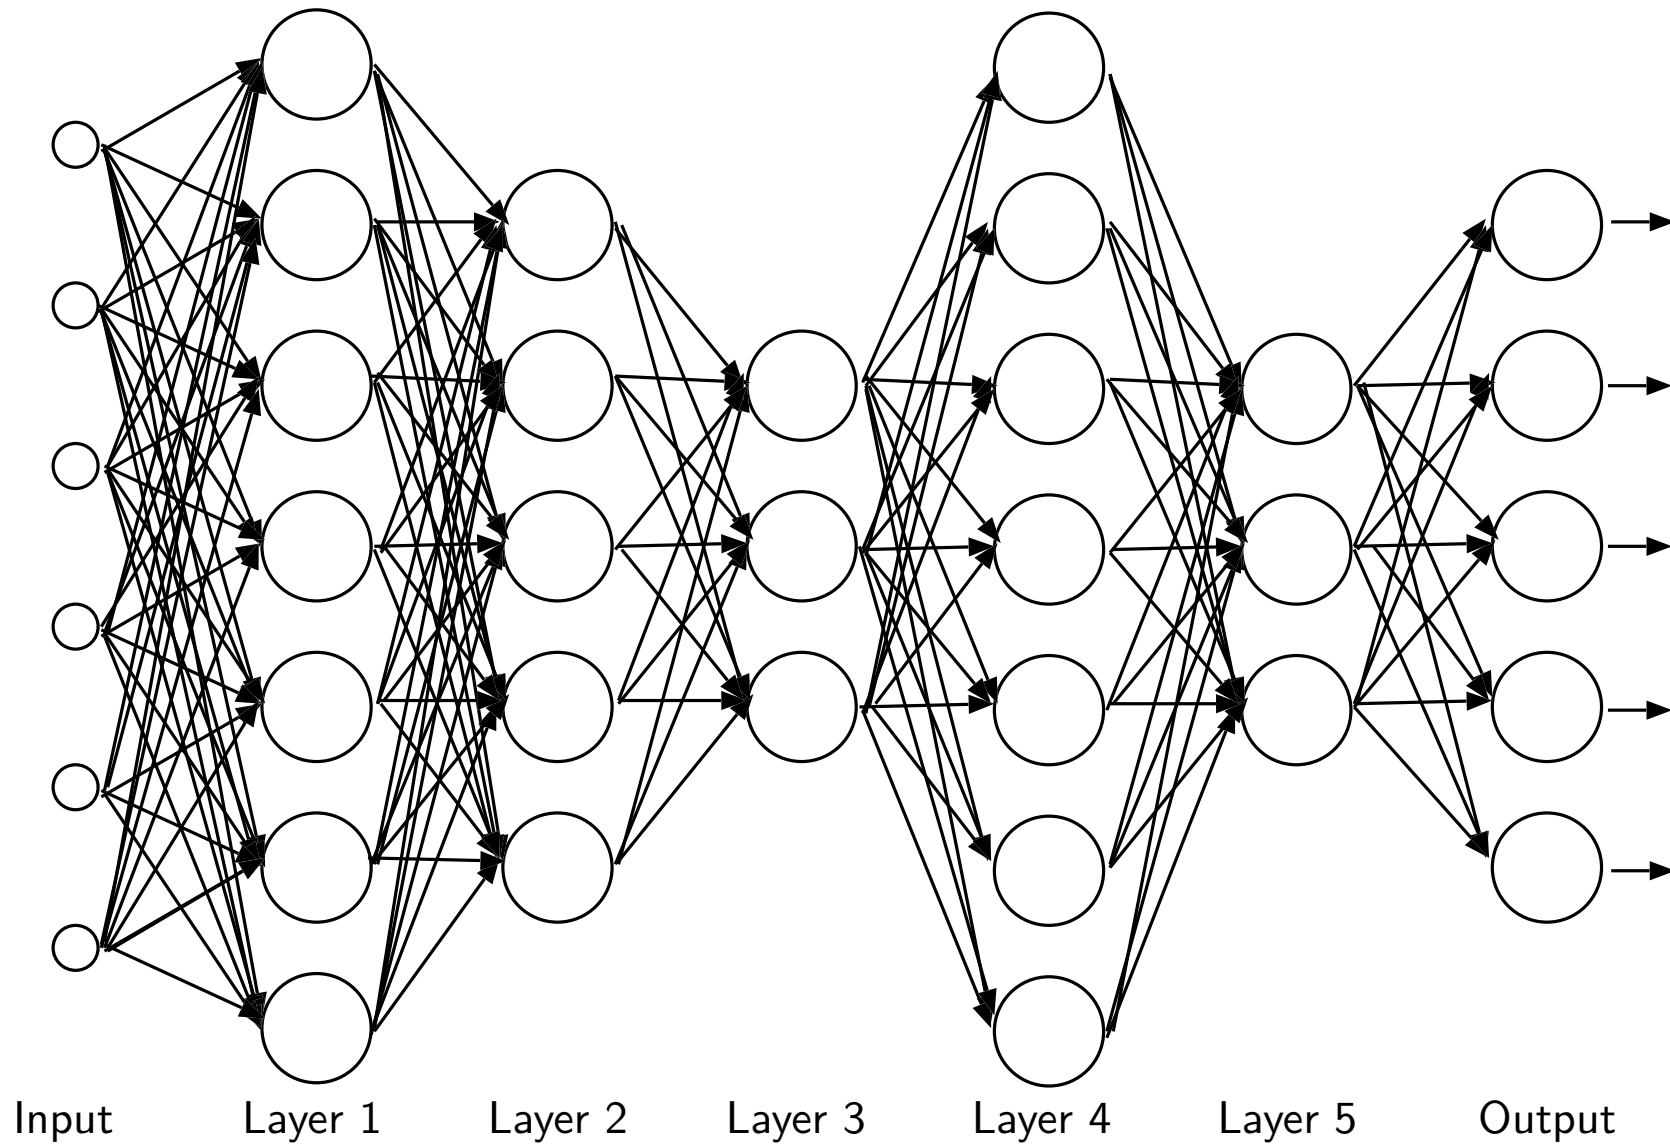
Validation approaches

- *Resubstitution method*
 - The training set = the test set
- *Hold-out method*
 - A training set and a validation set.
 - A test set for the evaluation.
- *Cross-validation*
 - Divide the training set in S parts.
 - For $i := 1$ to S
 - Use $S - 1$ parts as training set (and validation) and the rest as a test set.
 - The result of the evaluation is the average of the results on the S repetitions.
- *Leave-one-out method* (Cross-validation with $S = N$)

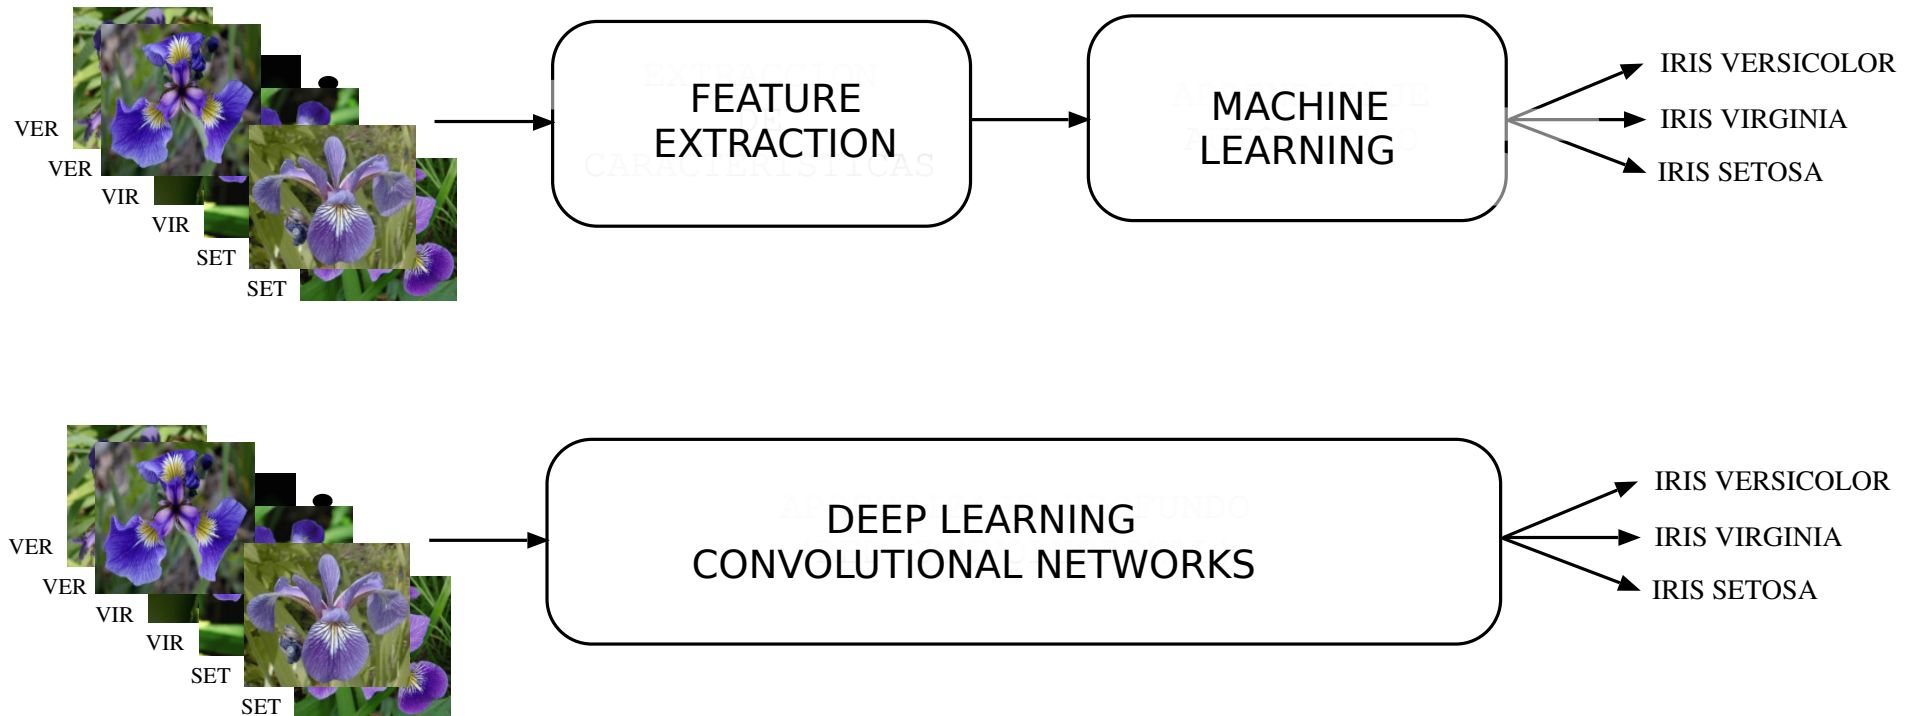
Index

- 1 Introduction ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 *A brief introduction to deep neural networks* ▷ 55
- 6 Bibliography ▷ 66

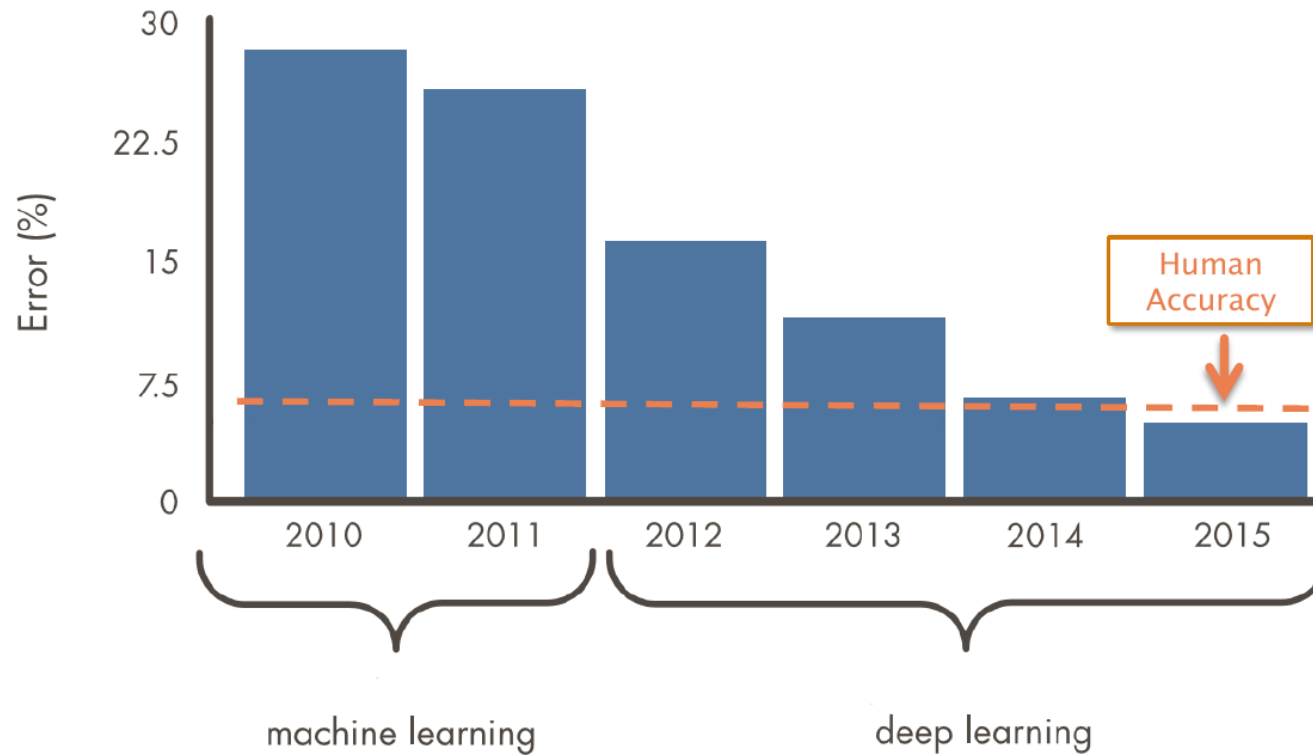
Deep neural network concept



Deep learning

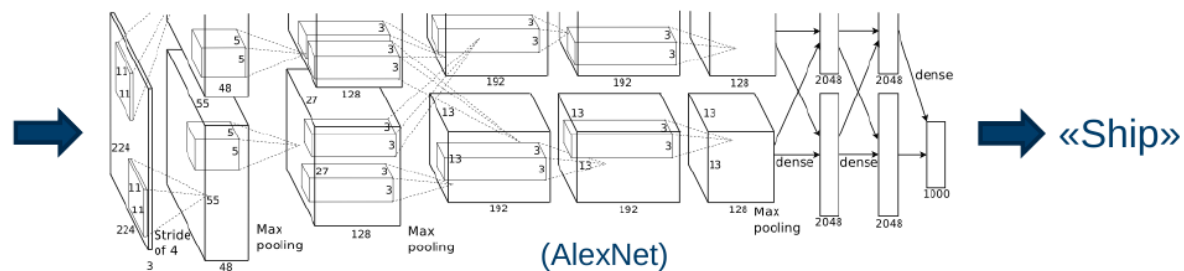
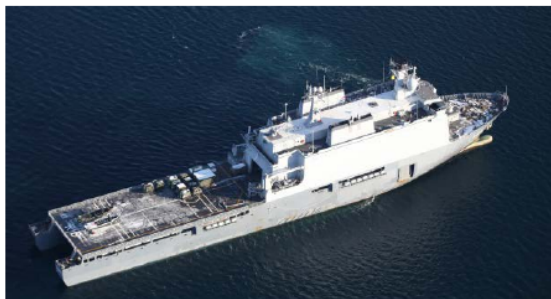


Deep learning [Daly 2017]



Source: ILSVRC Top-5 Error on ImageNet

Deep learning [Dyrda 2019]

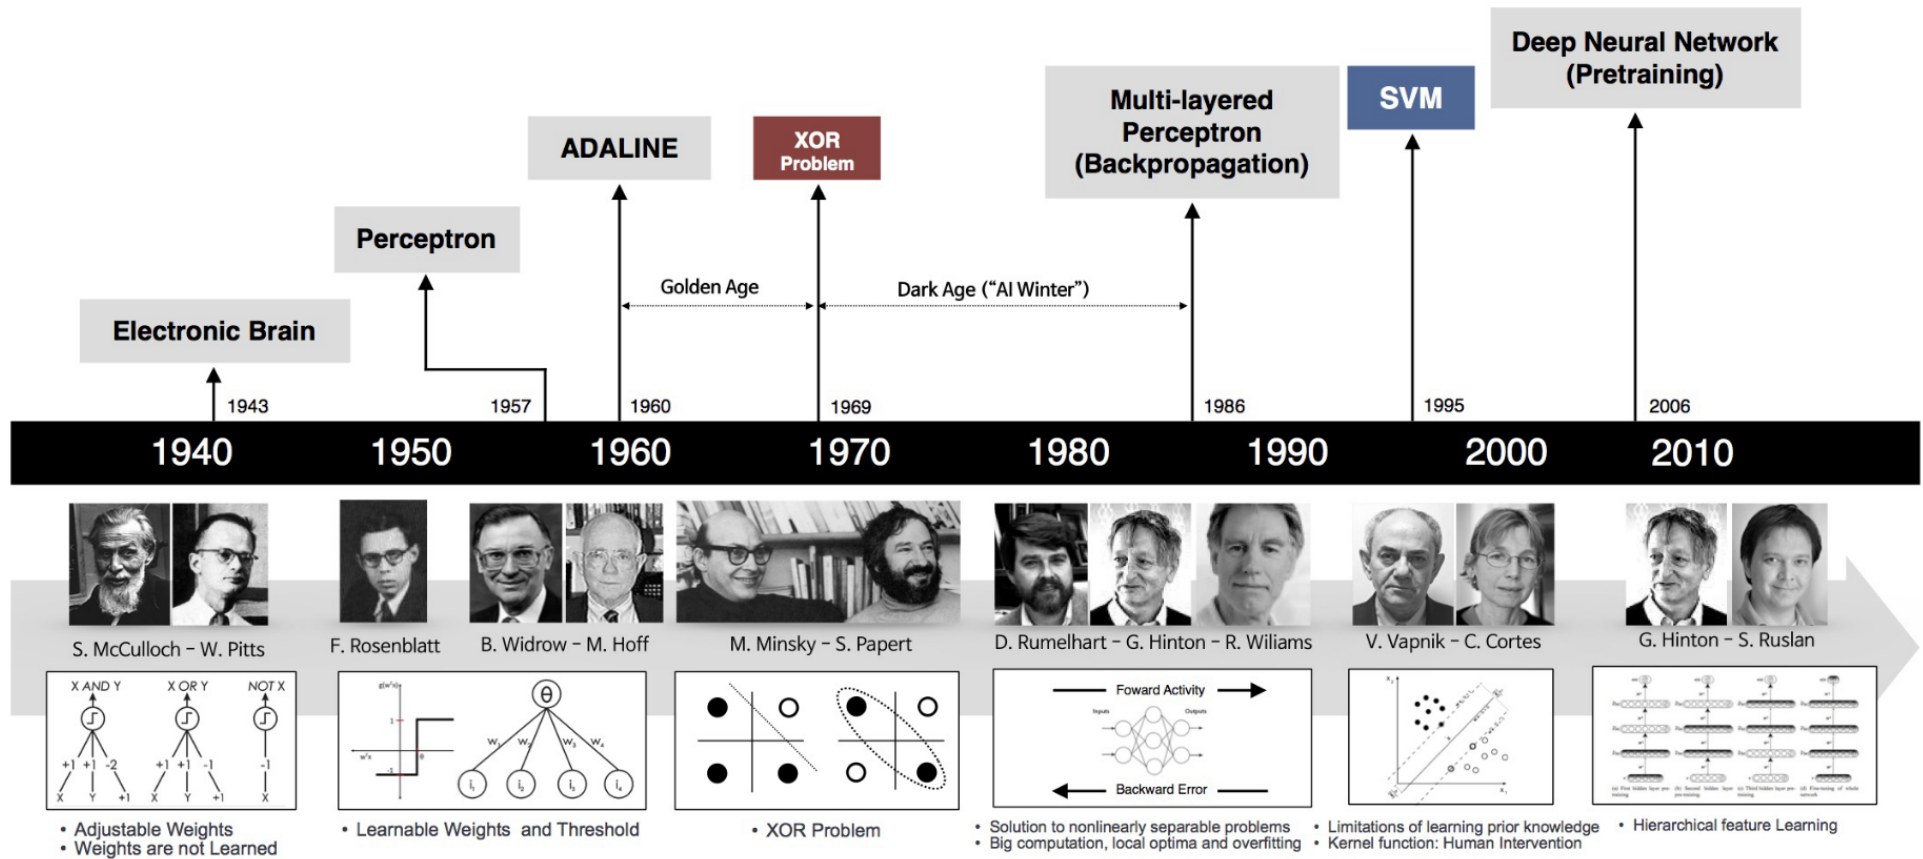


Millions of images

Millions of parameters

Thousands of classes

Deep learning [Serengil 2017]



Dynamic networks

- Recurrent networks:
 - Simple recurrent networks.
 - Elman network: recurrent + multiplayer perceptron,
 - Second order recurrent networks.
 - Long Short-Term Memory (LSTM)
 - Gated Recurrent Units (GRU)
- Feed-forward networks:
 - Convolutional networks.
 - **Transformer** (for text, images, tables, ...)
 - **Pre-trained networks**, usually based on Transformer: BERT (Google AI), GPT-3 (OpenAI), XML (Facebook), DALL-E 2 (OpenAI), BART (Facebook), Flamingo (DeepMind), ...
 - Prompting and few-shot (meta) learning.

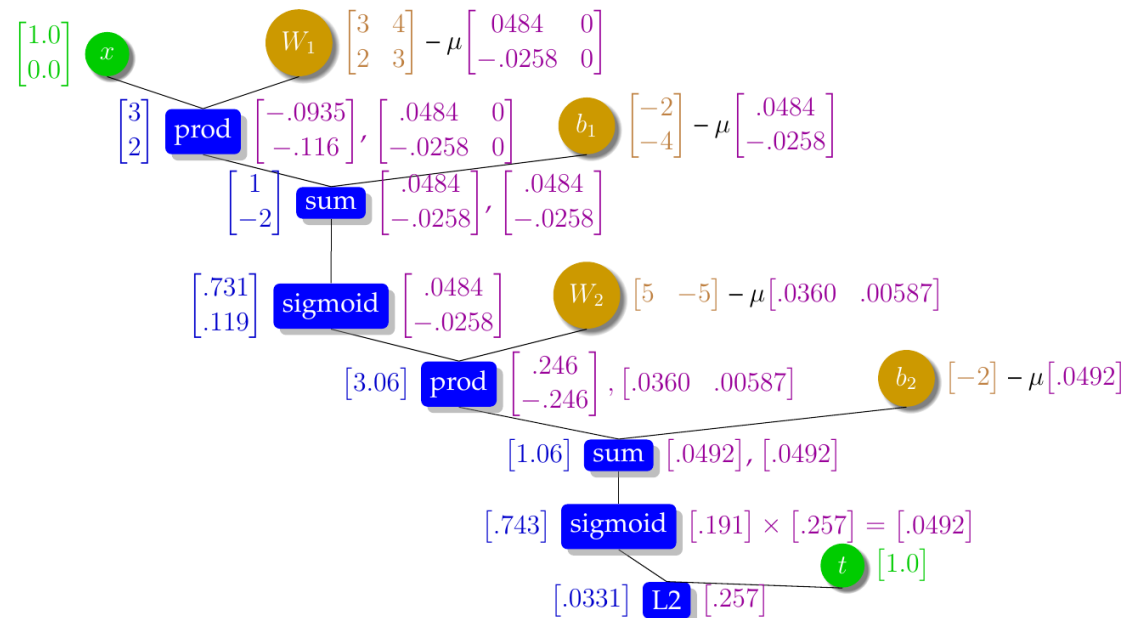
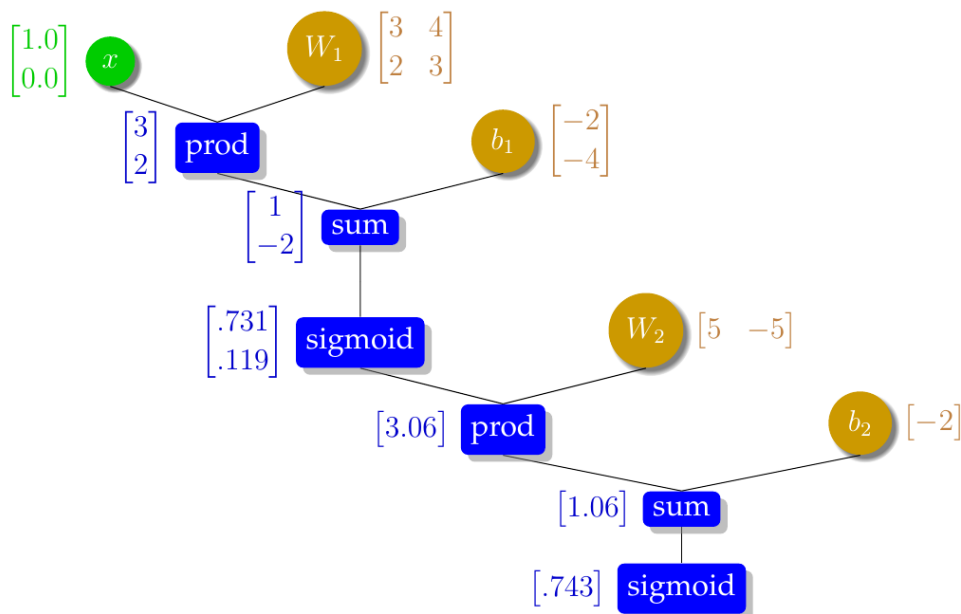
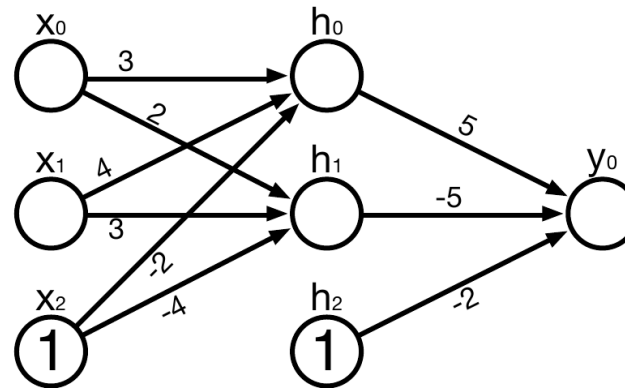
Avoiding vanishing and exploding gradients [Koehn 2020]

- Using activation functions such as Relu, LeakyRelu/PreLU, Maxout, ELU, ...
- **Dropout**: Randomly some nodes are ignored in each iteration.
- **Adaptive gradient clipping**.
- **Residual connections**: From $y = f(x)$ to $y = f(x) + x$.
- **Highway networks**: Residual networks+control gates.
- **Batch normalization**.
- **Layer normalization**.
- **Input normalization**.
- **Weight initialization**.
- **Regularization**.

Other issues

- Training issues:
 - Data augmentation.
- Computational issues:
 - Using **computational graphs** implemented in TensorFlow, PyTorch, ... (KoeHN 2020)
 - Use of graphics processing units (GPUs): the size of the minibatch is usual conditioned by the memory of the GPUs.

Computational graphs (Koehn 2020)



Computational graphs in PyTorch (<http://www.statmt.org/nmt-book/>)

```
import torch

# Data
W = torch.tensor([[3,4],[2,3]], requires_grad=True, dtype=torch.float)
b = torch.tensor([-2,-4], requires_grad=True, dtype=torch.float)
W2 = torch.tensor([5,-5], requires_grad=True, dtype=torch.float)
b2 = torch.tensor([-2], requires_grad=True, dtype=torch.float)
data = [ [ torch.tensor([0.,0.]), torch.tensor([0.]) ],
          [ torch.tensor([1.,0.]), torch.tensor([1.]) ],
          [ torch.tensor([0.,1.]), torch.tensor([1.]) ],
          [ torch.tensor([1.,1.]), torch.tensor([0.]) ] ]
mu = 0.1
```

```
for iteration in range(1000):
    # forward computation
    total_error = 0
    for item in data:
        x = item[0]
        t = item[1]
        s = W.mv(x) + b
        h = torch.nn.Sigmoid()(s)
        z = torch.dot(W2, h) + b2
        y = torch.nn.Sigmoid()(z)
        error = 1/2 * (t - y) ** 2
        total_error = total_error + error
    # backward computation
    total_error.backward()
    W.data = W - mu * W.grad.data
    b.data = b - mu * b.grad.data
    W2.data = W2 - mu * W2.grad.data
    b2.data = b2 - mu * b2.grad.data
    W.grad.data.zero_()
    b.grad.data.zero_()
    W2.grad.data.zero_()
    b2.grad.data.zero_()
    print("error: ",total_error.data/4)
```

Index

- 1 Introduction ▷ 3
- 2 Multilayer Perceptron ▷ 7
- 3 Error Back-propagation algorithm ▷ 22
- 4 Some issues about BackProp ▷ 41
- 5 A brief introduction to deep neural networks ▷ 55
- 6 *Bibliography* ▷ 66

Recommended bibliography

- Christopher M. Bishop. 2006. Pattern Recognition and Machine Learning (Information Science and Statistics). Springer-Verlag New York, Inc., Secaucus, NJ, USA.
- Christopher M. Bishop. 1995. Neural Networks for Pattern Recognition. Oxford University Press, Inc., New York, NY, USA.
- Ian Goodfellow, Yoshua Bengio and Aaron Courville. Deep Learning. MIT Press 2016. <http://www.deeplearningbook.org>
- Philipp Koehn. Neural Machine Translation. Cambridge University Press 2020