

Chapter 5. Deep Reinforcement Learning and Generative Models

Neural Networks

2022/2023

Máster Universitario en Inteligencia Artificial, Reconocimiento
de Formas e Imagen Digital

Departamento de Sistemas Informáticos y Computación

Index

- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Index

- 1 *Reinforcement Learning* ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Reinforcement Learning


- Goal: intelligent systems that interact with the environment
- Improving over time through trial and error
- RL have had some success in the past but still limited
- Deep RL rely on:
 - Neural networks as function approximation
 - Representation Learning
- Deep learning enables RL to scale to problems that were previously intractable:
 - settings with high-dimensional states (directly from pixels)
 - extremely huge action spaces

Reinforcement Learning

- Some examples:
 - Breakout: <https://www.youtube.com/watch?v=V1eYniJ0Rnk>
 - Super Mario Bros: <https://www.youtube.com/watch?v=qv6UV0Q0F44>

Reinforcement Learning

- Learning through interaction
- Trial and error learning
- Agent is a machine learning algorithm



Reinforcement Learning

RL as a Markov Decision Process (MDP):

- \mathcal{S} set of states
- $p(s_0)$ a distribution of starting states
- \mathcal{T} transitions, $\mathcal{T}(s_{t+1} \mid s_t, a_t)$
- \mathcal{R} reward function, $\mathcal{R}(s_t, a_t, s_{t+1})$
- γ is a discount factor, $\gamma \in [0, 1]$

Key concept: only s_t affects s_{t+1} no further past has to be considered


Reinforcement Learning

At time step t :

- Agent observes the environment state s_t
- Agent take an action a_t following a (stochastic) policy:

$$\pi : \mathcal{S} \rightarrow p(\mathcal{A} = a_t \mid \mathcal{S} = s_t)$$

- Agent and environment transition to new state s_{t+1}
- Environment provides a new reward r_{t+1}



- Goal: learn a *policy* that maximizes the cumulative reward

Reinforcement Learning

- Total Reward for T actions:

$$R = \sum_{t=0}^T \gamma^t r_{t+1}$$

- The goal of RL is to find an optimal policy π^* :

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}[R \mid \pi]$$

given an MDP, π^* exists.

Reinforcement Learning

- The optimal policy must be inferred by trial-and-error interaction with the environment. The only learning signal the agent receives is the reward
- The observations of the agent depend on its actions and can contain strong temporal correlations
- Agents must deal with long-range time dependencies: Often the consequences of an action only materialize after many transitions of the environment

Reinforcement Learning

- RL Algorithms:
 - Value functions: Indirect estimation of policy
 - Policy Search: Direct estimation of policy

Index

- 1 Reinforcement Learning ▷ 3
- 2 *Value Function Algorithms* ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Value Functions Algorithms

- Estimate the value (expected return) of being in a given state
- The state-value function:

$$V^\pi(s) = \mathbb{E}[R|s, \pi]$$

- the total reward starting in state s and following policy π
- the optimal state-value is the state-value with the optimal policy:

$$V^*(s) = \max_{\pi} V^\pi(s) \quad \forall s \in \mathcal{S}$$

- if we had V^* we could retrieve π^* by means of:

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{s_{t+1} \sim \mathcal{T}(s_{t+1}|s_t, a)} [V^*(s_{t+1})]$$

Value Functions Algorithms

- In this expression:

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{s_{t+1} \sim \mathcal{T}(s_{t+1}|s_t, a)} [V^*(s_{t+1})]$$

\mathcal{T} is unknown.

- we define $Q^\pi(s, a)$, the state-action-value or quality function:

$$Q^\pi(s, a) = \mathbb{E}[R \mid s, a, \pi]$$

- if we had $Q^\pi(s, a)$ the best action of the policy would be:

$$a^* = \operatorname{argmax}_a Q^\pi(s, a)$$

Value Functions Algorithms

How to learn $Q^\pi(s, a)$:

- Dynamic Programming (Bellman equation):

$$Q^\pi(s, a) = \mathbb{E}[r_{t+1} + \gamma Q^\pi(s_{t+1}, \pi(s_{t+1})) \mid s = s_t, a = a_t]$$

- optimal values decompose into a Bellman equation:

$$Q^*(s, a) = \mathbb{E}[r_{t+1} + \gamma \max_a Q^*(s_{t+1}, a) \mid s = s_t, a = a_t]$$

Value Functions Algorithms

How to learn $Q^*(s, a)$:

- Q-learning optimal values, iteratively:

$$Q(s, a) = Q(s, a) + \alpha \delta$$

- δ is the Temporal Difference error:

$$\delta = Y - Q(s, a)$$

in step t

$$Y = r_{t+1} + \gamma \max_a Q(s_{t+1}, a)$$

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha(r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$$

Value Functions Algorithms

- Q-learning algorithm:

Output: Q

initialize Q arbitrarily, e.g., to 0

Repeat

select s as an initial state

while(state s is not terminal) do

a = action for s derived by Q (epsilon-greedy)

take action a, observe r, sn

$Q(s,a) = Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s_n, a') - Q(s,a)]$

s=sn

end

Until convergence

end

Value Functions Algorithms

v

| | | | | | | | | | |
|-------|---|---|---|---|---|---|---|---|---|
| Start | | | ✗ | | | ✗ | | | ★ |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |

Actions:

R: right

JR: Jump right

Initial state: $v(0)$

Reward(s)

R

| | | | | | | | | | |
|-------|---|---|-----|---|---|-----|---|---|----|
| Start | | | ✗ | | | ✗ | | | ★ |
| 1 | 1 | 1 | -10 | 1 | 1 | -10 | 1 | 1 | 10 |

Value Functions Algorithms

| | | | | | | | | | | |
|---|-------|---|---|---|---|---|---|---|---|---|
| v | Start | | | ✗ | | | ✗ | | | ★ |
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |

Q function initial


| s a | R | JR |
|-------|---|----|
| 0 | 0 | 0 |
| 1 | 0 | 0 |
| 2 | 0 | 0 |
| 3 | 0 | 0 |
| 4 | 0 | 0 |
| 5 | 0 | 0 |
| 6 | 0 | 0 |
| 7 | 0 | 0 |
| 8 | 0 | 0 |
| 9 | 0 | 0 |

Q function obtained

| s a | R | JR |
|-------|-----|------|
| 0 | 10 | 10 |
| 1 | 10 | -9.5 |
| 2 | -10 | 10 |
| 3 | 0 | 0 |
| 4 | 10 | -10 |
| 5 | -10 | 10 |
| 6 | 0 | 0 |
| 7 | 10 | 10 |
| 8 | 10 | -8.3 |
| 9 | 0 | 0 |

Value Functions Algorithms


How to define the *state* of the agent? How to obtain $Q(s, a)$?



Value Functions Algorithms

Alternative way to learn $Q^\pi(s, a)$:

- Monte Carlo methods estimate the expected return from a state $\mathbb{E}[R \mid s, a, \pi]$ by averaging the return from multiple roll-outs of a policy



Index

- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 *Policy Search Algorithms* ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Policy Search

- No need a value function model
- Directly search for an optimal policy π^*
- Typically a **parameterized policy** and parameters (Φ) are updated to maximize $\mathbb{E}[R|\Phi]$
- Neural networks can be used to encode policies

Policy Search

Neural network outputs:

- **Continuous actions**
 - parameters of a probability distribution
 - e.g., mean and standard deviations
- **Discrete actions**
 - individual probabilities
 - multinomial distribution

In any case, actions are obtaining by sampling from these distributions that essentially define the transition dynamics \mathcal{T}

Policy Gradients

- Compute the expected return of a given policy in order to obtain the gradients w.r.t Φ :

$$\mathbb{E}[R \mid \pi]$$
$$\nabla_{\Phi} \mathbb{E}[R] = \mathbb{E}[R \nabla_{\Phi} \log \pi(a_t \mid s_t; \Phi)]$$

Update Φ by gradient ascent in the direction of:

$$R \nabla_{\Phi} \log \pi(a_t \mid s_t; \Phi)$$

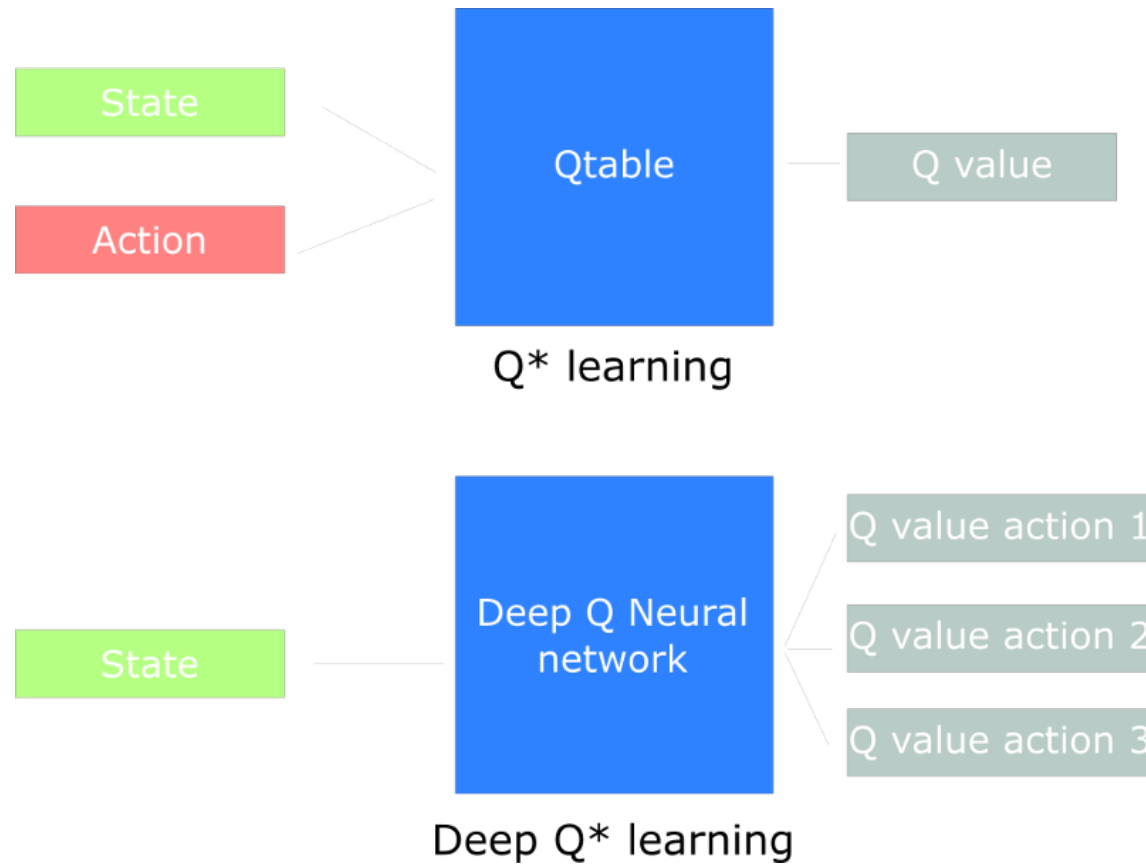
Index

- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 *Deep RL* ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Deep Reinforcement Learning

- Use a Deep Neural Network as a universal approximation function for:
 - state-action-value functions $Q(s, a)$
 - policy functions $a = \pi(s)$

Deep Q-Networks (DQN)



Deep Q-Networks (DQN)

- Universal function approximation

$$Q(s, a) \approx Q(s, a, \mathbf{w})$$

where \mathbf{w} are the weight of the neural network

- Train using temporal difference:

$$Q(s_t, a_t) = \mathbb{E}[r_{t+1} + \gamma \max_a Q(s_{t+1}, \pi(s_{t+1}))]$$

- Therefore:

- Input: (s_t, a_t) (just only s_t)
- Target: $r_{t+1} + \gamma \max_a Q(s_{t+1}, a)$
- Minimize MSE

$$(r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))^2$$

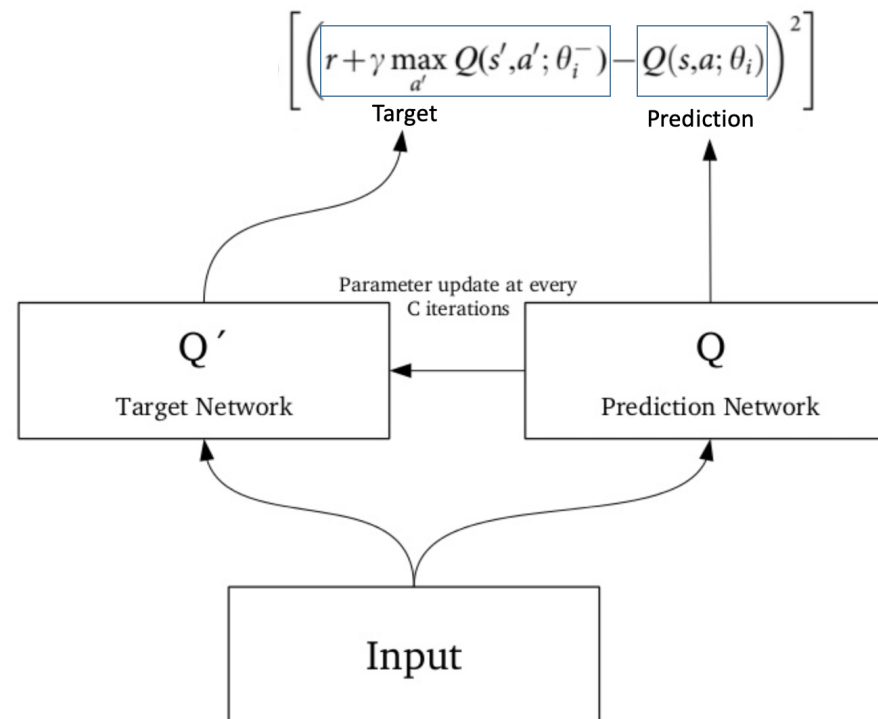
Deep Q-Networks (DQN)

- Experience replay
- re-use the tuples $s_t, a_t, r_{t+1}, s_{t+1}$ as a data set
- multiple passes with the same data is beneficial


Deep Q-Networks (DQN)

- Non-stationary targets

$$\delta = r_{t+1} + \gamma \max_a Q(s_{t+1}, a, \Phi^-) - Q(s_t, a_t, \Phi)$$

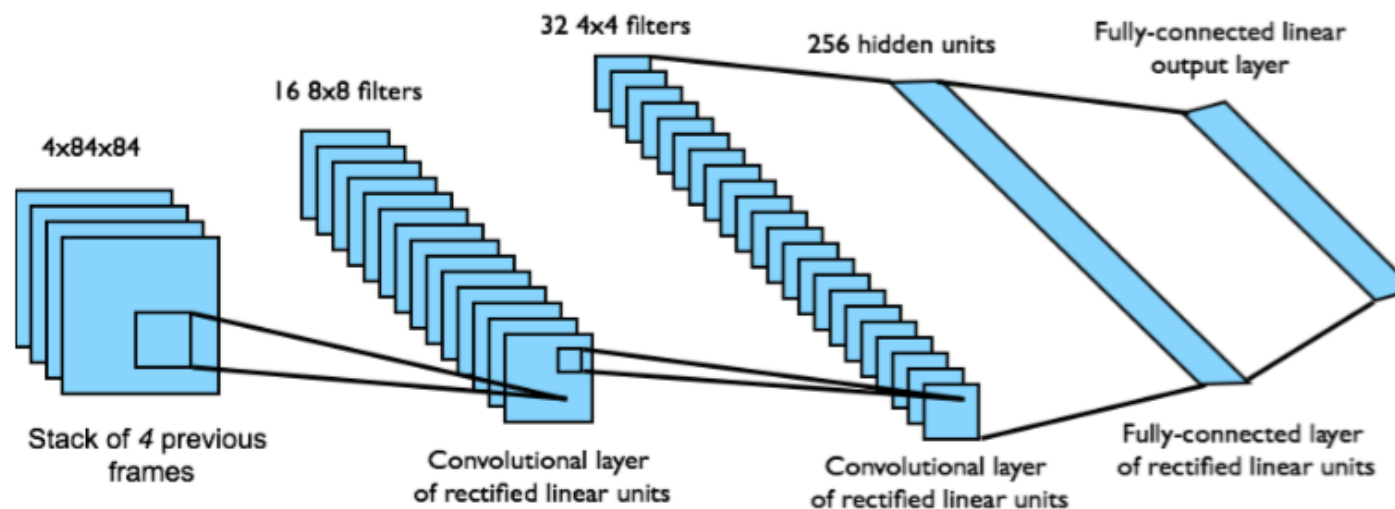


Deep Q-Networks (DQN)



Deep Q-Networks (DQN)

- ▶ End-to-end learning of values $Q(s, a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s, a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step



Network architecture and hyperparameters fixed across all games


Deep Policy Networks

Policy gradients

- Update Φ by gradient ascent in the direction of:


$$R \nabla_{\Phi} \log \pi(a_t \mid s_t; \Phi)$$

- where s could be an image, and a a particular action, e.g moving up



Deep Policy Networks


- High dimensional state representation: a frame with $210 * 160$ potential positions



- State must consider two consecutive frames

Deep Policy Networks

- Delayed reward:



Index


- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 *Generative Models* ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 Generative Adversarial Networks ▷ 49

Generative Models

- Autoencoders
- Denoising Autoencoders
- Variational Autoencoders
- Generative Adversarial Networks

Autoencoder

- Feed-forward neural network that reproduces the input at the output layer




- Training by minimizing the quadratic error

Autoencoders - Topology

- Undercomplete:
 - $\dim(\mathbf{z}) < \dim(\mathbf{x})$
 - compression
 - probably meaningful representation. **Latent space**
- Overcomplete:
 - $\dim(\mathbf{z}) > \dim(\mathbf{x})$,
 - hidden unit just copy input units
 - no meaningful representation
 - have to force sparsity

Autoencoders - weights first layer

- Training over MNIST digits



Denoising-Autoencoders


- Ideal representations should be robust to noise
- A denoising autoencoder is a feed-forward neural network that reproduce the input at the output layer:
 - One input layer \tilde{x}
 - One hidden layer h
 - One output layer layer \hat{x}

where \tilde{x} is a **corrupted** version of x

- The output \hat{x} is computed from a corrupted version of x
- But the loss function compares \hat{x} and x


Denoising-Autoencoders

- MNITS with 25% of noise



Denoising-Autoencoders

- MNITS with 50% of noise

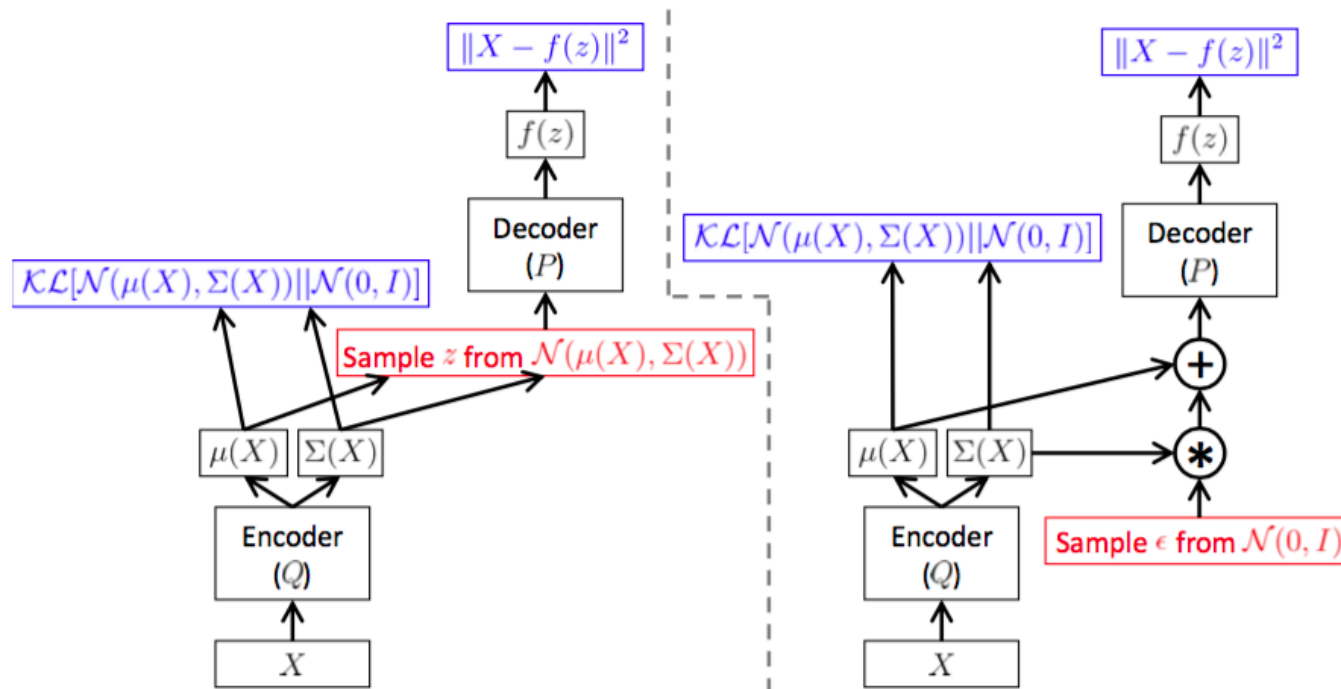


Index

- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 *Variational Autoencoders* ▷ 45
- 7 Generative Adversarial Networks ▷ 49


Variational Autoencoders

- The latent variables are drawn from a prior $p(\mathbf{z})$
- <https://arxiv.org/pdf/1606.05908.pdf>



Variational Autoencoders, Generation process

- Draw from $p(\mathbf{z})$



Variational Autoencoders, Generation process

- Encoder is a Convolutional Network
- Decoder is a De-Convolutional (convolutional transpose) Network



Index

- 1 Reinforcement Learning ▷ 3
- 2 Value Function Algorithms ▷ 12
- 3 Policy Search Algorithms ▷ 22
- 4 Deep RL ▷ 26
- 5 Generative Models ▷ 37
- 6 Variational Autoencoders ▷ 45
- 7 *Generative Adversarial Networks* ▷ 49

Generative Adversarial Networks

- A Generator Network G
- A Discriminator Network D
- During training G is trained to D
- During training D is trained discriminate between real data generated data (fake)

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

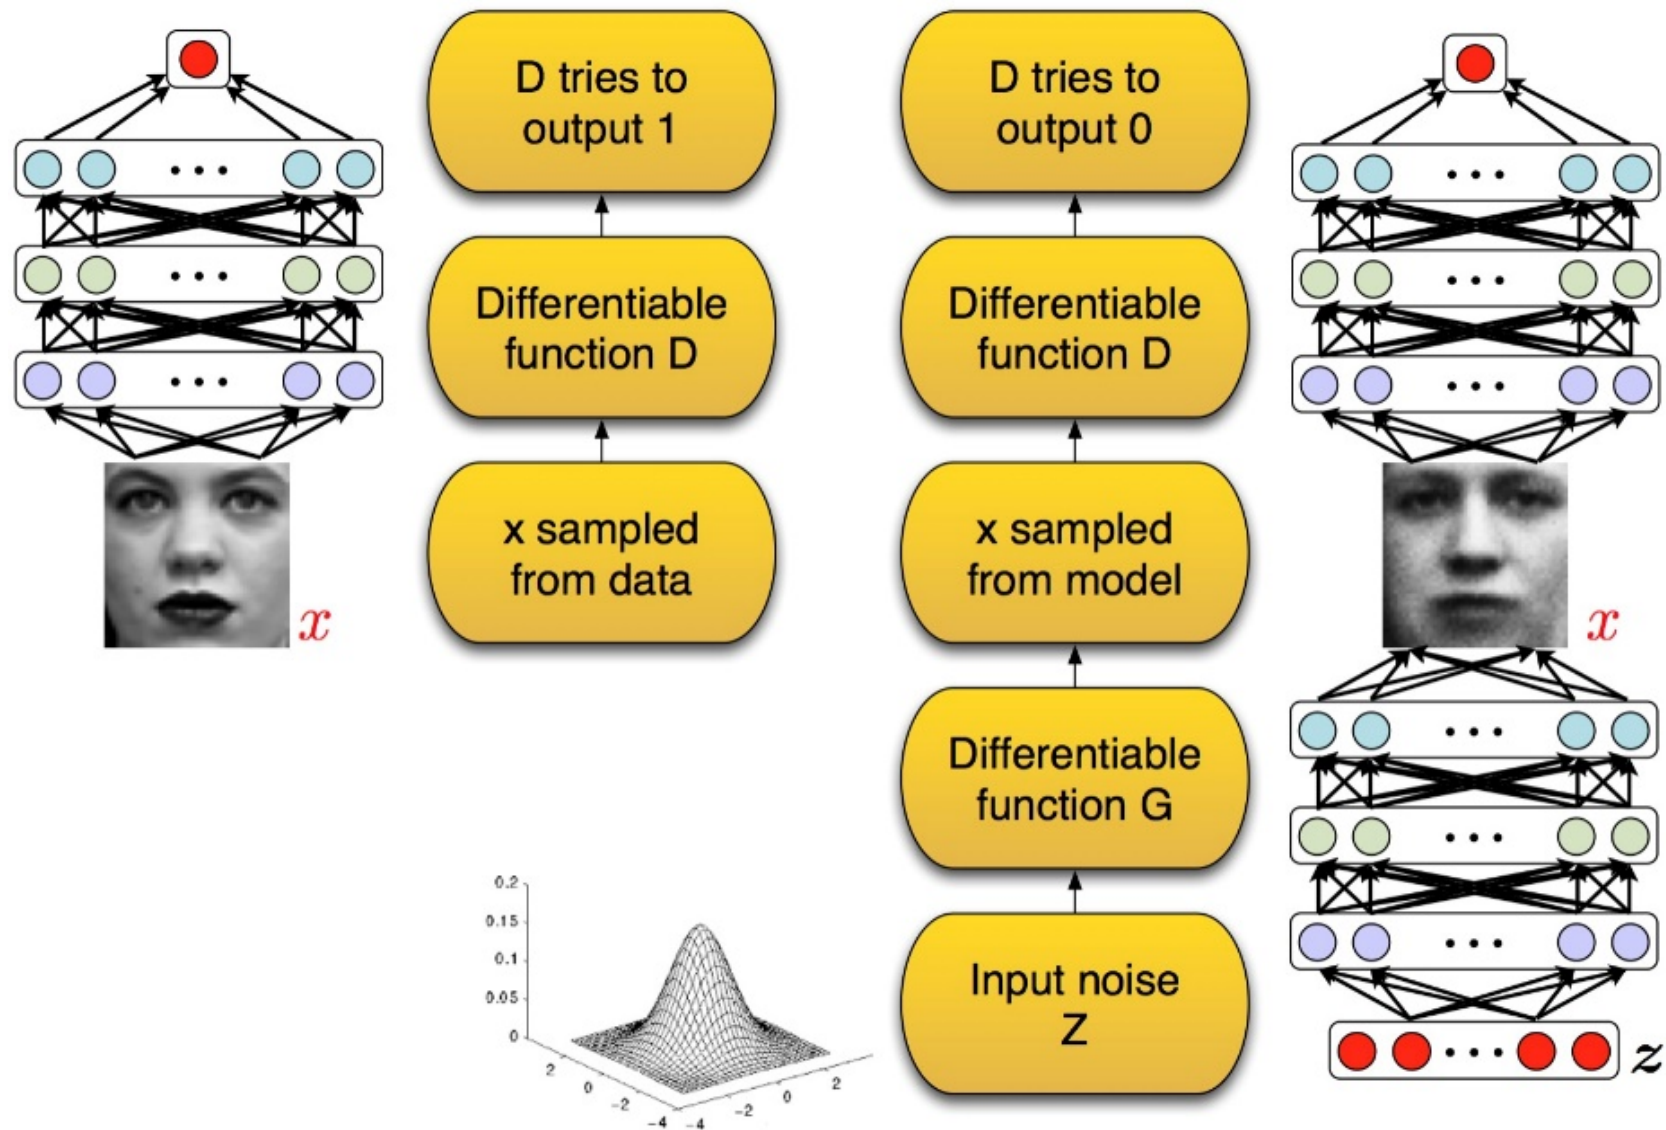
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

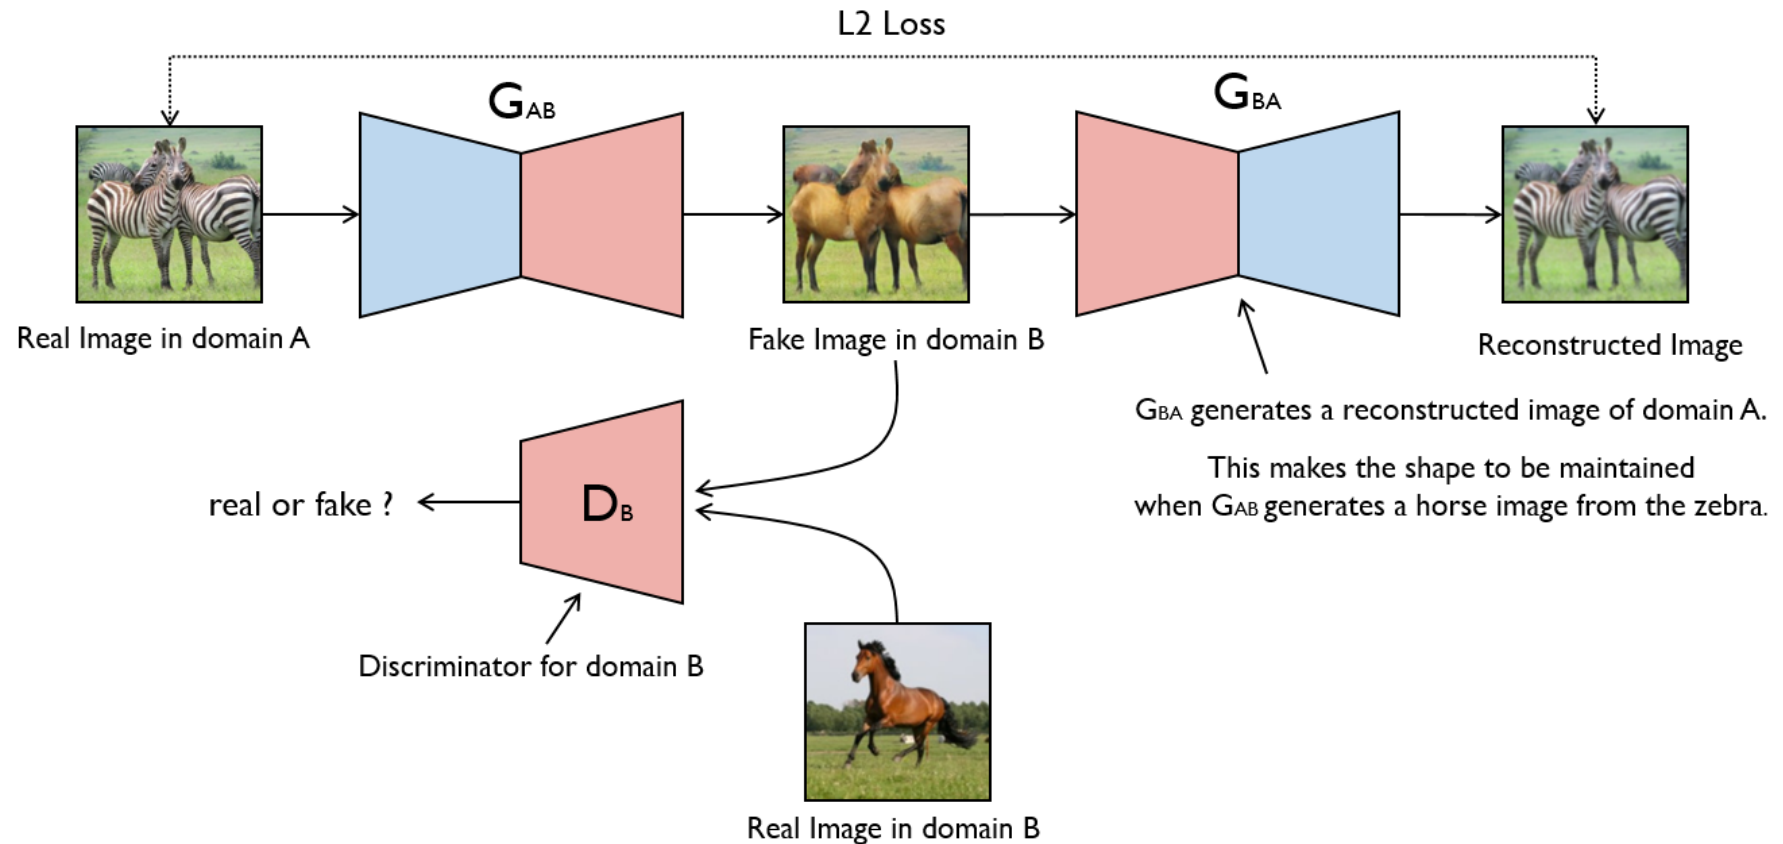
The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

<https://arxiv.org/abs/1406.2661>

Generative Adversarial Networks



CycleGan



Text and Image

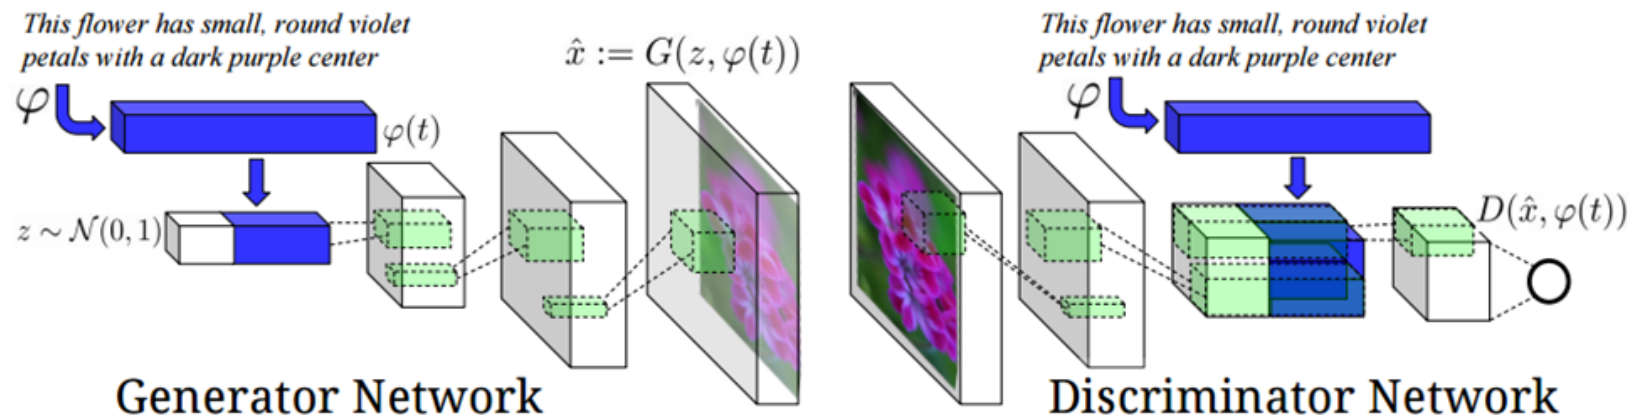


Figure 2. Our text-conditional convolutional GAN architecture. Text encoding $\varphi(t)$ is used by both generator and discriminator. It is projected to a lower-dimensions and depth concatenated with image feature maps for further stages of convolutional processing.

Network Architecture