


Convolutional Networks for Computer Vision

Roberto Paredes

Centro de Investigación
Pattern Recognition and Human Language Technologies
Universidad Politécnica de Valencia


- Introduction
- Convolution Operator
- Pooling Operator
- Reshape Layers
- Special Layers
- Convolutional Networks

- Deep Learning → Bridge the gap between raw representation and categories




Introduction

- ImageNet Challenge




Introduction. Image to text



Introduction. Image to text example




Introduction. Image to text example (unrolled)



- Introduction
- **Convolution Operator**
- Pooling Operator
- Special Layers
- Reshape Layers
- Convolutional Networks

Convolution Operator

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.



Convolution Operator


- The simplest case:
 - Size of input image: $I_R \times I_C$
 - Size of kernel: $k_r \times k_c$
 - Size of output image: $O_R \times O_C$
 - $O_R = (I_R - k_r) + 1$
 - $O_C = (I_C - k_c) + 1$
 - Convolution cost: $O_R \times O_C \times k_r \times k_c$

Convolution Operator

- Padding, same output size than input size
 - $O_R = I_R$
 - $O_C = I_C$
 - Add a frame of $k_r/2 \times k_c/2$ of 0's to the input image
- Stride, jump scanning the input image
- In general:
 - $O_R = \lfloor (I_R + 2 * pad - k_r) / stride_r + 1 \rfloor$
 - $O_C = \lfloor (I_C + 2 * pad - k_c) / stride_c + 1 \rfloor$


Convolution Operator

- The general case:
 - Input: Images (2D)
 - Apply more than one kernel (3D)
 - Output: Maps (3D)



Convolution Operator

- The general case:
 - Input: Maps (3D)
 - Apply more than one kernel (4D)
 - Output: Maps (3D)




Convolution Operator

- Implementation tricks: **LOWERING**
 - A convolution becomes a standard multiplication $I \times K$
- Implementation tricks: **Multi-threading**
 - For instance, split the batch into different threads
- Implementation tricks: **FFT**
 - A convolution is a multiplication in the frequency domain
- Implementation tricks: **Winograd** algorithm:
 - A fast method to obtain the convolution with less multiplication and addition operations

Convolution Operator as a layer

- Input map $M = 3$, $O_R = 32$ and $O_C = 32$
- Kernels $N = 64$, $K_r = 3$ and $k_c = 3$
- By default a convolutional layer as N bias




- Output map sizes?

- Introduction
- Convolution Operator
- **Pooling Operator**
- Reshape Layers
- Special Layers
- Convolutional Networks

- The main goal is to reduce the size of the maps:
 - Reduce the computational cost
 - Deal with multiscale
 - Capture higher level features

Pooling Operator


- Maxpool and Average Pool



- Normally $stride = size$

Pooling Operator


- Results after applying several pooling operators:



- Introduction
- Convolution Operator
- Pooling Operator
- **Reshape Layers**
- Special Layers
- Convolutional Network


Reshape Layers

- The goal of the reshape layer is to present the maps to the next layers Fully Connected layers
- The maps become a raw vector and the spatial relationship is not considered
- After the reshape several hidden layers can be stacked to reach the output layer, conforming a **Convolutional Network**:



Reshape Layers - Exercise

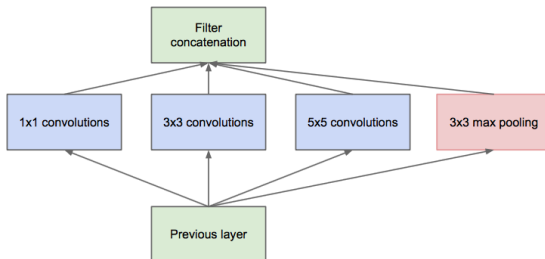
- Given the following CNN, How many parameters (weights) are?
- Where is the big amount of parameters?



- Introduction
- Convolution Operator
- Pooling Operator
- Reshape Layers
- **Special Layers**
- Convolutional Networks


Special Layers - Cat Layers

- A layer that cat **in depth** the maps of the layers that are connected to
- The sizes of the maps must be equivalent
- Used in the Inception Model (GoogleNet)



Special Layers - Agregation Layers


- Similar to cat layers
- A layer that sum the maps of the layers that are connected to
- The size and number of maps must be the same
- Used on the Residual Nets (Microsoft Research)



- Introduction
- Convolution Operator
- Pooling Operator
- Reshape Layers
- Special Layers
- **Convolutional Networks**


Convolutional Networks - AlexNet

- ILSVRC 2012 Winner



Convolutional Networks - GoogLeNet

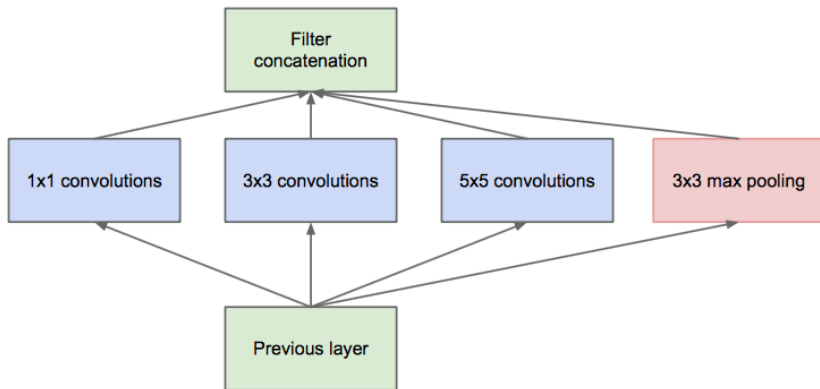
- ILSVRC 2014 Winner



- <http://www.cs.unc.edu/~wliu/papers/GoogLeNet.pdf>

Convolutional Networks - GoogleNet

- Inception model



Convolutional Networks - OxfordNet (VGG)


- ILSVRC2014 2nd, but best single model



- <http://arxiv.org/pdf/1409.1556v6.pdf>

Convolutional Networks - ResidualNet (MSR)

- ILSVRC 2015 Winner



- <http://arxiv.org/pdf/1512.03385v1.pdf>

Convolutional Networks - ResidualNet (MSR)

